



# MongoDB Developer Training



---

# MongoDB Developer Training

*Release 3.2*

**MongoDB, Inc.**

June 02, 2016

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Warm Up . . . . .	2
<b>2</b>	<b>Schema Design</b>	<b>4</b>
2.1	Schema Design Core Concepts . . . . .	4
2.2	Schema Evolution . . . . .	11
2.3	Common Schema Design Patterns . . . . .	15
2.4	Lab: Data Model for an E-Commerce Site . . . . .	19
<b>3</b>	<b>Indexes</b>	<b>21</b>
3.1	Index Fundamentals . . . . .	21
3.2	Compound Indexes . . . . .	27
3.3	Lab: Optimizing an Index . . . . .	32
3.4	Multikey Indexes . . . . .	33
3.5	Hashed Indexes . . . . .	36
3.6	Lab: Finding and Addressing Slow Operations . . . . .	38
3.7	Lab: Using <code>explain()</code> . . . . .	38
<b>4</b>	<b>Replica Sets</b>	<b>39</b>
4.1	Introduction to Replica Sets . . . . .	39
4.2	Write Concern . . . . .	41
4.3	Read Preference . . . . .	46
<b>5</b>	<b>Sharding</b>	<b>48</b>
5.1	Introduction to Sharding . . . . .	48
5.2	Balancing Shards . . . . .	55
5.3	Shard Tags . . . . .	57
5.4	Lab: Setting Up a Sharded Cluster . . . . .	59
<b>6</b>	<b>Aggregation</b>	<b>66</b>
6.1	Aggregation Tutorial . . . . .	66
6.2	Optimizing Aggregation . . . . .	74
6.3	Lab: Aggregation Framework . . . . .	76
<b>7</b>	<b>Application Engineering</b>	<b>78</b>
7.1	Introduction . . . . .	78
7.2	Java Driver Labs (MongoMart) . . . . .	79
<b>8</b>	<b>Reporting Tools and Diagnostics</b>	<b>81</b>
8.1	Performance Troubleshooting . . . . .	81

# 1 Introduction

*Warm Up (page 2)* Activities to get the class started

## 1.1 Warm Up

### Introductions

- Who am I?
- My role at MongoDB
- My background and prior experience

### Getting to Know You

- Who are you?
- What role do you play in your organization?
- What is your background?
- Do you have prior experience with MongoDB?

### MongoDB Experience

- Who has never used MongoDB?
- Who has some experience?
- Who has worked with production MongoDB deployments?
- Who is more of a developer?
- Who is more of operations person?

### 10gen

- MongoDB was initially created in 2008 as part of a hosted application stack.
- The company was originally called 10gen.
- As part of their overarching plan to create the 10gen platform, the company built a database.
- Suddenly everybody said: “I like that! Give me that database!”

## Origin of MongoDB

- 10gen became a database company.
- In 2013, the company rebranded as MongoDB, Inc.
- The founders have other startups to their credit: DoubleClick, ShopWiki, Gilt.
- The motivation for the database came from observing the following pattern with application development.
  - The user base grows.
  - The associated body of data grows.
  - Eventually the application outgrows the database.
  - Meeting performance requirements becomes difficult.

## 2 Schema Design

*Schema Design Core Concepts* (page 4) An introduction to schema design in MongoDB

*Schema Evolution* (page 11) Considerations for evolving a MongoDB schema design over an application's lifetime

*Common Schema Design Patterns* (page 15) Common design patterns for representing 1-1, 1-M, and M-M relationships and tree structures in MongoDB

*Lab: Data Model for an E-Commerce Site* (page 19) Schema design group exercise

### 2.1 Schema Design Core Concepts

#### Learning Objectives

Upon completing this module, students should understand:

- Basic schema design principles for MongoDB
- Tradeoffs for embedded documents in a schema
- Tradeoffs for linked documents in a schema
- The use of array fields as part of a schema design

#### What is a schema?

- Maps concepts and relationships to data
- Sets expectations for the data
- Minimizes overhead of iterative modifications
- Ensures compatibility

#### Example: Normalized Data Model

User:	Book:	Author:
- username	- title	- firstName
- firstName	- isbn	- lastName
- lastName	- language	
	- createdBy	
	- author	

## Example: Denormalized Version

User:	Book:
- username	- title
- firstName	- isbn
- lastName	- language
	- createdBy
	- author
	- firstName
	- lastName

## Schema Design in MongoDB

- Schema is defined at the application-level
- Design is part of each phase in its lifetime
- There is no magic formula

## Three Considerations

- The data your application needs
- Your application's read usage of the data
- Your application's write usage of the data

## Case Study

- A Library Web Application
- Different schemas are possible.

## Author Schema

```
{  "_id": int,
  "firstName": string,
  "lastName": string
}
```

## User Schema

```
{  "_id": int,
  "username": string,
  "password": string
}
```

## Book Schema

```
{  "_id": int,
  "title": string,
  "slug": string,
  "author": int,
  "available": boolean,
  "isbn": string,
  "pages": int,
  "publisher": {
    "city": string,
    "date": date,
    "name": string
  },
  "subjects": [ string, string ],
  "language": string,
  "reviews": [ { "user": int, "text": string },
               { "user": int, "text": string } ]
}
```

## Example Documents: Author

```
{  _id: 1,
  firstName: "F. Scott",
  lastName: "Fitzgerald"
}
```



## Example Documents: User

```
{
  _id: 1,
  username: "emily@10gen.com",
  password: "slsjfk4odk84k209dlkdj90009283d"
}
```

## Example Documents: Book

```
{
  _id: 1,
  title: "The Great Gatsby",
  slug: "9781857150193-the-great-gatsby",
  author: 1,
  available: true,
  isbn: "9781857150193",
  pages: 176,
  publisher: {
    name: "Everyman's Library",
    date: ISODate("1991-09-19T00:00:00Z"),
    city: "London"
  },
  subjects: ["Love stories", "1920s", "Jazz Age"],
  language: "English",
  reviews: [
    { user: 1, text: "One of the best..." },
    { user: 2, text: "It's hard to..." }
  ]
}
```

## Embedded Documents

- AKA sub-documents or embedded objects
- What advantages do they have?
- When should they be used?

## Example: Embedded Documents

```
{
  ...
  publisher: {
    name: "Everyman's Library",
    date: ISODate("1991-09-19T00:00:00Z"),
    city: "London"
  },
  subjects: ["Love stories", "1920s", "Jazz Age"],
  language: "English",
  reviews: [
    { user: 1, text: "One of the best..." },
    { user: 2, text: "It's hard to..." }
  ]
}
```

## Embedded Documents: Pros and Cons

- Great for read performance
- One seek to find the document
- At most, one sequential read to retrieve from disk
- Writes can be slow if constantly adding to objects

## Linked Documents

- What advantages does this approach have?
- When should they be used?

## Example: Linked Documents

```
{
  ...
  author: 1,
  reviews: [
    { user: 1, text: "One of the best..." },
    { user: 2, text: "It's hard to..." }
  ]
}
```

## Linked Documents: Pros and Cons

- More, smaller documents
- Can make queries by ID very simple
- Accessing linked documents requires extra seeks + reads.
- What effect does this have on the system?

## Arrays

- Array of scalars
- Array of documents

## Array of Scalars

```
{
  ...
  subjects: ["Love stories", "1920s", "Jazz Age"],
}
```

## Array of Documents

```
{
  ...
  reviews: [
    { user: 1, text: "One of the best..." },
    { user: 2, text: "It's hard to..." }
  ]
}
```

## Exercise: Users and Book Reviews

Design a schema for users and their book reviews. Usernames are immutable.

- Users
  - username (string)
  - email (string)
- Reviews
  - text (string)
  - rating (integer)
  - created\_at (date)

## Solution A: Users and Book Reviews

Reviews may be queried by user or book

```
// db.users (one document per user)
{
  _id: ObjectId("..."),
  username: "bob",
  email: "bob@example.com"
}

// db.reviews (one document per review)
{
  _id: ObjectId("..."),
  user: ObjectId("..."),
  book: ObjectId("..."),
  rating: 5,
  text: "This book is excellent!",
  created_at: ISODate("2012-10-10T21:14:07.096Z")
}
```

## Solution B: Users and Book Reviews

Optimized to retrieve reviews by user

```
// db.users, one document per user with all reviews
{
  _id: ObjectId("..."),
  username: "bob",
  email: "bob@example.com",
  reviews: [
    {
      book: ObjectId("..."),
      rating: 5,
      text: "This book is excellent!",
      created_at: ISODate("2012-10-10T21:14:07.096Z")
    }
  ]
}
```

## Solution C: Users and Book Reviews

Optimized to retrieve reviews by book

```
// db.users (one document per user)
{
  _id: ObjectId("..."),
  username: "bob",
  email: "bob@example.com"
}

// db.books, one document per book with all reviews
{
  _id: ObjectId("..."),
  // Other book fields...
  reviews: [
    {
      user: ObjectId("..."),
      rating: 5,
      text: "This book is excellent!",
      created_at: ISODate("2014-11-10T21:14:07.096Z")
    }
  ]
}
```

## Store Binary Files in MongoDB with GridFS

- Application may have a requirement for binary file storage
- GridFS is a specification for storing files larger than 16MB in MongoDB
- Handled automatically by most drivers
- “mongofiles” is the command line tool for working with GridFS

## How GridFS Works

- Files are split into chunks
- Default chunk size is 255k
- fs.files collection stores meta data for the file (name, size, etc.)
- fs.chunks collection stores chunks for binary file

## Schema Design Use Cases with GridFS

- Store large video files and stream chunks to a user
- Enterprise assets, replicated across data centers
- Medical record attachments (x-rays, reports, etc.)

## 2.2 Schema Evolution

### Learning Objectives

Upon completing this module, students should understand the basic philosophy of evolving a MongoDB schema during an application's lifetime:

- Development Phase
- Production Phase
- Iterative Modifications

### Development Phase

Support basic CRUD functionality:

- Inserts for authors and books
- Find authors by name
- Find books by basics of title, subject, etc.

## Development Phase: Known Query Patterns

```
// Find authors by last name.
db.authors.createIndex({ "lastName": 1 })

// Find books by slug for detail view
db.books.createIndex({ "slug": 1 })

// Find books by subject (multi-key)
db.books.createIndex({ "subjects": 1 })

// Find books by publisher (index on embedded doc)
db.books.createIndex({ "publisher.name": 1 })
```

## Production Phase

Evolve the schema to meet the application’s read and write patterns.

### Production Phase: Read Patterns

List books by author last name

```
authors = db.authors.find({ lastName: /^f.*/i }, { _id: 1 });
authorIds = authors.map(function(x) { return x._id; });
db.books.find({author: { $in: authorIds }});
```

### Addressing List Books by Last Name

“Cache” the author name in an embedded document.

```
{
  _id: 1,
  title: "The Great Gatsby",
  author: {
    firstName: "F. Scott",
    lastName: "Fitzgerald"
  }
  // Other fields follow...
}
```

Queries are now one step

```
db.books.find({ "author.firstName": /^f.*/i })
```

## Production Phase: Write Patterns

Users can review a book.

```
review = {
  user: 1,
  text: "I thought this book was great!",
  rating: 5
};

db.books.updateOne(
  { _id: 3 },
  { $push: { reviews: review } }
);
```

Caveats:

- Document size limit (16MB)
- Storage fragmentation after many updates/deletes

## Exercise: Recent Reviews

- Display the 10 most recent reviews by a user.
- Make efficient use of memory and disk seeks.

## Solution: Recent Reviews, Schema

Store users' reviews in monthly buckets.

```
// db.reviews (one document per user per month)
{
  _id: "bob-201412",
  reviews: [
    {
      _id: ObjectId("..."),
      rating: 5,
      text: "This book is excellent!",
      created_at: ISODate("2014-12-10T21:14:07.096Z")
    },
    {
      _id: ObjectId("..."),
      rating: 2,
      text: "I didn't really enjoy this book.",
      created_at: ISODate("2014-12-11T20:12:50.594Z")
    }
  ]
}
```

## Solution: Recent Reviews, Update

Adding a new review to the appropriate bucket

```
myReview = {
  _id: ObjectId("..."),
  rating: 3,
  text: "An average read.",
  created_at: ISODate("2012-10-13T12:26:11.502Z")
};

db.reviews.updateOne(
  { _id: "bob-201210" },
  { $push: { reviews: myReview } }
);
```

## Solution: Recent Reviews, Read

Display the 10 most recent reviews by a user

```
cursor = db.reviews.find(
  { _id: /^bob-/ },
  { reviews: { $slice: -10 } }
).sort({ _id: -1 }).batchSize(5);

num = 0;

while (cursor.hasNext() && num < 10) {
  doc = cursor.next();

  for (var i = 0; i < doc.reviews.length && num < 10; ++i, ++num) {
    printjson(doc.reviews[i]);
  }
}
```

## Solution: Recent Reviews, Delete

Deleting a review

```
db.reviews.updateOne(
  { _id: "bob-201210" },
  { $pull: { reviews: { _id: ObjectId("...") } } }
);
```



## 2.3 Common Schema Design Patterns

### Learning Objectives

Upon completing this module students should understand common design patterns for modeling:

- One-to-One Relationships
- One-to-Many Relationships
- Many-to-Many Relationships
- Tree Structures

### One-to-One Relationship

Let's pretend that authors only write one book.

#### One-to-One: Linking

Either side, or both, can track the relationship.

```
db.books.findOne()  
{  
  _id: 1,  
  title: "The Great Gatsby",  
  slug: "9781857150193-the-great-gatsby",  
  author: 1,  
  // Other fields follow...  
}  
  
db.authors.findOne({ _id: 1 })  
{  
  _id: 1,  
  firstName: "F. Scott",  
  lastName: "Fitzgerald"  
  book: 1,  
}
```

#### One-to-One: Embedding

```
db.books.findOne()  
{  
  _id: 1,  
  title: "The Great Gatsby",  
  slug: "9781857150193-the-great-gatsby",  
  author: {  
    firstName: "F. Scott",  
    lastName: "Fitzgerald"  
  }  
  // Other fields follow...  
}
```

## One-to-Many Relationship

In reality, authors may write multiple books.

### One-to-Many: Array of IDs

The “one” side tracks the relationship.

- Flexible and space-efficient
- Additional query needed for non-ID lookups

```
db.authors.findOne()  
{  
  _id: 1,  
  firstName: "F. Scott",  
  lastName: "Fitzgerald",  
  books: [1, 3, 20]  
}
```

### One-to-Many: Single Field with ID

The “many” side tracks the relationship.

```
db.books.find({ author: 1 })  
{  
  _id: 1,  
  title: "The Great Gatsby",  
  slug: "9781857150193-the-great-gatsby",  
  author: 1,  
  // Other fields follow...  
}  
  
{  
  _id: 3,  
  title: "This Side of Paradise",  
  slug: "9780679447238-this-side-of-paradise",  
  author: 1,  
  // Other fields follow...  
}
```

### One-to-Many: Array of Documents

```
db.authors.findOne()  
{  
  _id: 1,  
  firstName: "F. Scott",  
  lastName: "Fitzgerald",  
  books: [  
    { _id: 1, title: "The Great Gatsby" },  
    { _id: 3, title: "This Side of Paradise" }  
  ]  
  // Other fields follow...  
}
```

## Many-to-Many Relationship

Some books may also have co-authors.

### Many-to-Many: Array of IDs on Both Sides

```
db.books.findOne()
{
  _id: 1,
  title: "The Great Gatsby",
  authors: [1, 5]
  // Other fields follow...
}

db.authors.findOne()
{
  _id: 1,
  firstName: "F. Scott",
  lastName: "Fitzgerald",
  books: [1, 3, 20]
}
```

### Many-to-Many: Array of IDs on Both Sides

Query for all books by a given author.

```
db.books.find({ authors: 1 });
```

Query for all authors of a given book.

```
db.authors.find({ books: 1 });
```

### Many-to-Many: Array of IDs on One Side

```
db.books.findOne()
{
  _id: 1,
  title: "The Great Gatsby",
  authors: [1, 5]
  // Other fields follow...
}

db.authors.find({ _id: { $in: [1, 5] } })
{
  _id: 1,
  firstName: "F. Scott",
  lastName: "Fitzgerald"
}
{
  _id: 5,
  firstName: "Unknown",
  lastName: "Co-author"
}
```

## Many-to-Many: Array of IDs on One Side

Query for all books by a given author.

```
db.books.find({ authors: 1 });
```

Query for all authors of a given book.

```
book = db.books.findOne(
  { title: "The Great Gatsby" },
  { authors: 1 }
);

db.authors.find({ _id: { $in: book.authors } });
```

## Tree Structures

E.g., modeling a subject hierarchy.

### Allow users to browse by subject

```
db.subjects.findOne()
{
  _id: 1,
  name: "American Literature",
  sub_category: {
    name: "1920s",
    sub_category: { name: "Jazz Age" }
  }
}
```

- How can you search this collection?
- Be aware of document size limitations
- Benefit from hierarchy being in same document

### Alternative: Parents and Ancestors

```
db.subjects.find()
{ _id: "American Literature" }

{ _id: "1920s",
  ancestors: ["American Literature"],
  parent: "American Literature"
}

{ _id: "Jazz Age",
  ancestors: ["American Literature", "1920s"],
  parent: "1920s"
}

{ _id: "Jazz Age in New York",
  ancestors: ["American Literature", "1920s", "Jazz Age"],
  parent: "Jazz Age"
}
```

## Find Sub-Categories

```
db.subjects.find({ ancestors: "1920s" })
{
  _id: "Jazz Age",
  ancestors: ["American Literature", "1920s"],
  parent: "1920s"
}

{
  _id: "Jazz Age in New York",
  ancestors: ["American Literature", "1920s", "Jazz Age"],
  parent: "Jazz Age"
}
```

## Summary

- Schema design is different in MongoDB.
- Basic data design principles apply.
- It's about your application.
- It's about your data and how it's used.
- It's about the entire lifetime of your application.

## 2.4 Lab: Data Model for an E-Commerce Site

### Introduction

- In this group exercise, we're going to take what we've learned about MongoDB and develop a basic but reasonable data model for an e-commerce site.
- For users of RDBMSs, the most challenging part of the exercise will be figuring out how to construct a data model when joins aren't allowed.
- We're going to model for several entities and features.

### Product Catalog

- **Products.** Products vary quite a bit. In addition to the standard production attributes, we will allow for variations of product type and custom attributes. E.g., users may search for blue jackets, 11-inch macbooks, or size 12 shoes. The product catalog will contain millions of products.
- **Product pricing.** Current prices as well as price histories.
- **Product categories.** Every e-commerce site includes a category hierarchy. We need to allow for both that hierarchy and the many-to-many relationship between products and categories.
- **Product reviews.** Every product has zero or more reviews and each review can receive votes and comments.

## Product Metrics

- **Product views and purchases.** Keep track of the number of times each product is viewed and when each product is purchased.
- **Top 10 lists.** Create queries for top 10 viewed products, top 10 purchased products.
- **Graph historical trends.** Create a query to graph how a product is viewed/purchased over the past.
- **30 days with 1 hour granularity.** This graph will appear on every product page, the query must be very fast.

## Deliverables

- Sample document and schema for each collection
- Queries the application will use
- Index definitions

Break into groups of two or three and work together to create these deliverables.

## 3 Indexes

*Index Fundamentals* (page 21) An introduction to MongoDB indexes

*Compound Indexes* (page 27) Indexes on two or more fields

*Lab: Optimizing an Index* (page 32) Lab on optimizing a compound index

*Multikey Indexes* (page 33) Indexes on array fields

*Hashed Indexes* (page 36) Hashed indexes

*Lab: Finding and Addressing Slow Operations* (page 38) Lab on finding and addressing slow queries

*Lab: Using explain()* (page 38) Lab on using the explain operation to review execution stats

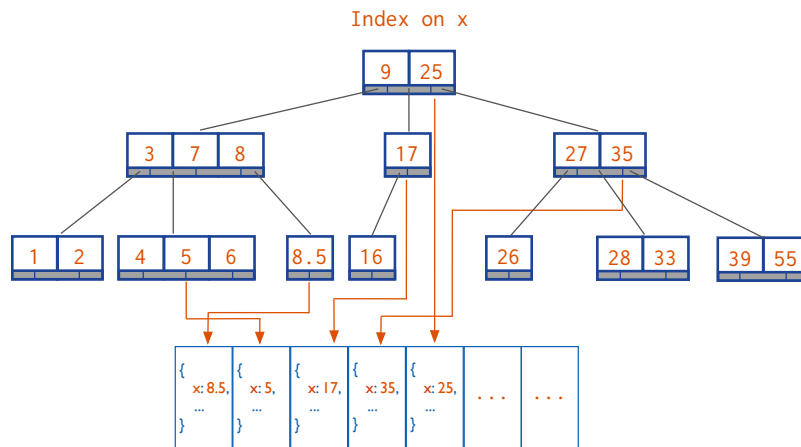
### 3.1 Index Fundamentals

#### Learning Objectives

Upon completing this module students should understand:

- The impact of indexing on read performance
- The impact of indexing on write performance
- How to choose effective indexes
- The utility of specific indexes for particular query patterns

#### Why Indexes?



## Types of Indexes

- Single-field indexes
- Compound indexes
- Multikey indexes
- Geospatial indexes
- Text indexes

## Exercise: Using `explain()`

Let's explore what MongoDB does for the following query by using `explain()`.

We are projecting only `user.name` so that the results are easy to read.

```
db.tweets.find( { "user.followers_count" : 1000 },  
               { "_id" : 0, "user.name": 1 } )
```

```
db.tweets.find( { "user.followers_count" : 1000 } ).explain()
```

## Results of `explain()`

With the default `explain()` verbosity, you will see results similar to the following:

```
{  
  "queryPlanner" : {  
    "plannerVersion" : 1,  
    "namespace" : "twitter.tweets",  
    "indexFilterSet" : false,  
    "parsedQuery" : {  
      "user.followers_count" : {  
        "$eq" : 1000  
      }  
    },  
  },  
}
```

## Results of `explain()` - Continued

```
  "winningPlan" : {  
    "stage" : "COLLSCAN",  
    "filter" : {  
      "user.followers_count" : {  
        "$eq" : 1000  
      }  
    },  
    "direction" : "forward"  
  },  
  "rejectedPlans" : [ ]  
},  
...  
}
```



## **explain() Verbosity Can Be Adjusted**

- **default:** determines the winning query plan but does not execute query
- **executionStats:** executes query and gathers statistics
- **allPlansExecution:** runs all candidate plans to completion and gathers statistics

### **explain("executionStats")**

```
> db.tweets.find( { "user.followers_count" : 1000 } )  
  .explain("executionStats")
```

Now we have query statistics:

```
..  
"executionStats" : {  
  "executionSuccess" : true,  
  "nReturned" : 8,  
  "executionTimeMillis" : 107,  
  "totalKeysExamined" : 0,  
  "totalDocsExamined" : 51428,  
  "executionStages" : {  
    "stage" : "COLLSCAN",  
    "filter" : {  
      "user.followers_count" : {  
        "$eq" : 1000  
      }  
    }  
  },  
}
```

### **explain("executionStats") - Continued**

```
  "nReturned" : 8,  
  "executionTimeMillisEstimate" : 100,  
  "works" : 51430,  
  "advanced" : 8,  
  "needTime" : 51421,  
  "needFetch" : 0,  
  "saveState" : 401,  
  "restoreState" : 401,  
  "isEOF" : 1,  
  "invalidates" : 0,  
  "direction" : "forward",  
  "docsExamined" : 51428  
}  
...  
}
```

## **explain("executionStats") Output**

- `nReturned` : number of documents returned by the query
- `totalDocsExamined` : number of documents touched during the query
- `totalKeysExamined` : number of index keys scanned
- A `totalKeysExamined` or `totalDocsExamined` value much higher than `nReturned` indicates we need a better index
- Based on `.explain()` output, this query would benefit from a better index

## **Other Operations**

In addition to `find()`, we often want to use `explain()` to understand how other operations will be handled.

- `aggregate()`
- `count()`
- `group()`
- `update()`
- `remove()`
- `findAndModify()`
- `insert()`

### **db.<COLLECTION>.explain()**

`db.<COLLECTION>.explain()` returns an `ExplainableCollection`.

```
> var explainable = db.tweets.explain()
> explainable.find( { "user.followers_count" : 1000 } )
```

equivalent to

```
> db.tweets.explain().find( { "user.followers_count" : 1000 } )
```

also equivalent to

```
> db.tweets.find( { "user.followers_count" : 1000 } ).explain()
```

## Using `explain()` for Write Operations

Simulate the number of writes that would have occurred and determine the index(es) used:

```
> db.tweets.explain("executionStats").remove( { "user.followers_count" : 1000 } )

> db.tweets.explain("executionStats").update( { "user.followers_count" : 1000 },
  { $set : { "large_following" : true } }, { multi: true } )
```

## Single-Field Indexes

- Single-field indexes are based on a single field of the documents in a collection.
- The field may be a top-level field.
- You may also create an index on fields in embedded documents.

## Creating an Index

The following creates a single-field index on `user.followers_count`.

```
db.tweets.createIndex( { "user.followers_count" : 1 } )
db.tweets.find( { "user.followers_count" : 1000 } ).explain()
```

`explain()` indicated there will be a substantial performance improvement in handling this type of query.

## Listing Indexes

List indexes for a collection:

```
db.tweets.getIndexes()
```

List index keys:

```
db.tweets.getIndexKeys()
```

## Indexes and Read/Write Performance

- Indexes improve read performance for queries that are supported by the index.
- Inserts will be slower when there are indexes that MongoDB must also update.
- The speed of updates may be improved because MongoDB will not need to do a collection scan to find target documents.
- An index is modified any time a document:
  - Is inserted (applies to *all* indexes)
  - Is deleted (applies to *all* indexes)
  - Is updated in such a way that its indexed field changes

## Index Limitations

- You can have up to 64 indexes per collection.
- You should NEVER be anywhere close to that upper bound.
- Write performance will degrade to unusable at somewhere between 20-30.

## Use Indexes with Care

- Every query should use an index.
- Every index should be used by a query.
- Any write that touches an indexed field will update every index that touches that field.
- Indexes require RAM.
- Be mindful about the choice of key.

## Additional Index Options

- Sparse
- Unique
- Background

## Sparse Indexes in MongoDB

- Sparse indexes only contain entries for documents that have the indexed field.

```
db.<COLLECTION>.createIndex(  
  { field_name : 1 },  
  { sparse : true } )
```

## Defining Unique Indexes

- Enforce a unique constraint on the index
  - On a per-collection basis
- Can't insert documents with a duplicate value for the field
  - Or update to a duplicate value
- No duplicate values may exist prior to defining the index

```
db.<COLLECTION>.createIndex(  
  { field_name : 1 },  
  { unique : true } )
```

## Building Indexes in the Background

- Building indexes in foreground is a blocking operation.
- Background index creation is non-blocking, however, takes longer to build.
- Initially larger, or less compact, than an index built in the foreground.

```
db.<COLLECTION>.createIndex(  
  { field_name : 1 },  
  { background : true } )
```

## 3.2 Compound Indexes

### Learning Objectives

Upon completing this module students should understand:

- What a compound index is.
- How compound indexes are created.
- The importance of considering field order when creating compound indexes.
- How to efficiently handle queries involving some combination of equality matches, ranges, and sorting.
- Some limitations on compound indexes.

### Introduction to Compound Indexes

- It is common to create indexes based on more than one field.
- These are called `compound indexes`.
- You may use up to 31 fields in a compound index.
- You may not use hashed index fields.

### The Order of Fields Matters

Specifically we want to consider how the index will be used for:

- Equality tests, e.g.,

```
db.movies.find( { "budget" : 7, "imdb_rating" : 8 } )
```

- Range queries, e.g.,

```
db.movies.find( { "budget" : 10, "imdb_rating" : { $lt : 9 } } )
```

- Sorting, e.g.,

```
db.movies.find( { "budget" : 10, "imdb_rating" : 6 }  
  ).sort( { "imdb_rating" : -1 } )
```

## Designing Compound Indexes

- Let's look at some guiding principles for building compound indexes.
- These will generally produce a good if not optimal index.
- You can optimize after a little experimentation.
- We will explore this in the context of a running example.

### Example: A Simple Message Board

Requirements:

- Find all messages in a specified timestamp range.
- Select for whether the messages are anonymous or not.
- Sort by rating from highest to lowest.

### Load the Data

```
a = [ { "timestamp" : 1, "username" : "anonymous", "rating" : 3 },
      { "timestamp" : 2, "username" : "anonymous", "rating" : 5 },
      { "timestamp" : 3, "username" : "sam", "rating" : 1 },
      { "timestamp" : 4, "username" : "anonymous", "rating" : 2 },
      { "timestamp" : 5, "username" : "martha", "rating" : 5 } ]
db.messages.insertMany(a)
```

### Start with a Simple Index

Start by building an index on { timestamp : 1 }

```
db.messages.createIndex( { timestamp : 1 }, { name : "myindex" } )
```

Now let's query for messages with timestamp in the range 2 through 4 inclusive.

```
db.messages.find( { timestamp : { $gte : 2, $lte : 4 } } ).explain("executionStats")
```

Analysis:

- Explain plan shows good performance, i.e. totalKeysExamined = n.
- However, this does not satisfy our query.
- Need to query again with { username : "anonymous" } as part of the query.

## Query Adding username

Let's add the user field to our query.

```
db.messages.find( { timestamp : { $gte : 2, $lte : 4 },  
                  username : "anonymous" } ).explain("executionStats")
```

totalKeysExamined > n.

## Include username in Our Index

```
db.messages.dropIndex( "myindex" );  
db.messages.createIndex( { timestamp : 1, username : 1 },  
                        { name : "myindex" } )  
db.messages.find( { timestamp : { $gte : 2, $lte : 4 },  
                  username : "anonymous" } ).explain("executionStats")
```

totalKeysExamined is still > n. Why?

**totalKeysExamined > n**

timestamp	username
1	"anonymous"
2	"anonymous"
3	"sam"
4	"anonymous"
5	"martha"

## A Different Compound Index

Drop the index and build a new one with user.

```
db.messages.dropIndex( "myindex" );  
db.messages.createIndex( { username : 1 , timestamp : 1 },  
                        { name : "myindex" } )  
  
db.messages.find( { timestamp : { $gte : 2, $lte : 4 },  
                  username : "anonymous" } ).explain("executionStats")
```

totalKeysExamined is 2. n is 2.

**totalKeysExamined == n**

username	timestamp
"anonymous"	1
"anonymous"	2
"anonymous"	4
"sam"	2
"martha"	5

## Let Selectivity Drive Field Order

- Order fields in a compound index from most selective to least selective.
- Usually, this means equality fields before range fields.
- When dealing with multiple equality values, start with the most selective.
- If a common range query is more selective instead (rare), specify the range component first.

## Adding in the Sort

Finally, let's add the sort and run the query

```
db.messages.find( {  
    timestamp : { $gte : 2, $lte : 4 },  
    username : "anonymous"  
} ).sort( { rating : -1 } ).explain("executionStats");
```

- Note that the winningPlan includes a SORT stage
- This means that MongoDB had to perform a sort in memory
- In memory sorts can degrade performance significantly
  - Especially if used frequently
  - In-memory sorts that use > 32 MB will abort

## In-Memory Sorts

Let's modify the index again to allow the database to sort for us.

```
db.messages.dropIndex( "myindex" );  
db.messages.createIndex( { username : 1, timestamp : 1, rating : 1 },  
    { name : "myindex" } );  
db.messages.find( {  
    timestamp : { $gte : 2, $lte : 4 },  
    username : "anonymous"  
} ).sort( { rating : -1 } ).explain("executionStats");
```

- The explain plan remains unchanged, because the sort field comes after the range fields.
- The index does not store entries in order by rating.
- Note that this requires us to consider a tradeoff.



## Avoiding an In-Memory Sort

Rebuild the index as follows.

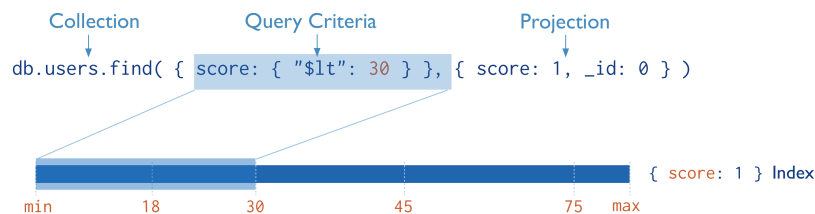
```
db.messages.dropIndex( "myindex" );
db.messages.createIndex( { username : 1, rating : 1, timestamp : 1 },
                        { name : "myindex" } );
db.messages.find( {
    timestamp : { $gte : 2, $lte : 4 },
    username : "anonymous"
} ).sort( { rating : -1 } ).explain("executionStats");
```

- We no longer have an in-memory sort, but need to examine more keys.
- `totalKeysExamined` is 3 and `n` is 2.
- This is the best we can do in this situation and this is fine.
- However, if `totalKeysExamined` is much larger than `n`, this might not be the best index.

## General Rules of Thumb

- Equality before range
- Equality before sorting
- Sorting before range

## Covered Queries



- When a query and projection include only the indexed fields, MongoDB will return results directly from the index.
- There is no need to scan any documents or bring documents into memory.
- These covered queries can be very efficient.

## Exercise: Covered Queries

```
db.testcol.drop()
for (i=1; i<=20; i++) {
  db.testcol.insertOne({ "_id" : i, "title" : i, "name" : i,
    "rating" : i, "budget" : i })
};
db.testcol.createIndex( { "title" : 1, "name" : 1, "rating" : 1 } )

// Not covered because _id is present.
db.testcol.find( { "title" : 3 },
  { "title" : 1, "name" : 1, "rating" : 1 }
).explain("executionStats")

// Not covered because other fields may exist in matching docs.
db.testcol.find( { "title" : 3 },
  { "_id" : 0, "budget" : 0 } ).explain("executionStats")

// Covered query!
db.testcol.find( { "title" : 3 },
  { "_id" : 0, "title" : 1, "name" : 1, "rating" : 1 }
).explain("executionStats")
```

## 3.3 Lab: Optimizing an Index

### Exercise: What Index Do We Need?

Run the the following Javascript file from the handouts.

```
mongo --shell localhost/performance performance.js
```

In the shell that launches execute the following method

```
performance.init()
```

The method above will build a sample data set in the “sensor\_readings” collection. What index is needed for this query?

```
db.sensor_readings.find( { tstamp: { $gte: ISODate("2012-08-01"),
    $lte: ISODate("2012-09-01") },
  active: true } ).limit(3)
```

### Exercise: Avoiding an In-Memory Sort

What index is needed for the following query to avoid an in-memory sort?

```
db.sensor_readings.find( { active: true } ).sort( { tstamp : -1 } )
```

## Exercise: Avoiding an In-Memory Sort, 2

What index is needed for the following query to avoid an in-memory sort?

```
db.sensor_readings.find(
  { x : { $in : [100, 200, 300, 400] } }
).sort( { tstamp : -1 } )
```

## 3.4 Multikey Indexes

### Learning Objectives

Upon completing this module, students should understand:

- What a multikey index is
- When MongoDB will use a multikey index to satisfy a query
- How multikey indexes work
- How multikey indexes handle sorting
- Some limitations on multikey indexes

### Introduction to Multikey Indexes

- A multikey index is an index on an array.
- An index entry is created on each value found in the array.
- Multikey indexes can support primitives, documents, or sub-arrays.
- There is nothing special that you need to do to create a multikey index.
- You create them using `createIndex()` just as you would with an ordinary single-field index.
- If there is an array as a value for an indexed field, the index will be multikey on that field.

### Example: Array of Numbers

```
db.race_results.drop()
db.race_results.createIndex( { "lap_times" : 1 } )
a = [ { "lap_times" : [ 3, 5, 2, 8 ] },
      { "lap_times" : [ 1, 6, 4, 2 ] },
      { "lap_times" : [ 6, 3, 3, 8 ] } ]
db.race_results.insertMany( a )

// Used the index
db.race_results.find( { lap_times : 1 } ).explain()

// One document found.
// Index not used, because it is naive to position.
db.race_results.find( { "lap_times.2" : 3 } ).explain()
```

## Exercise: Array of Documents, Part 1

Create a collection and add an index on the `x` field:

```
db.blog.drop()
b = [ { "comments" : [
  { "name" : "Bob", "rating" : 1 },
  { "name" : "Frank", "rating" : 5.3 },
  { "name" : "Susan", "rating" : 3 } ] },
  { "comments" : [
    { name : "Megan", "rating" : 1 } ] },
  { "comments" : [
    { "name" : "Luke", "rating" : 1.4 },
    { "name" : "Matt", "rating" : 5 },
    { "name" : "Sue", "rating" : 7 } ] } ]
db.blog.insertMany(b)

db.blog.createIndex( { "comments" : 1 } )
// vs
db.blog.createIndex( { "comments.rating" : 1 } )

// for this query
db.blog.find( { "comments.rating" : 5 } )
```

## Exercise: Array of Documents, Part 2

For each of the three queries below:

- How many documents will be returned?
- Will it use our multi-key index? Why or why not?
- If a query will not use the index, which index will it use?

```
db.blog.find( { "comments" : { "name" : "Bob", "rating" : 1 } } )
db.blog.find( { "comments" : { "rating" : 1 } } )
db.blog.find( { "comments.rating" : 1 } )
```

## Exercise: Array of Arrays, Part 1

Add some documents and create an index simulating a player in a game moving on an X,Y grid.

```
db.player.drop()
db.player.createIndex( { "last_moves" : 1 } )
c = [ { "last_moves" : [ [ 1, 2 ], [ 2, 3 ], [ 3, 4 ] ] },
  { "last_moves" : [ [ 3, 4 ], [ 4, 5 ] ] },
  { "last_moves" : [ [ 4, 5 ], [ 5, 6 ] ] },
  { "last_moves" : [ [ 3, 4 ] ] },
  { "last_moves" : [ [ 4, 5 ] ] } ]
db.player.insertMany(c)
db.player.find()
```

## Exercise: Array of Arrays, Part 2

For each of the queries below:

- How many documents will be returned?
- Does the query use the multi-key index? Why or why not?
- If the query does not use the index, what is an index it could use?

```
db.player.find( { "last_moves" : [ 3, 4 ] } )
db.player.find( { "last_moves" : 3 } )
db.player.find( { "last_moves.1" : [ 4, 5 ] } )
db.player.find( { "last_moves.2" : [ 2, 3 ] } )
```

## How Multikey Indexes Work

- Each array element is given one entry in the index.
- So an array with 17 elements will have 17 entries – one for each element.
- Multikey indexes can take up much more space than standard indexes.

## Multikey Indexes and Sorting

- If you sort using a multikey index:
  - A document will appear at the first position where a value would place the document.
  - It will not appear multiple times.
- This applies to array values generally.
- It is not a specific property of multikey indexes.

## Exercise: Multikey Indexes and Sorting

```
db.testcol.drop()
a = [ { x : [ 1, 11 ] }, { x : [ 2, 10 ] }, { x : [ 3 ] },
      { x : [ 4 ] }, { x : [ 5 ] } ]
db.testcol.insert(a)

db.testcol.createIndex( { x : 1 } )

// x : [ 1, 11 ] array comes first. It contains the lowest value.
db.testcol.find().sort( { x : 1 } )

// x : [ 1, 11 ] array still comes first. Contains the highest value.
db.testcol.find().sort( { x : -1 } )
```

## Limitations on Multikey Indexes

- You cannot create a compound index using more than one array-valued field.
- This is because of the combinatorics.
- For a compound index on two array-valued fields you would end up with  $N * M$  entries for one document.
- You cannot have a hashed multikey index.
- You cannot have a shard key use a multikey index.
- We discuss shard keys in another module.
- The index on the `_id` field cannot become a multikey index.

## Example: Multikey Indexes on Multiple Fields

```
db.testcol.drop()
db.testcol.createIndex( { x : 1, y : 1 } )

// no problems yet
db.testcol.insertOne( { _id : 1, x : 1, y : 1 } )

// still OK
db.testcol.insertOne( { _id : 2, x : [ 1, 2 ], y : 1 } )

// still OK
db.testcol.insertOne( { _id : 3, x : 1, y : [ 1, 2 ] } )

// Won't work
db.testcol.insertOne( { _id : 4, x : [ 1, 2 ], y : [ 1, 2 ] } )
```

## 3.5 Hashed Indexes

### Learning Objectives

Upon completing this module, students should understand:

- What a hashed index is
- When to use a hashed index

## What is a Hashed Index?

- Hashed indexes are based on field values like any other index.
- The difference is that the values are hashed and it is the hashed value that is indexed.
- The hashing function collapses sub-documents and computes the hash for the entire value.
- MongoDB can use the hashed index to support equality queries.
- Hashed indexes do not support multi-key indexes, i.e. indexes on array fields.
- Hashed indexes do not support range queries.

## Why Hashed Indexes?

- In MongoDB, the primary use for hashed indexes is to support sharding a collection using a hashed shard key.
- In some cases, the field we would like to use to shard data would make it difficult to scale using sharding.
- Using a hashed shard key to shard a collection ensures an even distribution of data and overcomes this problem.
- See [Shard a Collection Using a Hashed Shard Key](#)<sup>1</sup> for more details.
- We discuss sharding in detail in another module.

## Limitations

- You may not create compound indexes that have hashed index fields.
- You may not specify a unique constraint on a hashed index.
- You can create both a hashed index and a non-hashed index on the same field.

## Floating Point Numbers

- MongoDB hashed indexes truncate floating point numbers to 64-bit integers before hashing.
- Do not use a hashed index for floating point numbers that cannot be reliably converted to 64-bit integers.
- MongoDB hashed indexes do not support floating point values larger than  $2^{53}$ .

---

<sup>1</sup><http://docs.mongodb.org/manual/tutorial/shard-collection-with-a-hashed-shard-key/>

## Creating a Hashed Index

Create a hashed index using an operation that resembles the following. This operation creates a hashed index for the active collection on the a field.

```
db.active.createIndex( { a: "hashed" } )
```

## 3.6 Lab: Finding and Addressing Slow Operations

### Exercise: Determine Indexes Needed

- In a mongo shell run `performance.b()`. This will run in an infinite loop printing some output as it runs various statements against the server.
- Now imagine we have detected a performance problem and suspect there is a slow operation running.
- Find the slow operation and terminate it. Every slow operation is assumed to run for 100ms or more.
- In order to do this, open a second window (or tab) and run a second instance of the mongo shell.
- What indexes can we introduce to make the slow queries more efficient? Disregard the index created in the previous exercise.

## 3.7 Lab: Using `explain()`

### Exercise: `explain("executionStats")`

Drop all indexes from previous exercises:

```
mongo performance
> db.sensor_readings.dropIndexes()
```

Create an index for the “active” field:

```
db.sensor_readings.createIndex({ "active" : 1 } )
```

How many index entries and documents are examined for the following query? How many results are returned?

```
db.sensor_readings.find(
    { "active": false, "_id": { $gte: 99, $lte: 1000 } }
).explain("executionStats")
```



## 4 Replica Sets

*Introduction to Replica Sets (page 39)* An introduction to replication and replica sets

*Write Concern (page 41)* Balancing performance and durability of writes

*Read Preference (page 46)* Configuring clients to read from specific members of a replica set

### 4.1 Introduction to Replica Sets

#### Learning Objectives

Upon completing this module, students should understand:

- Striking the right balance between cost and redundancy
- The many scenarios replication addresses and why
- How to avoid downtime and data loss using replication

#### Use Cases for Replication

- High Availability
- Disaster Recovery
- Functional Segregation

#### High Availability (HA)

- Data still available following:
  - Equipment failure (e.g. server, network switch)
  - Datacenter failure
- This is achieved through automatic failover.

#### Disaster Recovery (DR)

- We can duplicate data across:
  - Multiple database servers
  - Storage backends
  - Datacenters
- Can restore data from another node following:
  - Hardware failure
  - Service interruption

## Functional Segregation

There are opportunities to exploit the topology of a replica set:

- Based on physical location (e.g. rack or datacenter location)
- For analytics, reporting, data discovery, system tasks, etc.
- For backups

## Large Replica Sets

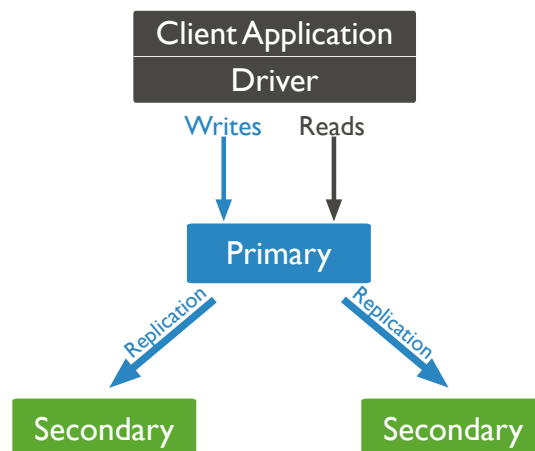
Functional segregation can be further exploited by using large replica sets.

- 50 node replica set limit
- Useful for deployments with a large number of data centers or offices
- Read only workloads can position secondaries in data centers around the world (closer to application servers)

## Replication is Not Designed for Scaling

- Can be used for scaling reads, but generally not recommended.
- Drawbacks include:
  - Eventual consistency
  - Not scaling writes
  - Potential system overload when secondaries are unavailable
- Consider sharding for scaling reads and writes.

## Replica Sets



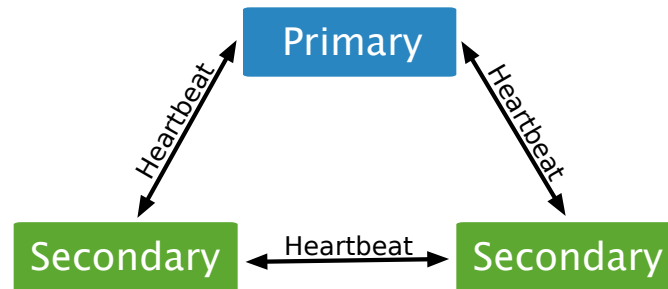
## Primary Server

- Clients send writes the primary only.
- MongoDB, Inc. maintains client drivers in many programming languages like Java, C#, Python, Ruby, and PHP.
- MongoDB drivers are replica set aware.

## Secondaries

- A secondary replicates operations from another node in the replica set.
- Secondaries usually replicate from the primary.
- Secondaries may also replicate from other secondaries. This is called replication chaining.
- A secondary may become primary as a result of a failover scenario.

## Heartbeats



## The Oplog

- The operations log, or oplog, is a special capped collection that is the basis for replication.
- The oplog maintains one entry for each document affected by every write operation.
- Secondaries copy operations from the oplog of their sync source.

## 4.2 Write Concern

### Learning Objectives

Upon completing this module students should understand:

- How and when rollback occurs in MongoDB.
- The tradeoffs between durability and performance.
- Write concern as a means of ensuring durability in MongoDB.
- The different levels of write concern.

## What happens to the write?

- A write is sent to a primary.
- The primary acknowledges the write to the client.
- The primary then becomes unavailable before a secondary can replicate the write

## Answer

- Another member might be elected primary.
- It will not have the last write that occurred before the previous primary became unavailable.
- When the previous primary becomes available again:
  - It will note it has writes that were not replicated.
  - It will put these writes into a `rollback file`.
  - A human will need to determine what to do with this data.
- This is default behavior in MongoDB and can be controlled using `write concern`.


## Balancing Durability with Performance

- The previous scenario is a specific instance of a common distributed systems problem.
- For some applications it might be acceptable for writes to be rolled back.
- Other applications may have varying requirements with regard to durability.
- Tunable write concern:
  - Make critical operations persist to an entire MongoDB deployment.
  - Specify replication to fewer nodes for less important operations.

## Defining Write Concern

- Clients may define the write concern per write operation, if necessary.
- Standardize on specific levels of write concerns for different classes of writes.
- In the discussion that follows we will look at increasingly strict levels of write concern.


**Write Concern: { w : 1 }**



**Example: { w : 1 }**

```
db.edges.insertOne( { from : "tom185", to : "mary_p" },  
                    { writeConcern : { w : 1 } } )
```

**Write Concern: { w : 2 }**



### Example: { w : 2 }

```
db.customer.updateOne( { user : "mary_p" },
  { $push : { shoppingCart:
    { _id : 335443, name : "Brew-a-cup",
      price : 45.79 } } },
  { writeConcern : { w : 2 } } )
```

### Other Write Concerns

- You may specify any integer as the value of the w field for write concern.
- This guarantees that write operations have propagated to the specified number of members.
- E.g., { w : 3 }, { w : 4 }, etc.

### Write Concern: { w : "majority" }

- Ensures the primary completed the write (in RAM).
- Ensures write operations have propagated to a majority of a replica set's **voting** members.
- Avoids hard coding assumptions about the size of your replica set into your application.
- Using majority trades off performance for durability.
- It is suitable for critical writes and to avoid rollbacks.

### Example: { w : "majority" }

```
db.products.updateOne({ _id : 335443 },
  { $inc : { inStock : -1 } },
  { writeConcern : { w : "majority" } })
```

### Quiz: Which write concern?

Suppose you have a replica set with 7 data nodes. Your application has critical inserts for which you do not want rollbacks to happen. Secondaries may be taken down from to time for maintenance, leaving you with a potential 4 server replica set. Which write concern is best suited for these critical inserts?

- { w : 1 }
- { w : 2 }
- { w : 3 }
- { w : 4 }
- { w : "majority" }

## Further Reading

See [Write Concern Reference](#)<sup>2</sup> for more details on write concern configurations, including setting timeouts and identifying specific replica set members that must acknowledge writes (i.e. [tag sets](#)<sup>3</sup>).

---

<sup>2</sup><http://docs.mongodb.org/manual/reference/write-concern>

<sup>3</sup><http://docs.mongodb.org/manual/tutorial/configure-replica-set-tag-sets/#replica-set-configuration-tag-sets>

## 4.3 Read Preference

### What is Read Preference?

- Read preference allows you to specify the nodes in a replica set to read from.
- Clients only read from the primary by default.
- There are some situations in which a client may want to read from:
  - Any secondary
  - A specific secondary
  - A specific type of secondary
- Only read from a secondary if you can tolerate possibly stale data, as not all writes might have replicated.

### Use Cases

- Running systems operations without affecting the front-end application.
- Providing local reads for geographically distributed applications.
- Maintaining availability during a failover.

### Not for Scaling

- In general, do *not* read from secondaries to provide extra capacity for reads.
- [Sharding](http://docs.mongodb.org/manual/sharding)<sup>4</sup> increases read and write capacity by distributing operations across a group of machines.
- Sharding is a better strategy for adding capacity.

### Read Preference Modes

MongoDB drivers support the following read preferences. Note that hidden nodes will never be read from when connected via the replica set.

- **primary**: Default. All operations read from the primary.
- **primaryPreferred**: Read from the primary but if it is unavailable, read from secondary members.
- **secondary**: All operations read from the secondary members of the replica set.
- **secondaryPreferred**: Read from secondary members but if no secondaries are available, read from the primary.
- **nearest**: Read from member of the replica set with the least network latency, regardless of the member's type.

---

<sup>4</sup><http://docs.mongodb.org/manual/sharding>



## Tag Sets

- There is also the option to use tag sets.
- You may tag nodes such that queries that contain the tag will be routed to one of the servers with that tag.
- This can be useful for running reports, say for a particular data center or nodes with different hardware (e.g. hard disks vs SSDs).

For example, in the mongo shell:

```
conf = rs.conf()
conf.members[0].tags = { dc : "east", use : "production" }
conf.members[1].tags = { dc : "east", use : "reporting" }
conf.members[2].tags = { use : "production" }
rs.reconfig(conf)
```

## 5 Sharding

*Introduction to Sharding (page 48)* An introduction to sharding

*Balancing Shards (page 55)* Chunks, the balancer, and their role in a sharded cluster

*Shard Tags (page 57)* How tag-based sharding works

*Lab: Setting Up a Sharded Cluster (page 59)* Deploying a sharded cluster

### 5.1 Introduction to Sharding

#### Learning Objectives

Upon completing this module, students should understand:

- What problems sharding solves
- When sharding is appropriate
- The importance of the shard key and how to choose a good one
- Why sharding increases the need for redundancy

#### Contrast with Replication

- In an earlier module, we discussed Replication.
- This should never be confused with sharding.
- Replication is about high availability and durability.
  - Taking your data and constantly copying it
  - Being ready to have another machine step in to field requests.


#### Sharding is Concerned with Scale

- What happens when a system is unable to handle the application load?
- It is time to consider scaling.
- There are 2 types of scaling we want to consider:
  - Vertical scaling
  - Horizontal scaling

## Vertical Scaling

- Adding more RAM, faster disks, etc.
- When is this the solution?
- First, consider a concept called the *working set*.

## The Working Set



## Limitations of Vertical Scaling

- There is a limit to how much RAM one machine can support.
- There are other bottlenecks such as I/O, disk access and network.
- Cost may limit our ability to scale up.
- There may be requirements to have a large working set that no single machine could possibly support.
- This is when it is time to scale horizontally.


## Sharding Overview

- MongoDB enables you to scale horizontally through sharding.
- Sharding is about adding more capacity to your system.
- MongoDB's sharding solution is designed to perform well on commodity hardware.
- The details of sharding are abstracted away from applications.
- Queries are performed the same way as if sending operations to a single server.
- Connections work the same by default.

## When to Shard

- If you have more data than one machine can hold on its drives
- If your application is write heavy and you are experiencing too much latency.
- If your working set outgrows the memory you can allocate to a single machine.

## Dividing Up Your Dataset



## Sharding Concepts

To understanding how sharding works in MongoDB, we need to understand:

- Shard Keys
- Chunks

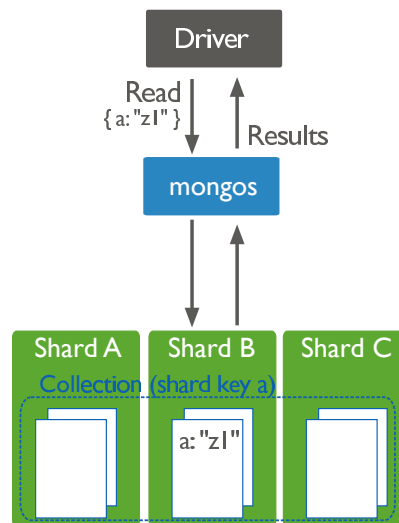
### Shard Key

- You must define a shard key for a sharded collection.
- Based on one or more fields (like an index)
- Shard key defines a space of values
- Think of the key space like points on a line
- A key range is a segment of that line

## Shard Key Ranges

- A collection is partitioned based on shard key ranges.
- The shard key determines where documents are located in the cluster.
- It is used to route operations to the appropriate shard.
- For reads and writes
- Once a collection is sharded, you cannot change a shard key.


## Targeted Query Using Shard Key



## Chunks

- MongoDB partitions data into `chunks` based on shard key ranges.
- This is bookkeeping metadata.
- MongoDB attempts to keep the amount of data balanced across shards.
- This is achieved by migrating chunks from one shard to another as needed.
- There is nothing in a document that indicates its chunk.
- The document does not need to be updated if its assigned chunk changes.

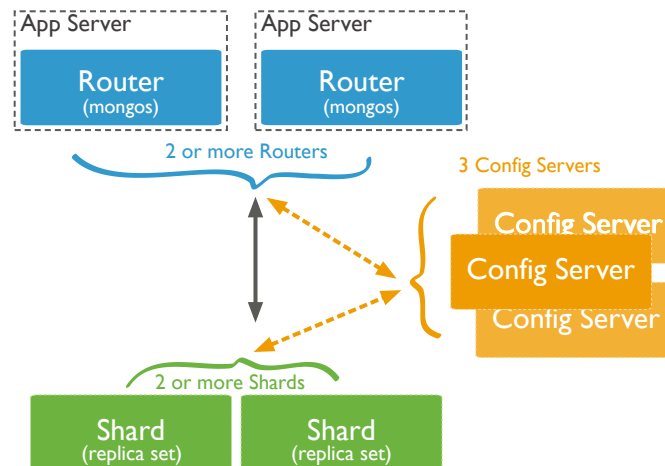
## Sharded Cluster Architecture



### Mongos

- A mongos is responsible for accepting requests and returning results to an application driver.
- In a sharded cluster, nearly all operations go through a mongos.
- A sharded cluster can have as many mongos routers as required.
- It is typical for each application server to have one mongos.
- Always use more than one mongos to avoid a single point of failure.

### Config Servers



## Config Server Hardware Requirements

- Quality network interfaces
- A small amount of disk space (typically a few GB)
- A small amount of RAM (typically a few GB)
- The larger the sharded cluster, the greater the config server hardware requirements.

## Possible Imbalance?

- Depending on how you configure sharding, data can become unbalanced on your sharded cluster.
  - Some shards might receive more inserts than others.
  - Some shards might have documents that grow more than those in other shards.
- This may result in too much load on a single shard.
  - Reads and writes
  - Disk activity
- This would defeat the purpose of sharding.

## Balancing Shards

- If a chunk grows too large MongoDB will split it into two chunks.
- The MongoDB balancer keeps chunks distributed across shards in equal numbers.
- However, a balanced sharded cluster depends on a good shard key.

## With a Good Shard Key

You might easily see that:

- Reads hit only 1 or 2 shards per query.
- Writes are distributed across all servers.
- Your disk usage is evenly distributed across shards.
- Things stay this way as you scale.

## **With a Bad Shard Key**

You might see that:

- Your reads hit every shard.
- Your writes are concentrated on one shard.
- Most of your data is on just a few shards.
- Adding more shards to the cluster will not help.

## **Choosing a Shard Key**

Generally, you want a shard key:

- That has high cardinality
- That is used in the majority of read queries
- For which the values read and write operations use are randomly distributed
- For which the majority of reads are routed to a particular server

## **More Specifically**

- Your shard key should be consistent with your query patterns.
- If reads usually find only one document, you only need good cardinality.
- If reads retrieve many documents:
  - Your shard key supports locality
  - Matching documents will reside on the same shard

## **Cardinality**

- A good shard key will have high cardinality.
- A relatively small number of documents should have the same shard key.
- Otherwise operations become isolated to the same server.
- Because documents with the same shard key reside on the same shard.
- Adding more servers will not help.
- Hashing will not help.



## Non-Monotonic

- A good shard key will generate new values non-monotonically.
- Datetimes, counters, and ObjectIds make bad shard keys.
- Monotonic shard keys cause all inserts to happen on the same shard.
- Hashing will solve this problem.
- However, doing range queries with a hashed shard key will perform a scatter-gather query across the cluster.

## Shards Should be Replica Sets

- As the number of shards increases, the number of servers in your deployment increases.
- This increases the probability that one server will fail on any given day.
- With redundancy built into each shard you can mitigate this risk.

## 5.2 Balancing Shards

### Learning Objectives

Upon completing this module students should understand:

- Chunks and the balancer
- The status of chunks in a newly sharded collection
- How chunk splits automatically occur
- Advantages of pre-splitting chunks
- How the balancer works

### Chunks and the Balancer

- Chunks are groups of documents.
- The shard key determines which chunk a document will be contained in.
- Chunks can be split when they grow too large.
- The balancer decides where chunks go.
- It handles migrations of chunks from one server to another.


## Chunks in a Newly Sharded Collection

- The range of a chunk is defined by the shard key values of the documents the chunk contains.
- When a collection is sharded it starts with just one chunk.
- The first chunk for a collection will have the range:

```
{ $minKey : 1 } to { $maxKey : 1 }
```

- All shard key values from the smallest possible to the largest fall in this chunk's range.

## Chunk Splits



## Pre-Splitting Chunks

- You may pre-split data before loading data into a sharded cluster.
- Pre-splitting is useful if:
  - You plan to do a large data import early on
  - You expect a heavy initial server load and want to ensure writes are distributed

## Start of a Balancing Round

- A balancing round may be initiated by any mongos in the cluster.
- This happens when the difference in the number of chunks between two shards becomes too large.
- Specifically, the difference between the shard with the most chunks and the shard with the fewest.
- A balancing round starts when the imbalance reaches:
  - 2 when the cluster has < 20 chunks
  - 4 when the cluster has 20-79 chunks
  - 8 when the cluster has 80+ chunks

## Balancing is Resource Intensive

- Chunk migration requires copying all the data in the chunk from one shard to another.
- MongoDB can migrate only a single chunk at a time.
- MongoDB creates splits only after an insert operation.
- For these reasons, it is possible to define a balancing window to ensure the balancer will only run during scheduled times.

## Chunk Migration Steps

1. The balancer process sends the `moveChunk` command to the source shard.
2. The source shard continues to process reads/writes for that chunk during the migration.
3. The destination shard requests documents in the chunk and begins receiving copies.
4. After receiving all documents, the destination shard receives any changes to the chunk.
5. Then the destination shard tells the config db that it has the chunk.
6. The destination shard will now handle all reads/writes.
7. The source shard deletes its copy of the chunk.

## Concluding a Balancing Round

- Each chunk will move:
  - From the shard with the most chunks
  - To the shard with the fewest
- A balancing round ends when all shards differ by at most one chunk.

## 5.3 Shard Tags

### Learning Objectives

Upon completing this module students should understand:

- The purpose for shard tags
- Advantages of using shard tags
- Potential drawbacks of shard tags

## Tags - Overview

- Shard tags allow you to “tie” data to one or more shards.
- A shard tag describes a range of shard key values.
- If a chunk is in the shard tag range, it will live on a shard with that tag.

### Example: DateTime

- Documents older than one year need to be kept, but are rarely used.
- You tag those ranges as “LTS” for Long Term Storage.
- Tag specific shards to hold LTS documents.
- These shards can be on cheaper, slower machines.
- Invest in high-performance servers for more frequently accessed data.

### Example: Location

- You are required to keep certain data in its home country.
- You include the country in the shard tag.
- Maintain data centers within each country that house the appropriate shards.
- Meets the country requirement but allows all servers to be part of the same system.

### Example: Premium Tier

- You have customers who want to pay for a “premium” tier.
- The shard key permits you to distinguish one customer’s documents from all others.
- Tag the document ranges for each customer so that their documents will be located on shards of the appropriate tier.
- Shards tagged as premium tier run on high performance servers.
- Other shards run on commodity hardware.
- See [Manage Shard Tags](#)<sup>5</sup>

---

<sup>5</sup><http://docs.mongodb.org/manual/tutorial/administer-shard-tags/>

## Tags - Caveats

- Because tagged chunks will only be on certain servers, if you tag more than those servers can handle, you'll have a problem.
  - You're not only worrying about your overall server load, you're worrying about server load for each of your tags.
- Your chunks will evenly distribute themselves across the available chunks. You cannot control things more fine grained than your tags.

## 5.4 Lab: Setting Up a Sharded Cluster

### Learning Objectives

Upon completing this module students should understand:

- How to set up a sharded cluster including:
  - Replica Sets as Shards
  - Config Servers
  - Mongos processes
- How to enable sharding for a database
- How to shard a collection
- How to determine where data will go

### Our Sharded Cluster

- In this exercise, we will set up a cluster with 3 shards.
- Each shard will be a replica set with 3 members (including one arbiter).
- We will insert some data and see where it goes.

### Sharded Cluster Configuration

- Three shards:
  1. A replica set on ports 27107, 27108, 27109
  2. A replica set on ports 27117, 27118, 27119
  3. A replica set on ports 27127, 27128, 27129
- Three config servers on ports 27217, 27218, 27219
- Two mongos servers at ports 27017 and 27018

## Build Our Data Directories

On Linux or MacOS, run the following in the terminal to create the data directories we'll need.

```
mkdir -p ~/data/cluster/config/{c0,c1,c2}
mkdir -p ~/data/cluster/shard0/{m0,m1,arb}
mkdir -p ~/data/cluster/shard1/{m0,m1,arb}
mkdir -p ~/data/cluster/shard2/{m0,m1,arb}
mkdir -p ~/data/cluster/{s0,s1}
```

On Windows, run the following commands instead:

```
md c:\data\cluster\config\c0 c:\data\cluster\config\c1 c:\data\cluster\config\c2
md c:\data\cluster\shard0\m0 c:\data\cluster\shard0\m1 c:\data\cluster\shard0\arb
md c:\data\cluster\shard1\m0 c:\data\cluster\shard1\m1 c:\data\cluster\shard1\arb
md c:\data\cluster\shard2\m0 c:\data\cluster\shard2\m1 c:\data\cluster\shard2\arb
md c:\data\cluster\s0 c:\data\cluster\s1
```

## Initiate a Replica Set (Linux/MacOS)

```
mongod --replSet shard0 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard0/m0 \
  --logpath ~/data/cluster/shard0/m0/mongod.log \
  --fork --port 27107

mongod --replSet shard0 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard0/m1 \
  --logpath ~/data/cluster/shard0/m1/mongod.log \
  --fork --port 27108

mongod --replSet shard0 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard0/arb \
  --logpath ~/data/cluster/shard0/arb/mongod.log \
  --fork --port 27109

mongo --port 27107 --eval "\
  rs.initiate(); sleep(3000);\
  rs.add('$HOSTNAME:27108');\
  rs.addArb('$HOSTNAME:27109')"
```

## Initiate a Replica Set (Windows)

```
mongod --replSet shard0 --smallfiles --nojournal --noprealloc \
  --dbpath c:\data\cluster\shard0\m0 \
  --logpath c:\data\cluster\shard0\m0\mongod.log \
  --port 27107 --oplogSize 10

mongod --replSet shard0 --smallfiles --nojournal --noprealloc \
  --dbpath c:\data\cluster\shard0\m1 \
  --logpath c:\data\cluster\shard0\m1\mongod.log \
  --port 27108 --oplogSize 10

mongod --replSet shard0 --smallfiles --nojournal --noprealloc \
  --dbpath c:\data\cluster\shard0\arb \
  --logpath c:\data\cluster\shard0\arb\mongod.log \
  --port 27109 --oplogSize 10
```

```
mongo --port 27107 --eval "\
rs.initiate(); sleep(3000);\
rs.add ('<HOSTNAME>:27108');\
rs.addArb('<HOSTNAME>:27109')"
```

## Spin Up a Second Replica Set (Linux/macOS)

```
mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
--dbpath ~/data/cluster/shard1/m0 \
--logpath ~/data/cluster/shard1/m0/mongod.log \
--fork --port 27117
```

```
mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
--dbpath ~/data/cluster/shard1/m1 \
--logpath ~/data/cluster/shard1/m1/mongod.log \
--fork --port 27118
```

```
mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
--dbpath ~/data/cluster/shard1/arb \
--logpath ~/data/cluster/shard1/arb/mongod.log \
--fork --port 27119
```

```
mongo --port 27117 --eval "\
rs.initiate(); sleep(3000);\
rs.add ('$HOSTNAME:27118');\
rs.addArb('$HOSTNAME:27119')"
```

## Spin Up a Second Replica Set (Windows)

```
mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
--dbpath c:\data\cluster\shard1\m0 \
--logpath c:\data\cluster\shard1\m0\mongod.log \
--port 27117 --oplogSize 10
```

```
mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
--dbpath c:\data\cluster\shard1\m1 \
--logpath c:\data\cluster\shard1\m1\mongod.log \
--port 27118 --oplogSize 10
```

```
mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
--dbpath c:\data\cluster\shard1\arb \
--logpath c:\data\cluster\shard1\arb\mongod.log \
--port 27119 --oplogSize 10
```

```
mongo --port 27117 --eval "\
rs.initiate(); sleep(3000);\
rs.add ('<HOSTNAME>:27118');\
rs.addArb('<HOSTNAME>:27119')"
```

## A Third Replica Set (Linux/macOS)

```
mongod --replSet shard2 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard2/m0 \
  --logpath ~/data/cluster/shard2/m0/mongod.log \
  --fork --port 27127

mongod --replSet shard2 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard2/m1 \
  --logpath ~/data/cluster/shard2/m1/mongod.log \
  --fork --port 27128

mongod --replSet shard2 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard2/arb \
  --logpath ~/data/cluster/shard2/arb/mongod.log \
  --fork --port 27129

mongo --port 27127 --eval "\
  rs.initiate(); sleep(3000);\
  rs.add    ('$HOSTNAME:27128');\
  rs.addArb('$HOSTNAME:27129')"
```

## A Third Replica Set (Windows)

```
mongod --replSet shard2 --smallfiles --nojournal --noprealloc \
  --dbpath c:\data\cluster\shard2\m0 \
  --logpath c:\data\cluster\shard2\m0\mongod.log \
  --port 27127 --oplogSize 10

mongod --replSet shard2 --smallfiles --nojournal --noprealloc \
  --dbpath c:\data\cluster\shard2\m1 \
  --logpath c:\data\cluster\shard2\m1\mongod.log \
  --port 27128 --oplogSize 10

mongod --replSet shard2 --smallfiles --nojournal --noprealloc \
  --dbpath c:\data\cluster\shard2\arb \
  --logpath c:\data\cluster\shard2\arb\mongod.log \
  --port 27129 --oplogSize 10

mongo --port 27127 --eval "\
  rs.initiate(); sleep(3000);\
  rs.add    ('<HOSTNAME>:27128');\
  rs.addArb('<HOSTNAME>:27129')"
```



## Status Check

- Now we have three replica sets running.
- We have one for each shard.
- They do not know about each other yet.
- To make them a sharded cluster we will:
  - Build our config databases
  - Launch our mongos processes
  - Add each shard to the cluster
- To benefit from this configuration we also need to:
  - Enable sharding for a database
  - Shard at least one collection within that database

## Launch Config Servers (Linux/MacOS)

```
mongod --smallfiles --nojournal --noprealloc \  
  --dbpath ~/data/cluster/config/c0 \  
  --logpath ~/data/cluster/config/c0/mongod.log \  
  --fork --port 27217 --configsvr
```

```
mongod --smallfiles --nojournal --noprealloc \  
  --dbpath ~/data/cluster/config/c1 \  
  --logpath ~/data/cluster/config/c1/mongod.log \  
  --fork --port 27218 --configsvr
```

```
mongod --smallfiles --nojournal --noprealloc \  
  --dbpath ~/data/cluster/config/c2 \  
  --logpath ~/data/cluster/config/c2/mongod.log \  
  --fork --port 27219 --configsvr
```

## Launch Config Servers (Windows)

```
mongod --smallfiles --nojournal --noprealloc \  
  --dbpath c:\data\cluster\config\c0 \  
  --logpath c:\data\cluster\config\c0\mongod.log \  
  --port 27217 --configsvr
```

```
mongod --smallfiles --nojournal --noprealloc \  
  --dbpath c:\data\cluster\config\c1 \  
  --logpath c:\data\cluster\config\c1\mongod.log \  
  --port 27218 --configsvr
```

```
mongod --smallfiles --nojournal --noprealloc \  
  --dbpath c:\data\cluster\config\c2 \  
  --logpath c:\data\cluster\config\c2\mongod.log \  
  --port 27219 --configsvr
```



## Observe What Happens

Connect to either mongos using a mongo shell and frequently issue:

```
sh.status()
```

## 6 Aggregation

*Aggregation Tutorial (page 66)* An introduction to the the aggregation framework, pipeline concept, and stages

*Optimizing Aggregation (page 74)* Resource management in the aggregation pipeline

*Lab: Aggregation Framework (page 76)* Aggregation labs

### 6.1 Aggregation Tutorial

#### Learning Objectives

Upon completing this module students should understand:

- The concept of the aggregation pipeline
- The stages of the aggregation pipeline
- How to use aggregation operators
- The fundamentals of using aggregation for data analysis
- Group aggregation operators
- Using the same operator in multiple stages of an aggregation pipeline

#### Aggregation Basics

- Use the aggregation framework to transform and analyze data in MongoDB collections.
- For those who are used to SQL, aggregation can be similar to `GROUP BY`.
- The aggregation framework is based on the concept of a pipeline.

#### The Aggregation Pipeline

- An aggregation pipeline is analogous to a UNIX pipeline.
- Each stage of the pipeline:
  - Receives a set of documents as input.
  - Performs an operation on those documents.
  - Produces a set of documents for use by the following stage.
- A pipeline has the following syntax:

```
db.<COLLECTION>.aggregate( [ { stage1 }, { stage2 }, ... ],  
                           { options } )
```

## Aggregation Stages

- `$match`: Similar to `find()`
- `$project`: Shape documents
- `$sort`: Like the cursor method of the same name
- `$skip`: Like the cursor method of the same name
- `$limit`: Like the cursor method of the same name
- `$unwind`: Used for working with arrays
- `$group`: Used to aggregate field values from multiple documents
- `$out`: Creates a new collection from the output of an aggregation pipeline

### The Match Stage

- The `$match` operator works like the query phase of `find()`
- Documents in the pipeline that match the query document will be passed to subsequent stages.
- `$match` is often the first operator used in an aggregation stage.
- Like other aggregation operators, `$match` can occur multiple times in a single pipeline.

### The Project Stage

- `$project` allows you to shape the documents into what you need for the next stage.
  - The simplest form of shaping is using `$project` to select only the fields you are interested in.
  - `$project` can also create new fields from other fields in the input document.
    - \* *E.g.*, you can pull a value out of an embedded document and put it at the top level.
    - \* *E.g.*, you can create a ratio from the values of two fields as pass along as a single field.
- `$project` produces 1 output document for every input document it sees.

## A Twitter Dataset

- Let's look at some examples that illustrate the MongoDB aggregation framework.
- These examples operate on a collection of tweets.
  - As with any dataset of this type, it's a snapshot in time.
  - It may not reflect the structure of Twitter feeds as they look today.

## Tweets Data Model

```
{
  "text" : "Something interesting ...",
  "entities" : {
    "user_mentions" : [
      {
        "screen_name" : "somebody_else",
        ...
      }
    ],
    "urls" : [ ],
    "hashtags" : [ ]
  },
  "user" : {
    "friends_count" : 544,
    "screen_name" : "somebody",
    "followers_count" : 100,
    ...
  },
}
```

## Analyzing Tweets

- Imagine the types of analyses one might want to do on tweets.
- It's common to analyze the behavior of users and the networks involved.
- Our examples will focus on this type of analysis

## Friends and Followers

- Let's look again at two stages we touched on earlier:
  - `$match`
  - `$project`
- In our dataset:
  - `friends` are those a user follows.
  - `followers` are others that follow a users.
- Using these operators we will write an aggregation pipeline that will:
  - Ignore anyone with no friends and no followers.
  - Calculate who has the highest followers to friends ratio.

## Exercise: Friends and Followers

```
db.tweets.aggregate( [
  { $match: { "user.friends_count": { $gt: 0 },
             "user.followers_count": { $gt: 0 } } },
  { $project: { ratio: { $divide: ["$user.followers_count",
                                   "$user.friends_count"] },
              screen_name : "$user.screen_name" } },
  { $sort: { ratio: -1 } },
  { $limit: 1 } ] )
```

## Exercise: \$match and \$project

- Of the users in the “Brasilia” timezone who have tweeted 100 times or more, who has the largest number of followers?
- Time zone is found in the “time\_zone” field of the user object in each tweet.
- The number of tweets for each user is found in the “statuses\_count” field.
- A result document should look something like the following:

```
{ _id      : ObjectId('52fd2490bac3fa1975477702'),
  followers : 2597,
  screen_name: 'marbles',
  tweets    : 12334
}
```

## The Group Stage

- For those coming from the relational world, \$group is similar to the SQL GROUP BY statement.
- \$group operations require that we specify which field to group on.
- Documents with the same identifier will be aggregated together.
- With \$group, we aggregate values using [accumulators](http://docs.mongodb.org/manual/meta/aggregation-quick-reference/#accumulators)<sup>6</sup>.

## Tweet Source

- The tweets in our twitter collection have a field called `source`.
- This field describes the application that was used to create the tweet.
- Let’s write an aggregation pipeline that identifies the applications most frequently used to publish tweets.

---

<sup>6</sup><http://docs.mongodb.org/manual/meta/aggregation-quick-reference/#accumulators>

## Exercise: Tweet Source

```
db.tweets.aggregate( [
  { "$group" : { "_id" : "$source",
                 "count" : { "$sum" : 1 } } },
  { "$sort" : { "count" : -1 } }
] )
```

## Group Aggregation Accumulators

Accumulators available in the group stage:

- \$sum
- \$avg
- \$first
- \$last
- \$max
- \$min
- \$push
- \$addToSet

## Rank Users by Number of Tweets

- One common task is to rank users based on some metric.
- Let's look at who tweets the most.
- Earlier we did the same thing for tweet source.
  - Group together all tweets by a user for every user in our collection
  - Count the tweets for each user
  - Sort in decreasing order
- Let's add the list of tweets to the output documents.
- Need to use an accumulator that works with arrays.
- Can use either \$addToSet or \$push.



## Exercise: Adding List of Tweets

For each user, aggregate all their tweets into a single array.

```
db.tweets.aggregate( [
  { "$group" : { "_id" : "$user.screen_name",
                 "tweet_texts" : { "$push" : "$text" },
                 "count" : { "$sum" : 1 } } },
  { "$sort" : { "count" : -1 } },
  { "$limit" : 3 }
] )
```

## The Unwind Stage

- In many situations we want to aggregate using values in an array field.
- In our tweets dataset we need to do this to answer the question:
  - “Who includes the most user mentions in their tweets?”
- User mentions are stored within an embedded document for entities.
- This embedded document also lists any urls and hashtags used in the tweet.

## Example: User Mentions in a Tweet

```
...
"entities" : {
  "user_mentions" : [
    {
      "indices" : [
        28,
        44
      ],
      "screen_name" : "LatinsUnitedGSX",
      "name" : "Henry Ramirez",
      "id" : 102220662
    }
  ],
  "urls" : [ ],
  "hashtags" : [ ]
},
...
```

## Using \$unwind

Who includes the most user mentions in their tweets?

```
db.tweets.aggregate(
  { $unwind: "$entities.user_mentions" },
  { $group: { _id: "$user.screen_name",
              count: { $sum: 1 } } },
  { $sort: { count: -1 } },
  { $limit: 1 })
```

## Data Processing Pipelines

- The aggregation framework allows you to create a data processing pipeline.
- You can include as many stages as necessary to achieve your goal.
- For each stage consider:
  - What input that stage must receive
  - What output it should produce.
- Many tasks require us to include more than one stage using a given operator.

## Most Unique User Mentions

- We frequently need multiple group stages to achieve our goal.
- We just looked at a pipeline to find the tweeter that mentioned the most users.
- Let's change this so that it is more of a question about a tweeter's active network.
- We might ask which tweeter has mentioned the most unique users in their tweets.

## Same Operator (\$group), Multiple Stages

Which tweeter has mentioned the most unique users in their tweets?

```
db.tweets.aggregate( [
  { $unwind: "$entities.user_mentions" },
  { $group: {
    _id: "$user.screen_name",
    mset: { $addToSet: "$entities.user_mentions.screen_name" } } },
  { $unwind: "$mset" },
  { $group: { _id: "$_id", count: { $sum: 1 } } },
  { $sort: { count: -1 } },
  { $limit: 1 }
] )
```

## The Sort Stage

- Uses the `$sort` operator
- Works like the `sort()` cursor method
- 1 to sort ascending; -1 to sort descending
- E.g, `db.testcol.aggregate( [ { $sort : { b : 1, a : -1 } } ] )`

## The Skip Stage

- Uses the `$skip` operator
- Works like the `skip()` cursor method.
- Value is an integer specifying the number of documents to skip.
- E.g, the following will pass all but the first 3 documents to the next stage in the pipeline.
  - `db.testcol.aggregate( [ { $skip : 3 }, ... ] )`

## The Limit Stage

- Used to limit the number of documents passed to the next aggregation stage.
- Works like the `limit()` cursor method.
- Value is an integer.
- E.g., the following will only pass 3 documents to the stage that comes next in the pipeline.
  - `db.testcol.aggregate( [ { $limit: 3 }, ... ] )`

## The Out Stage

- Used to create a new collection from the output of the aggregation pipeline.
- Can only be the last stage in the pipeline.
- If a collection by the name already exists, it replaces that collection.
- Syntax is `{ $out : "collection_name" }`

## 6.2 Optimizing Aggregation

### Learning Objectives

Upon completing this module students should understand:

- Aggregation pipeline options
- Key aspects of resource management during the aggregation pipeline
- How to order aggregation stages to maximize speed and minimize resource usage
- How MongoDB automatically reorders pipeline stages to improve efficiency
- Changes in the aggregation framework from MongoDB 2.4 to 2.6.

### Aggregation Options

- You may pass an options document to `aggregate()`.
- Syntax:

```
db.<COLLECTION>.aggregate( [ { stage1 }, { stage2 }, ... ], { options } )
```

- Following are some of the fields that may be passed in the options document.
  - `allowDiskUse : true` - permit the use of disk for memory-intensive queries
  - `explain : true` - display how indexes are used to perform the aggregation.

### Aggregation Limits

- An aggregation pipeline cannot use more than 100 MB of RAM.
- `allowDiskUse : true` allows you to get around this limit.
- The follow operators do not require the entire dataset to be in memory:
  - `$match`, `$skip`, `$limit`, `$unwind`, and `$project`
  - Stages for these operators are not subject to the 100 MB limit.
  - `$unwind` can, however, dramatically increase the amount of memory used.
- `$group` and `$sort` might require all documents in memory at once.

## Limits Prior to MongoDB 2.6

- `aggregate()` returned results in a single document up to 16 MB in size.
- The upper limit on pipeline memory usage was 10% of RAM.

## Optimization: Reducing Documents in the Pipeline

- These operators can reduce the number of documents in the pipeline:
  - `$match`
  - `$skip`
  - `$limit`
- They should be used as early as possible in the pipeline.

## Optimization: Sorting

- `$sort` can take advantages of indexes.
- Must be used before any of the following to do this:
  - `$group`
  - `$unwind`
  - `$project`
- After these stages, the fields or their values change.
- `$sort` requires a full scan of the input documents.

## Automatic Optimizations

MongoDB will perform some optimizations automatically. For example:

- If a `$project` stage is used late in the pipeline it may be used to eliminate those fields earlier if possible.
- A `$sort` followed by a `$match` will be executed as a `$match` followed by a `$sort` to reduce the number of documents to be sorted.
- A `$skip` followed by a `$limit` will be executed as a `$limit` followed by a `$skip`, with the `$limit` parameter increased by the `$skip` amount to allow `$sort + $limit` coalescence.
- See: [Aggregation Pipeline Optimization](http://docs.mongodb.org/manual/core/aggregation-pipeline-optimization/)<sup>7</sup>

## 6.3 Lab: Aggregation Framework

### Exercise: Working with Array Fields

Use the aggregation framework to find the name of the individual who has made the most comments on a blog.

Start by importing the necessary data if you have not already.

```
# for version <= 2.6.x
mongoimport -d blog -c posts --drop posts.json
# for version > 3.0
mongoimport -d blog -c posts --drop --batchSize=100 posts.json
```

To help you verify your work, the author with the fewest comments is Mariela Sherer and she commented 387 times.

### Exercise: Repeated Aggregation Stages

Import the `zips.json` file from the data handouts provided:

```
mongoimport -d sample -c zips --drop zips.json
```

Consider together cities in the states of California (CA) and New York (NY) with populations over 25,000. Calculate the average population of this sample of cities.

Please note:

- Different states might have the same city name.
- A city might have multiple zip codes.

---

<sup>7</sup><http://docs.mongodb.org/manual/core/aggregation-pipeline-optimization/>

## Exercise: Projection

Calculate the total number of people who live in a zip code in the US where the city starts with a digit.

`$project` can extract the first digit from any field. E.g.,

```
db.zips.aggregate([
  {$project:
    {
      first_char: { $substr: ["$city", 0, 1] },
    }
  })
```

## Exercise: Descriptive Statistics

From the `grades` collection, find the class (display the `class_id`) with the highest average student performance on **exams**. To solve this problem you'll want an average of averages.

First calculate the average exam score of each student in each class. Then determine the average class exam score using these values. If you have not already done so, import the `grades` collection as follows.

```
mongoimport -d sample -c grades --drop grades.json
```

Before you attempt this exercise, explore the `grades` collection a little to ensure you understand how it is structured.

For additional exercises, consider other statistics you might want to see with this data and how to calculate them.

## 7 Application Engineering

*Introduction (page 78)* MongoMart Introduction

*Java Driver Labs (MongoMart) (page 79)* Build an e-commerce site backed by MongoDB (Java)

### 7.1 Introduction

#### What is MongoMart

MongoMart is an on-line store for buying MongoDB merchandise. We'll use this application to learn more about interacting with MongoDB through the driver.

#### MongoMart Demo of Fully Implemented Version

- View Items
- View Items by Category
- Text Search
- View Item Details
- Shopping Cart

#### View Items

- <http://localhost:8080>
- Pagination and page numbers
- Click on a category

#### View Items by Category

- <http://localhost:8080/?category=Apparel>
- Pagination and page numbers
- “All” is listed as a category, to return to all items listing



## Text Search

- <http://localhost:8080/search?query=shirt>
- Search for any word or phrase in item title, description or slogan
- Pagination

## View Item Details

- <http://localhost:8080/item?id=1>
- Star rating based on reviews
- Add a review
- Related items
- Add item to cart

## Shopping Cart

- <http://localhost:8080/cart>
- Adding an item multiple times increments quantity by 1
- Change quantity of any item
- Changing quantity to 0 removes item

## 7.2 Java Driver Labs (MongoMart)

### Introduction

- In this lab, we'll set up and optimize an application called MongoMart. MongoMart is an on-line store for buying MongoDB merchandise.

### Lab: Setup and Connect to the Database

- Import the “item” collection to a standalone MongoDB server (without replication) as noted in the README.md file of the /data directory of MongoMart
- Become familiar with the structure of the Java application in /java/src/main/java/mongomart/
- Modify the MongoMart.java class to properly connect to your local database instance

### **Lab: Populate All Necessary Database Queries**

- After running the MongoMart.java class, navigate to “localhost:8080” to view the application
- Initially, all data is static and the application does not query the database
- Modify the ItemDao.java and CartDao.java classes to ensure all information comes from the database (do not modify the method return types or parameters)

### **Lab: Use a Local Replica Set with a Write Concern**

- It is important to use replication for production MongoDB instances, however, Lab 1 advised us to use a standalone server.
- Convert your local standalone mongod instance to a three node replica set named “shard1”
- Modify MongoMart’s MongoDB connection string to include at least two nodes from the replica set
- Modify your application’s write concern to MAJORITY for all writes to the “cart” collection, any writes to the “item” collection should continue using the default write concern of W:1

## 8 Reporting Tools and Diagnostics

*Performance Troubleshooting (page 81)* An introduction to reporting and diagnostic tools for MongoDB

### 8.1 Performance Troubleshooting

#### Learning Objectives

Upon completing this module students should understand basic performance troubleshooting techniques and tools including:

- `mongostat`
- `mongotop`
- `db.setProfilingLevel()`
- `db.currentOp()`
- `db.<COLLECTION>.stats()`
- `db.serverStatus()`

#### **mongostat and mongotop**

- `mongostat` samples a server every second.
  - See current ops, pagefaults, network traffic, etc.
  - Does not give a view into historic performance; use MMS for that.
- `mongotop` looks at the time spent on reads/writes in each collection.

#### **Exercise: mongostat (setup)**

In one window, perform the following commands.

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  arr = [];
  for (j=1; j<=1000; j++) {
    doc = { _id: (1000 * (i-1) + j), a: i, b: j, c: (1000 * (i-1) + j) };
    arr.push(doc)
  };
  db.testcol.insertMany(arr);
  var x = db.testcol.find( { b : 255 } );
  x.next();
  var x = db.testcol.find( { _id : 1000 * (i-1) + 255 } );
  x.next();
  var x = "asdf";
  db.testcol.updateOne( { a : i, b : 255 }, { $set : { d : x.pad(1000) } });
  print(i)
}
```

### Exercise: mongostat (run)

- In another window/tab, run mongostat.
- You will see:
  - Inserts
  - Queries
  - Updates

### Exercise: mongostat (create index)

- In a third window, create an index when you see things slowing down:

```
db.testcol.createIndex( { a : 1, b : 1 } )
```
- Look at mongostat.
- Notice that things are going significantly faster.
- Then, let's drop that and build another index.

```
db.testcol.dropIndexes()  
db.testcol.createIndex( { b : 1, a : 1 } )
```

### Exercise: mongotop

Perform the following then, in another window, run mongotop.

```
db.testcol.drop()  
for (i=1; i<=10000; i++) {  
  arr = [];  
  for (j=1; j<=1000; j++) {  
    doc = { _id: (1000*(i-1)+j), a: i, b: j, c: (1000*(i-1)+j) };  
    arr.push(doc)  
  };  
  db.testcol.insertMany(arr);  
  var x = db.testcol.find( {b: 255} ); x.next();  
  var x = db.testcol.find( { _id: 1000*(i-1)+255 } ); x.next();  
  var x = "asdf";  
  db.testcol.updateOne( {a: i, b: 255}, { $set: {d: x.pad(1000)}} );  
  print(i)  
}
```

## **db.currentOp()**

- currentOp is a tool that asks what the db is doing at the moment.
- currentOp is useful for finding long-running processes.
- Fields of interest:
  - microsecs\_running
  - op
  - query
  - lock
  - waitingForLock

## **Exercise: db.currentOp()**

Do the following then, connect with a separate shell, and repeatedly run `db.currentOp()`.

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  arr = [];
  for (j=1; j<=1000; j++) {
    doc = {_id: (1000*(i-1)+j), a: i, b: j, c: (1000*(i-1)+j)};
    arr.push(doc)
  };
  db.testcol.insertMany(arr);
  var x = db.testcol.find( {b: 255} ); x.next();
  var x = db.testcol.find( {_id: 1000*(i-1)+255 } ); x.next();
  var x = "asdf";
  db.testcol.updateOne( {a: i, b: 255}, {$set: {d: x.pad(1000)}});
  print(i)
}
```

## **db.<COLLECTION>.stats()**

- Used to view the current stats for a collection.
- Everything is in bytes; use the multiplier parameter to view in KB, MB, etc
- You can also use `db.stats()` to do this at scope of the entire database

## Exercise: Using Collection Stats

Look at the output of the following:

```
db.testcol.drop()
db.testcol.insertOne( { a : 1 } )
db.testcol.stats()
var x = "asdf"
db.testcol2.insertOne( { a : x.pad(10000000) } )
db.testcol2.stats()
db.stats()
```

## The Profiler

- Off by default.
- To reset, `db.setProfilerLevel(0)`
- At setting 1, it captures “slow” queries.
- You may define what “slow” is.
- Default is 100ms: `db.setProfilerLevel(1)`
- E.g., to capture 20 ms: `db.setProfilerLevel(1, 20)`

## The Profiler (continued)

- If the profiler level is 2, it captures all queries.
  - This will severely impact performance.
  - Turns all reads into writes.
- Always turn the profiler off when done (set level to 0)
- Creates `db.system.profile` collection

## Exercise: Exploring the Profiler

Perform the following, then look in your `db.system.profile`.

```
db.setProfilingLevel(0)
db.testcol.drop()
db.system.profile.drop()
db.setProfilingLevel(2)
db.testcol.insertOne( { a : 1 } )
db.testcol.find()
var x = "asdf"
db.testcol.insertOne( { a : x.pad(10000000) } ) // ~10 MB
db.setProfilingLevel(0)
db.system.profile.find().pretty()
```

## **db.serverStatus()**

- Takes a snapshot of server status.
- By taking diffs, you can see system trends.
- Most of the data that MMS gets is from here.

## **Exercise: Using db.serverStatus()**

- Open up two windows. In the first, type:

```
db.testcol.drop()  
var x = "asdf"  
for (i=0; i<=10000000; i++) {  
    db.testcol.insertOne( { a : x.pad(100000) } )  
}
```

- In the second window, type periodically:

```
var x = db.serverStatus(); x.metrics.document
```

## **Analyzing Profiler Data**

- Enable the profiler at default settings.
- Run for 5 seconds.
- Slow operations are captured.
- The issue is there is not a proper index on the message field.
- You will see how fast documents are getting inserted.
- It will be slow b/c the documents are big.

## **Performance Improvement Techniques**

- Appropriate write concerns
- Bulk operations
- Good schema design
- Good Shard Key choice
- Good indexes

## Performance Tips: Write Concern

- Increasing the write concern increases data safety.
- This will have an impact on performance, however.
- This is especially true when there are network issues.

## Bulk Operations

- Using bulk operations (including `insertMany` and `updateMany` ) can improve performance, especially when using write concern greater than 1.
- These enable the server to amortize acknowledgement.
- Can be done with both `insertMany` and `updateMany` .

## Exercise: Comparing `insertMany` with `mongostat`

Let's spin up a 3-member replica set:

```
mkdir -p /data/replset/{1,2,3}
mongod --logpath /data/replset/1/mongod.log \
      --dbpath /data/replset/1 --replSet mySet --port 27017 --fork
mongod --logpath /data/replset/2/mongod.log \
      --dbpath /data/replset/2 --replSet mySet --port 27018 --fork
mongod --logpath /data/replset/3/mongod.log \
      --dbpath /data/replset/3 --replSet mySet --port 27019 --fork

echo "conf = {_id: 'mySet', members: [{_id: 0, host: 'localhost:27017'}, \
    {_id: 1, host: 'localhost:27018'}, {_id: 2, host: 'localhost:27019'}]}; \
rs.initiate(conf)" | mongo
```

## `mongostat`, `insertOne` with `{w: 1}`

Perform the following, with `writeConcern : 1` and `insertOne()`:

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  for (j=1; j<=1000; j++) {
    db.testcol.insertOne( { _id : (1000 * (i-1) + j),
                          a : i, b : j, c : (1000 * (i-1) + j) },
                          { writeConcern : { w : 1 } } );
  };
  print(i);
}
```

Run `mongostat` and see how fast that happens.



## Multiple insertOne s with {w: 3}

Increase the write concern to 3 (safer but slower):

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  for (j=1; j<=1000; j++) {
    db.testcol.insertOne(
      { _id: (1000 * (i-1) + j), a: i, b: j, c: (1000 * (i-1)+ j)},
      { writeConcern: { w: 3 } }
    );
  };
  print(i);
}
```

Again, run mongostat.

## mongostat, insertMany with {w: 3}

- Finally, let's use insertMany to our advantage:
- Note that writeConcern is still { w: 3 }

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  arr = []
  for (j=1; j<=1000; j++) {
    arr.push(
      { _id: (1000 * (i-1) + j), a: i, b: j, c: (1000 * (i-1)+ j) }
    );
  };
  db.testcol.insertMany( arr, { writeConcern : { w : 3 } } );
  print(i);
}
```

## Schema Design

- The structure of documents affects performance.
- Optimize for your application's read/write patterns.
- We want as few requests to the database as possible to perform a given application task.
- See the data modeling section for more information.

## Shard Key Considerations

- Choose a shard key that distributes load across your cluster.
- Create a shard key such that only a small number of documents will have the same value.
- Create a shard key that has a high degree of randomness.
- Your shard key should enable a mongos to target a single shard for a given query.

## Indexes and Performance

- Reads and writes that don't use an index will cripple performance.
- In compound indexes, order matters:
  - Sort on a field that comes before any range used in the index.
  - You can't skip fields; they must be used in order.
  - Revisit the indexing section for more detail.





**Find out more**

[mongodb.com](http://mongodb.com) | [mongodb.org](http://mongodb.org)  
[university.mongodb.com](http://university.mongodb.com)

**Having trouble?**

File a JIRA ticket:  
[jira.mongodb.org](http://jira.mongodb.org)

**Follow us on twitter**

[@MongoDBInc](https://twitter.com/MongoDBInc)  
[@MongoDB](https://twitter.com/MongoDB)