

# Drivers

[http://docs.mongodb.org/  
ecosystem/drivers/](http://docs.mongodb.org/ecosystem/drivers/)

# Outline

- Role of the driver and implementation
- Some internals
- Shell tips & tricks!

# Role of the driver

One reason why MongoDB was successfully adopted by developers was the investment in almost a dozen drivers up front.

This allowed developers to work with objects that were already familiar to them:

- Python -> dictionaries
- Ruby -> hashes
- Perl -> associative arrays

Drivers provide the ability to persist known objects in the DB without having to learn too many new things.

# Drivers and MongoDB 3.0

- 2 releases around the same time in Q1-2015
  - A. Needed to support the new SCRAM authentication scheme for MongoDB 3.0
    - Users must upgrade drivers to use 3.0 if *-auth* is enabled
  - B. Drivers are all supporting the same specifications
    - <https://github.com/mongodb/specifications/tree/master/source>
    - Before, they implemented functionality differently

# Common Specifications

- CRUD
  - <https://github.com/mongodb/specifications/blob/master/source/crud/crud.rst>
- Authentication
  - <https://github.com/mongodb/specifications/blob/master/source/auth/auth.rst>
- Server Discovery and Monitoring
  - <https://github.com/mongodb/specifications/blob/master/source/server-discovery-and-monitoring/server-discovery-and-monitoring-summary.rst>
- Server Selection
  - <https://github.com/mongodb/specifications/blob/master/source/server-selection/server-selection.rst>
- Index Management
  - <https://github.com/mongodb/specifications/blob/master/source/index-management.rst>

# What does a driver need to do

- BSON support
- Wire Protocol
- Keeping track of the cluster status
- Implements the following commands:
  - Find
  - getMore
  - killCursors

# BSON

- BSON is the language of the DB; it's what's sent over the wire, stored in the DB...
- [bsonspec.org](http://bsonspec.org)

# API for CRUD+ operations

- Essentials
  - Basic operations: query, insert, update, remove, ensureIndex, findOne, limit, sort
  - Fetch more data from a cursor (dbGetMore)
  - Sending the KillCursors operation
- Nice to haves
  - Database \$cmd support and helpers
  - getLastError(), count(), eval(), explain(), hint(), \$where



# Wire Protocol

- Clients communicate with the database server through a regular TCP/IP socket
- There are two types of messages, client requests and database responses, each having a slightly different structure.
- Commands from the driver:
  - OP\_QUERY
  - OP\_GET\_MORE
  - OP\_KILL\_CURSORS
  - OP\_INSERT
  - OP\_UPDATE
  - OP\_DELETE
  - OP\_COMMAND
  - OP\_COMMANDREPLY
- Wire Protocol Versions:
  - 2: write commands
  - 3: auth protocol
  - 4: read concern
- Spec: <https://docs.mongodb.org/manual/reference/mongodb-wire-protocol/>

# Standard Message Header

```
struct MsgHeader {  
    int32    messageLength; // total message size, including this  
    int32    requestId;      // identifier for this message  
    int32    responseTo;     // requestId from the original request  
                                // (used in reponses from db)  
    int32    opCode;         // request type - see table below  
}
```

Opcode Name	opCode value	Comment
OP_REPLY	1	Reply to a client request. responseTo is set
OP_MSG	1000	generic msg command followed by a string
OP_UPDATE	2001	update document
OP_INSERT	2002	insert new document
RESERVED	2003	formerly used for OP_GET_BY_OID
OP_QUERY	2004	query a collection
OP_GET_MORE	2005	Get more data from a query. See Cursors
OP_DELETE	2006	Delete documents
OP_KILL_CURSORS	2007	Tell database client is done with a cursor

# getLastError

- App sends directives and then getLastError (GLE) to find out what happened
- Resulted in strange behaviors since you had to keep connections open to send GLE without sending other operations first
- Drivers abstracted this away, but still made two calls under the hood: one for the op and one for GLE
- Since 2.6, most write methods are returning a status at the end of command.
  - Op\_Command => 2010
  - Op\_CommandReply => 2011

# How about all the other commands?

```
> db.runCommand
function ( obj, extra ){
  if ( typeof( obj ) == "string" ){
    var n = {};
    n[obj] = 1;
    obj = n;
    if ( extra && typeof( extra ) == "object" ) {
      for ( var x in extra ) {
        n[x] = extra[x];
      }
    }
  }
  return this.getCollection( "$cmd" ).findOne( obj );
}
```

A command is just a findOne on some special \$cmd collection

# Look at example messages

- <http://docs.mongodb.org/meta-driver/latest/legacy/mongodb-wire-protocol/#client-request-messages>
- The types used in these documents (cstring, int32, etc.) are the same as those defined in the BSON specification.
- Integer constants are in capitals (e.g. ZERO for the integer value of 0).

# Cursors

- Cursor id returned in query response, driver calls *getmore*
- Tailable cursors, e.g. *OpLog*
- Cursors timeout after 10 minutes by default
  - In Sharded Clusters, chunk migration cleaning is not completed until cursors are closed
- But cursors are dumb, so you must be smart!

# Connection handling/pools

- Pool of connections that can be shared by many many threads
  - Replica Set vs Sharded Cluster
- Automatically connect to proper server, and failover, when connecting to a replica set
- Driver does *\*NOT\** retry a failed operation
  - “Shit will happen” – Hannes Magnuson



# MongoDB supported drivers

Documentation	Releases	Source	API	JIRA	Online Course
C	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	
C++11	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	
C++ (legacy)	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	
C#	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	<a href="#">Course</a>
Java	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	<a href="#">Course</a>
Node.js	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	<a href="#">Course</a>
Perl	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	
PHP	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	
Python	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	<a href="#">Course</a>
Motor	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	
Ruby	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	
Scala	<a href="#">Releases</a>	<a href="#">Source</a>	<a href="#">API</a>	<a href="#">JIRA</a>	

# Object Mappers

- Names
  - ORM – Object Relational Mapping tool
  - ODM – Object Data-store Mapping tool
  - OM – Object Mapping tool
  - DRM – Document Resource Mapping tool
- Some examples per language
  - Java: Morphia\*, Spring MongoDB
  - Node.js: Mongoose

# Driver/Server compatibility

- [https://docs.google.com/a/10gen.com/spreadsheet/ccc?key=0AkOLPg1iZn3qdFc0T25hd0Y3VGZyTkR2aE9JQmw2UXc&usp=drive\\_web#gid=1](https://docs.google.com/a/10gen.com/spreadsheet/ccc?key=0AkOLPg1iZn3qdFc0T25hd0Y3VGZyTkR2aE9JQmw2UXc&usp=drive_web#gid=1)

		Java Driver							
		Server Version							
		1.8	2.0	2.2	2.4	2.6	3.0		
Driver Version	2.9								
	2.10								
	2.11								
	2.12								
	2.13								
	3.0 (planned)								
		Language Version							
		Java 5	Java 6	Java 7	Java 8				
Driver Version	2.7								
	2.8								
	2.9								
	2.10								
	2.11								
	2.12								
	2.13 (planned)								
3.0 (planned)									

# Shell tips & tricks!

- The shell is a special case of the driver after all ;)

# Internal - did we cover?

- Replica Set vs Sharded Cluster pool of connections
- BSON sorting