

Overview

Contents

- MongoDB Releases
- Features
- Storage Engines

MongoDB Releases

- Naming
 - A.B.C
 - odd second numbers (B) for developer releases
 - even second numbers (B) for production releases
- Goals
 - 1 or 2 releases (A.B) per year
 - patches for some time (end of life)
 - support for
 - MongoDB: 18 months (end of support)
 - OpsManager: 12 months

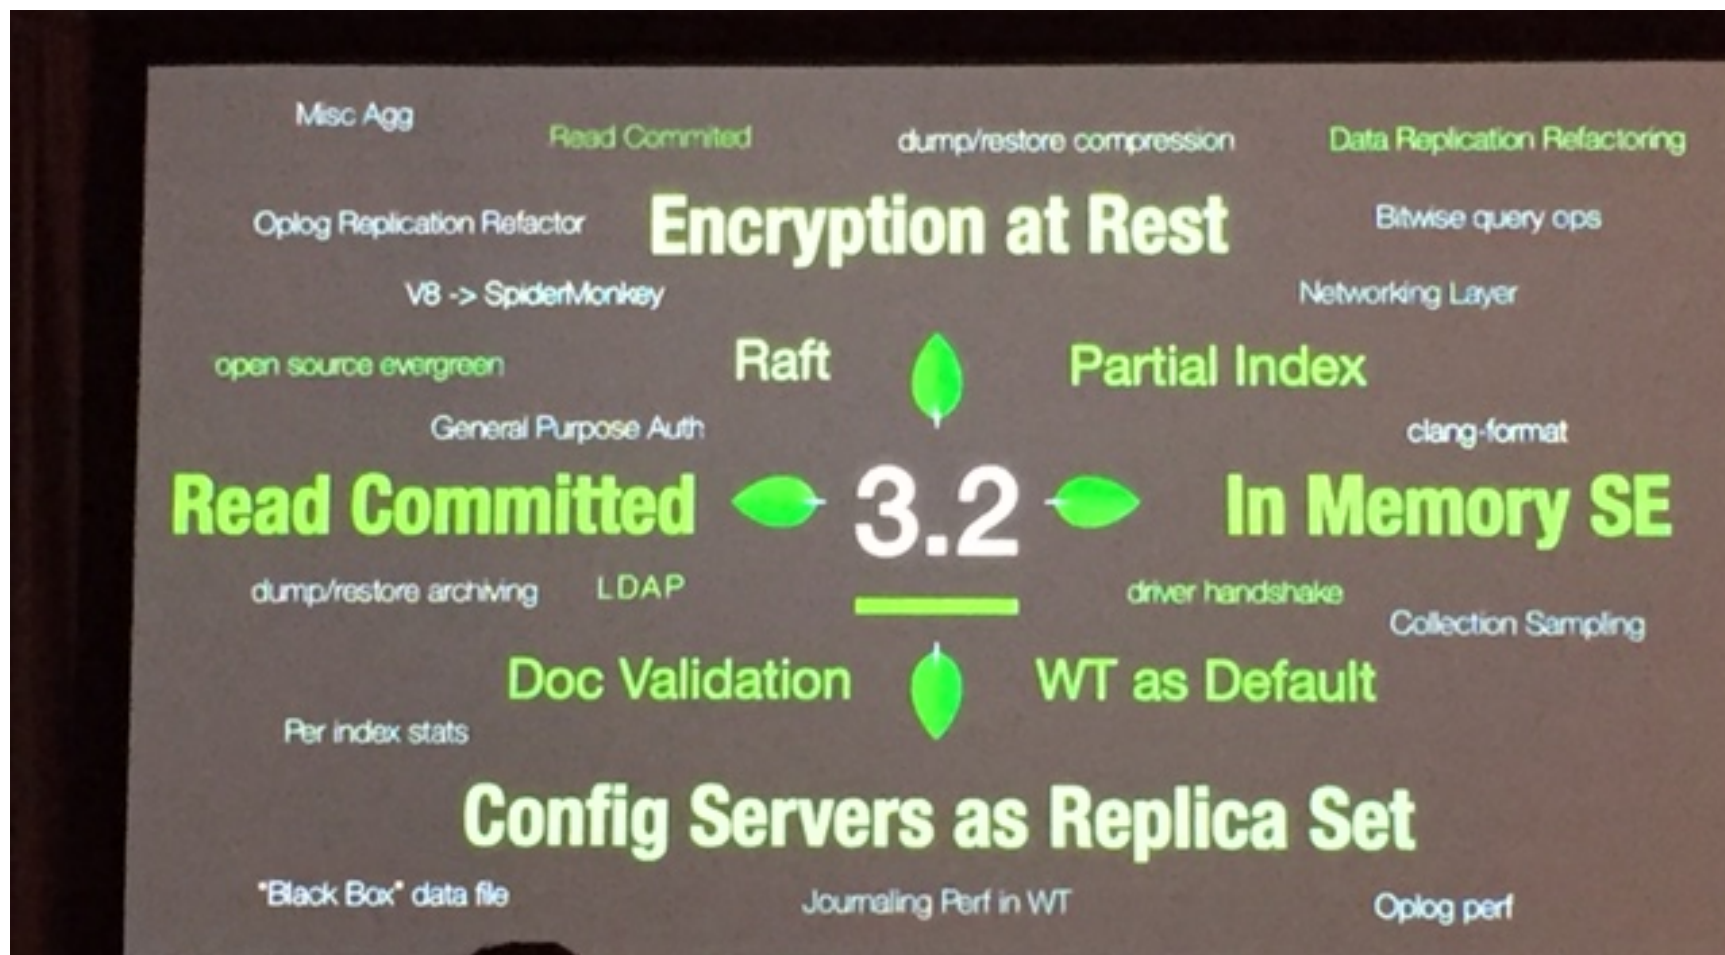
History of the MongoDB Product

- 2009 – 1.0
- 2010 – 1.2 master slave replication
- 2010 – 1.4 geo queries
- 2010 – 1.6 replica sets and sharding
- 2011 – 1.8 journaling (25 people)
- 2011 – 2.0 concurrency and fixes on everything above
- 2012 – 2.2 database level locking (100 people)
- 2013 – 2.4 text search and hashed sharding (325 people)
- 2014 – 2.6 security, auditing and new query planner (400 people)
- 2015 – 3.0 collection level locking in MMAPV1, document level locking with WiredTiger and OpsManager (MMS 1.6/OpsManager 1.8)
- 2015 – 3.2 WT as default, new storage engines (memory, encrypted), config servers as replica set (500 people) and OpsManager 2.0
- 2016 (Nov 1?) – 3.4, features list to be available internally in Feb.

MongoDB Upgrades

- We officially support:
 - MongoDB:
 - 2.4 (up to 03/31/2016)
 - 2.6 (up to 10/31/2016), sending notification on April 7, 2016
 - 3.0 and 3.2
 - OpsManager
 - 1.6 (up to 03/2016)
 - 1.8 and 2.0
- Some laggards on 1.8, 2.0 and 2.2
 - Product called “Extended Lifecycle Support Add-On”
- We often ask to upgrade
 - example: if running WT on 3.0, must use the latest (but not 3.0.10!)
- Complicated upgrades
 - Recommend the Consulting Package “Major Version Upgrade”

New features in 3.2



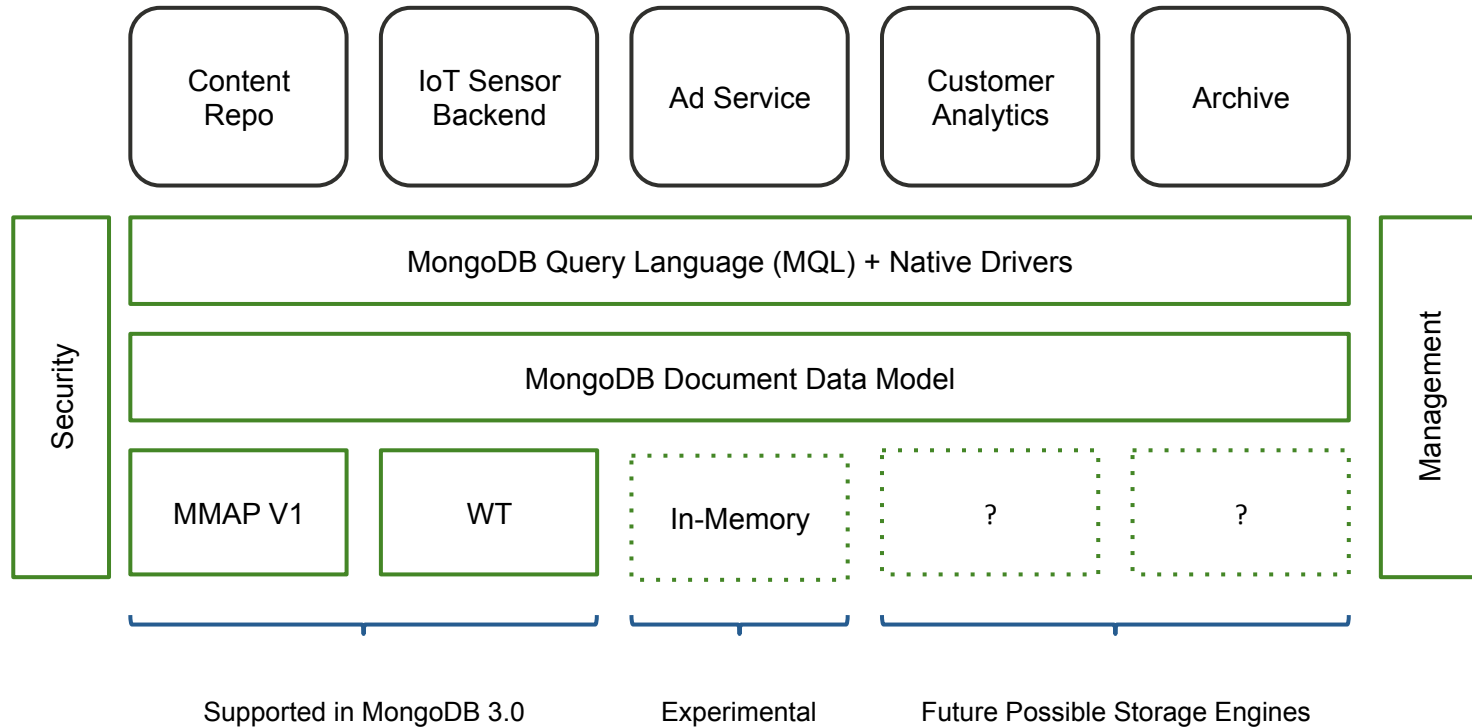
MongoDB 3.3/3.4

- Current plan (CONFIDENTIAL!!!) :
 - <https://docs.google.com/presentation/d/1OVAKk-aZKVttoyWJ5u8zNDUWajk-21D5ob-13qF1Wg4/edit#slide=id.p>
- Download the development build at:
<https://www.mongodb.com/customer-evaluation-downloads-development-versions>

How our customers use MongoDB?

- Stats from Cloud Manager users (Feb 2016)
 - (usually, does not include largest deployments)
 - 15% use sharding
 - 1.5% use tag-aware sharding
 - Sharded cluster 3X the size of non-sharded ones

Architecture



Main Features

- Flexible Schema
- Storage Engines
- Replication
- Sharding
- API

Flexible Schema

BSON (bsonspec.org)

- Has everything that JSON has and some extra stuff:
 - Binary data, ObjectId, UTC datetime, Timestamp, regex...
- So what is a document, really?
 - A document is a 32 bit integer followed by an element list and a null byte
 - An element list is an element followed by another element list or zero bytes
 - An element is a single byte (type) followed by the field name and the datum, i.e. the value of the field...

BSON continued

element	::=	"\x01" e_name double	Floating point
		"\x02" e_name string	UTF-8 string
		"\x03" e_name document	Embedded document
		"\x04" e_name document	Array
		"\x05" e_name binary	Binary data
		"\x06" e_name	Undefined — <i>Deprecated</i>
		"\x07" e_name (byte*12)	ObjectId
		"\x08" e_name "\x00"	Boolean "false"
		"\x08" e_name "\x01"	Boolean "true"
		"\x09" e_name int64	UTC datetime
		"\x0A" e_name	Null value
		"\x0B" e_name cstring cstring	Regular expression - The f stored in alphabetical order etc. locale dependent, 's' fo
		"\x0C" e_name string (byte*12)	DBPointer — <i>Deprecated</i>
		"\x0D" e_name string	JavaScript code
		"\x0E" e_name string	Deprecated
		"\x0F" e_name code_w_s	JavaScript code w/ scope
		"\x10" e_name int32	32-bit Integer
		"\x11" e_name int64	Timestamp
		"\x12" e_name int64	64-bit integer
		"\xFF" e_name	Min key
		"\x7F" e_name	Max key



BSON Notes

BSON does not include types for versioning, checksums, fixed or arbitrary decimal types... **kinda important if you're a bank (SERVER-1393, SERVER-4996)!!**

Strings are required to be valid UTF-8 encoded... **sucks if your hostname is "普查西" (SERVER-5215)**

or if you like your strings with non-English characters to sort properly (SERVER-1920)...

_id

- Each document must have an immutable `_id` field that is unique to the collection
- ObjectId by convention
 - 12 byte integer
 - a 4-byte value representing the seconds since the Unix epoch,
 - a 3-byte machine identifier
 - first three bytes of the (md5) hash of the machine host name, or of the mac/network address, or the virtual machine id (ex: `Digest::MD5.digest(Socket.gethostname)[0, 3]`)
 - a 2-byte process id, and
 - a 3-byte counter, starting with a random value.
 - Monotonically increasing, contains a date at the beginning
- Can store different types, but don't do that!
 - How will they sort? Sorting order for data types, e.g. double before string...

Storage Engines at a glance

MMAPv1

- MongoDB >= 1.0
- OS manages memory
- No compression
- All collections and indexes within one set of db files
- Padding/PowerOf2Sizes, in-place updates
- Supports 32-bit and 64-bit

WiredTiger

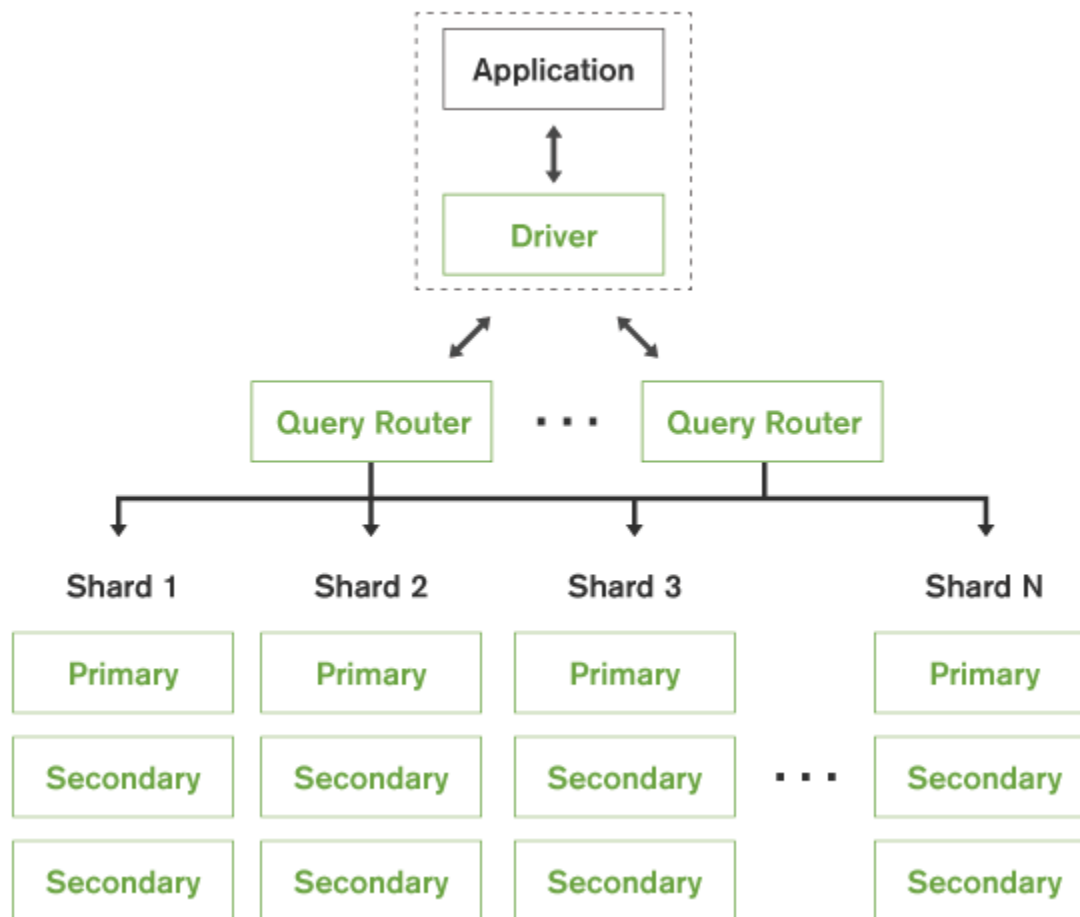
- MongoDB >= 3.0
- Server manages memory
- Compression (snappy, zlib)
- One collection or index per file
- No padding, always reallocate for writes
- Supports 64-bit only

Other storage engines

- inMemory
 - Beta in 3.2.0
- encrypted
- ephemeralForTest
 - Experiment in 3.0, kept as an example for the API

Replication

Sharding



References

- <https://www.mongodb.com/support-policy>
- <https://www.mongodb.com/products/consulting>
- <https://docs.mongodb.org/manual/release-notes/3.2>
- <https://www.mongodb.com/customer-evaluation-downloads-development-versions>
- <http://bsonspec.org/spec.html>