
MongoDB Administrators Training

Release 2.6

MongoDB, Inc.

November 25, 2014

Contents

1	Introduction	10
1.1	Warm Up	10
	Introductions	10
	Getting to Know You	10
	MongoDB Experience	11
	10gen	11
	Origin of MongoDB	11
1.2	MongoDB Overview	11
	Learning Objectives	11
	MongoDB is a Document Database	12
	An Example MongoDB Document	12
	Vertical Scaling	13
	Scaling with MongoDB	13
	Database Landscape	14
	MongoDB Deployment Models	15
1.3	MongoDB Stores Documents	16
	Learning Objectives	16
	JSON	16
	A Simple JSON Object	16
	JSON Keys and Values	16
	Example Field Values	17
	BSON	17
	BSON Hello World	17
	A More Complex BSON Example	18
	Documents, Collections, and Databases	18
	The <code>_id</code> Field	18
	ObjectIds	19
	Storing BSON Documents	19
	Padding Factor	19
	usePowerOf2Sizes	20
1.4	Exercise: Installing MongoDB	20
	Learning Objectives	20
	Production Releases	21
	Installing MongoDB	21
	Setup	21

Launch a <code>mongod</code>	22
Import Exercise Data	22
Launch a Mongo Shell	22
Explore Databases	23
Exploring Collections	23
Admin Commands	23
The MongoDB Data Directory	24
2 CRUD	25
2.1 Creating and Deleting Documents	25
Learning Objectives	25
Creating New Documents	25
Exercise: Inserting a Document	25
Implicit <code>_id</code> Assignment	26
Exercise: Assigning <code>_ids</code>	26
Inserts will fail if...	26
Exercise: Inserts will fail if...	26
Bulk Inserts	27
Ordered Bulk Insert	27
Exercise: Ordered Bulk Insert	27
Unordered Bulk Insert	28
Exercise: Unordered Bulk Insert	28
The Shell is a JavaScript Interpreter	28
Exercise: Creating Data in the Shell	29
Deleting Documents	29
Using <code>remove()</code>	29
Exercise: Removing Documents	29
Dropping a Collection	30
Exercise: Dropping a Collection	30
Dropping a Database	30
Exercise: Dropping a Database	30
2.2 Reading Documents	31
Learning Objectives	31
The <code>find()</code> Method	31
Query by Example	31
Exercise: Querying by Example	31
Querying Arrays	32
Exercise: Querying Arrays	32
Querying with Dot Notation	32
Exercise: Querying with Dot Notation	33
Exercise: Arrays and Dot Notation	33
Cursors	34
Exercise: Introducing Cursors	34
Exercise: Cursor Objects in the Mongo Shell	34
Cursor Methods	35
Exercise: Using <code>count()</code>	35
Exercise: Using <code>sort()</code>	35
The <code>skip()</code> Method	36
The <code>limit()</code> Method	36
The <code>distinct()</code> Method	36
Exercise: Using <code>distinct()</code>	36

2.3	Query Operators	37
	Learning Objectives	37
	Comparison Query Operators	37
	Exercise: Comparison Operators	37
	Logical Query Operators	38
	Exercise: Logical Operators (Setup)	38
	Exercise: Logical Operators	38
	Element Query Operators	39
	Exercise: Element Operators	39
	Array Query Operators	39
	Exercise: Array Operators	39
2.4	Updating Documents	40
	Learning Objectives	40
	The <code>update()</code> Method	40
	Parameters to <code>update()</code>	41
	<code>\$set</code> and <code>\$unset</code>	41
	Exercise: <code>\$set</code> and <code>\$unset</code>	41
	Update Operators	42
	Exercise: Update Operators	42
	<code>update()</code> Defaults to one Document	42
	Updating Multiple Documents	43
	Exercise: Multi-Update	43
	Array Operators	43
	Exercise: Array Operators	44
	The Positional <code>\$</code> Operator	44
	Exercise: The Positional <code>\$</code> Operator	45
	Upserts	45
	Upsert Mechanics	45
	Exercise: Upserts	45
	<code>save()</code>	46
	Exercise: <code>save()</code>	46
	Be Careful with <code>save()</code>	46
3	Indexes	48
3.1	Index Fundamentals	48
	Learning Objectives	48
	Why Indexes?	48
	Types of Indexes	49
	Exercise: Using <code>explain()</code>	49
	Results of <code>explain()</code>	50
	Understanding <code>explain()</code> Output	50
	Single-Field Indexes	51
	Creating an Index	51
	Indexes and Read/Write Performance	51
	Index Limitations	52
	Use Indexes with Care	52
3.2	Compound Indexes	52
	Learning Objectives	52
	Introduction to Compound Indexes	53
	The Order of Fields Matters	53
	Designing Compound Indexes	53

Example: A Simple Message Board	54
Load the Data	54
Start with a Simple Index	54
Query Adding <code>username</code>	54
Include <code>username</code> in Our Index	55
<code>ncanned > n</code>	55
A Different Compound Index	55
<code>nscanned == n == 2</code>	56
Let Selectivity Drive Field Order	56
Adding in the Sort	56
In-Memory Sorts	57
Avoiding an In-Memory Sort	57
General Rules of Thumb	58
3.3 Multikey Indexes	58
Learning Objectives	58
Introduction to Multikey Indexes	58
Example: Array of Numbers	58
Exercise: Array of Documents, Part 1	59
Exercise: Array of Documents, Part 2	59
Exercise: Array of Arrays, Part 1	60
Exercise: Array of Arrays, Part 2	60
How Multikey Indexes Work	61
Multikey Indexes and Sorting	61
Exercise: Multikey Indexes and Sorting	61
Limitations on Multikey Indexes	62
Example: Multikey Indexes on Multiple Fields	62
3.4 Hashed Indexes	62
Learning Objectives	62
What is a Hashed Index?	63
Why Hashed Indexes?	63
Limitations	63
Floating Point Numbers	63
Creating a Hashed Index	64
3.5 Geospatial Indexes	64
Learning Objectives	64
Introduction to Geospatial Indexes	64
Easiest to Start with 2 Dimensions	64
Location Field	65
Find Nearby Documents	65
Flat vs. Spherical Indexes	65
Flat Geospatial Index	65
Spherical Geospatial Index	66
Creating a 2d Index	66
Exercise: Creating a 2d Index	66
Inserting Documents with a 2d Index	67
Exercise: Inserting Documents with 2d Fields	67
Querying Documents Using a 2d Index	67
Example: Find Based on 2d Coords	68
Creating a 2dsphere Index	68
The geoJSON Specification	68
geoJSON Considerations	68

Simple Types of 2dsphere Objects	69
Polygons	69
Other Types of 2dsphere Objects	69
Exercise: Inserting geoJSON Objects (1)	70
Exercise: Inserting geoJSON Objects (2)	70
Exercise: Inserting geoJSON Objects (3)	71
Exercise: Creating a 2dsphere Index	71
Querying 2dsphere Objects	72
3.6 TTL Indexes	72
Learning Objectives	72
TTL Index Basics	72
Creating a TTL Index	72
Exercise: Creating a TTL Index	73
Exercise: Check the Collection	73
3.7 Text Indexes	73
Learning Objectives	73
What is a Text Index?	74
Creating a Text Index	74
Exercise: Creating a Text Index	74
Text Indexes are Similar to Multikey Indexes	74
Exercise: Inserting Texts	74
Querying a Text Index	75
Exercise: Querying a Text Index	75
Exercise: Querying Using Two Words	75
Search for a Phrase	75
Text Search Score	75
4 Replica Sets	76
4.1 Introduction to Replica Sets	76
Learning Objectives	76
Use Cases for Replication	76
High Availability (HA)	76
Disaster Recovery (DR)	77
Functional Segregation	77
Replication is Not Designed for Scaling	77
Replica Sets	78
Primary Server	79
Secondaries	79
Heartbeats	79
The Oplog	80
4.2 Elections in Replica Sets	80
Learning Objectives	80
Members and Votes	80
Calling Elections	81
Selecting a New Primary	81
Priority	81
Optime	82
Connections	82
When will a primary step down?	82
Exercise: Elections in Failover Scenarios	82

Scenario A: 3 Data Nodes in 1 DC	83
Scenario B: 3 Data Nodes in 2 DCs	83
Scenario C: 4 Data Nodes in 2 DCs	84
Scenario D: 5 Nodes in 2 DCs	84
Scenario E: 3 Data Nodes in 3 DCs	85
Scenario F: 5 data nodes in 3 DCs	85
4.3 Replica Set Roles and Configuration	86
Learning Objectives	86
Example: A Five-Member Replica Set Configuration	86
Configuration	86
Principal Data Center	86
Data Center 2	87
What about dc1-3 and dc2-2?	87
What about dc2-2?	87
4.4 The Oplog: Statement Based Replication	88
Learning Objectives	88
Binary Replication	88
Tradeoffs	88
Statement-Based Replication	89
Example	89
Replication Based on the Oplog	89
Create a Replica Set	90
ReplSetTest	90
Start the Replica Set	90
Status Check	90
Connect to the Primary	91
Create some Inventory Data	91
Perform an Update	91
View the Oplog	91
Operations in the Oplog are Idempotent	92
The Oplog Window	92
Sizing the Oplog	92
4.5 Write Concern	93
Learning Objectives	93
What happens to the write?	93
Answer	93
Balancing Durability with Performance	93
Defining Write Concern	94
Write Concern: { w : 1 }	94
Example: { w : 1 }	94
Write Concern: { w : 2 }	95
Example: { w : 2 }	95
Other Write Concerns	95
Write Concern: { w : "majority" }	96
Example: { w : "majority" }	96
Quiz: Which write concern?	96
Further Reading	96
4.6 Read Preference	97
What is Read Preference?	97
Use Cases	97

Not for Scaling	97
Read Preference Modes	98
Tag Sets	98
4.7 Exercise: Setting up a Replica Set	98
Overview	98
Create Data Directories	99
Launch Each Member	99
Status	99
Connect to a MongoDB Instance	100
Configure the Replica Set	100
Write to the Primary	100
Read from a Secondary	100
Review the Oplog	101
Changing Replica Set Configuration	101
Verifying Configuration Change	101
Further Reading	102
5 Sharding	103
5.1 Introduction to Sharding	103
Learning Objectives	103
Contrast with Replication	103
Sharding is Concerned with Scale	103
Vertical Scaling	104
The Working Set	104
Limitations of Vertical Scaling	104
Sharding Overview	105
A Model that Does Not Scale	105
A Scalable Model	105
Sharding Basics	106
Sharded Cluster Architecture	106
Mongos	107
Config Servers	108
Config Server Hardware Requirements	108
When to Shard	109
Possible Imbalance?	109
Balancing Shards	109
What is a Shard Key?	109
Targeted Query Using Shard Key	110
With a Good Shard Key	110
With a Bad Shard Key	110
Choosing a Shard Key	111
More Specifically	111
Cardinality	111
Non-Monotonic	111
Shards Should be Replica Sets	112
5.2 Balancing Shards	112
Learning Objectives	112
Chunks and the Balancer	112
Chunks in a Newly Sharded Collection	112
Chunk Splits	113
Pre-Splitting Chunks	113

Start of a Balancing Round	114
Balancing is Resource Intensive	114
Chunk Migration Steps	114
Concluding a Balancing Round	115
5.3 Shard Tags	115
Learning Objectives	115
Tags - Overview	115
Example: DateTime	115
Example: Location	116
Example: Premium Tier	116
Tags - Caveats	116
5.4 Exercise: Setting Up a Sharded Cluster	117
Learning Objectives	117
Our Sharded Cluster	117
Sharded Cluster Configuration	117
Build Our Data Directories	118
Initiate a Replica Set	118
Spin Up a Second Replica Set	118
A Third Replica Set	119
Status Check	119
Launch Config Servers	120
Launch the Mongos Processes	120
Add All Shards	120
Enable Sharding and Shard a Collection	120
Observe What Happens	121
6 Security	122
6.1 Security	122
Learning Objectives	122
Overview	122
Authentication Options	123
Authorization via MongoDB	123
Network Exposure Options	123
Encryption (SSL)	123
Native MongoDB Auth	124
Exercise: Create an Admin User, Part 1	124
Exercise: Create an Admin User, Part 2	125
Using MongoDB Roles	125
Exercise: Creating a readWrite User, Part 1	126
Exercise: Creating a readWrite User, Part 2	126
MongoDB Custom User Roles	126
7 Performance Troubleshooting	127
7.1 Performance Troubleshooting	127
Learning Objectives	127
mongostat and mongotop	127
Exercise: mongostat (setup)	127
Exercise: mongostat (run)	128
Exercise: mongostat (create index)	128
Exercise: mongotop	128

db.currentOp()	129
Exercise: db.currentOp()	129
db.collection.stats()	129
Exercise: Using Collection Stats	130
The Profiler	130
The Profiler (continued)	130
Exercise: Exploring the Profiler	130
db.serverStatus()	131
Exercise: Using db.serverStatus()	131
Analyzing profiler data	131
Performance Improvement Techniques	132
Performance Tips: Write Concern	132
Bulk Operations	132
Exercise: Comparing bulk inserts with mongostat	132
mongostat, bulk inserts with {w: 1}	133
Bulk inserts with {w: 3}	133
mongostat, bulk inserts with {w: 3}	133
Schema Design	134
Shard Key Considerations	134
Indexes and Performance	134
8 Backup and Recovery	135
8.1 Backup and Recovery	135
Disasters Do Happen	135
Human Disasters	136
Terminology: RPO vs. RTO	136
Terminology: DR vs. HA	136
Quiz	137
Backup Options	137
Document Level Backups - mongodump	137
mongodump	137
File System Level	138
Ensure Consistency	138
File System Backups: Pros and Cons	138
Document Level - mongorestore	138
File System Restores	139
Backup Sharded Cluster	139
Restore Sharded Cluster	139
Tips and Tricks	139
Backup Options	140
MMS Backup	140
Sharded Clusters	140
Under the Hood	140
Key Benefits	141
Point in Time Backups	141
Easy to Restore	141
Unlimited Restores	141
Fully Managed	141
Getting Started	141

1 Introduction

Warm Up (page 10) Activities to get the class started

MongoDB Overview (page 11) MongoDB philosophy and features.

MongoDB Stores Documents (page 16) The structure of data in MongoDB.

Exercise: Installing MongoDB (page 20) Install mongod on experiment with a few operations.

1.1 Warm Up

Introductions

- Who am I?
 - My role at MongoDB
 - My background and prior experience
-

Note:

- Tell the students about yourself:
 - Your role
 - Prior experience
-

Getting to Know You

- Who are you?
 - What role do you play in your organization?
 - What is your background?
 - Do you have prior experience with MongoDB?
-

Note:

- Ask students to go around the room and introduce themselves.
 - Make sure the names match the roster of attendees.
 - Ask about what roles the students play in their organization and note on attendance sheet.
 - Ask what software stacks students are using.
 - With MongoDB and in general.
 - Note this information as well.
-

MongoDB Experience

- Who has never used MongoDB?
- Who has some experience?
- Who has worked with production MongoDB deployments?
- Who is more of a developer?
- Who is more of operations person?

10gen

- MongoDB was initially created in 2008 as part of a hosted application stack.
- The company was originally called 10gen.
- As part of their overarching plan to create the 10gen platform, the company built a database.
- Suddenly everybody said: “I like that! Give me that database!”

Origin of MongoDB

- 10gen became a database company.
- In 2013, the company rebranded as MongoDB, Inc.
- The founders have other startups to their credit: DoubleClick, ShopWiki, Gilt.
- The motivation for the database came from observing the following pattern with application development.
 - The user base grows.
 - The associated body of data grows.
 - Eventually the application outgrows the database.
 - Meeting performance requirements becomes difficult.

1.2 MongoDB Overview

Learning Objectives

Upon completing this module students should understand:

- MongoDB vs. relational databases and key/value stores
- Vertical vs. horizontal scaling
- The role of MongoDB in the development stack
- The structure of documents in MongoDB
- Array fields
- Embedded documents
- Fundamentals of BSON

MongoDB is a Document Database

Documents are associative arrays like:

- Python dictionaries
- Ruby hashes
- PHP arrays
- JSON objects

An Example MongoDB Document


A MongoDB document expressed using JSON syntax.

```
{
  "a" : 3,
  "b" : [3, 2, 7],
  "c" : {
    "d" : 4 ,
    "e" : "asdf",
    "f" : true,
    "h" : ISODate("2014-10-23T01:19:40.732Z")
  }
}
```

Note:

- Where relational databases store rows, MongoDB stores documents.
 - Documents are hierarchical data structures.
 - This is a fundamental departure from relational databases where rows are flat.
-


Vertical Scaling



Note: Another difference is in terms of scalability. With an RDBMS:

- If you need to support a larger workload, you buy bigger machine.
 - The problem is, machines are not priced linearly.
 - The largest machines cost much more than commodity hardware.
 - If your application is successful, you may find you simply cannot buy a large enough a machine to support your workload.
-


Scaling with MongoDB



Note:

- MongoDB is designed to be horizontally scalable (linear).
 - MongoDB scales by enabling you to shard your data.
 - When you need more performance, you just buy another machine and add it to your cluster.
 - MongoDB is highly performant on commodity hardware.
-

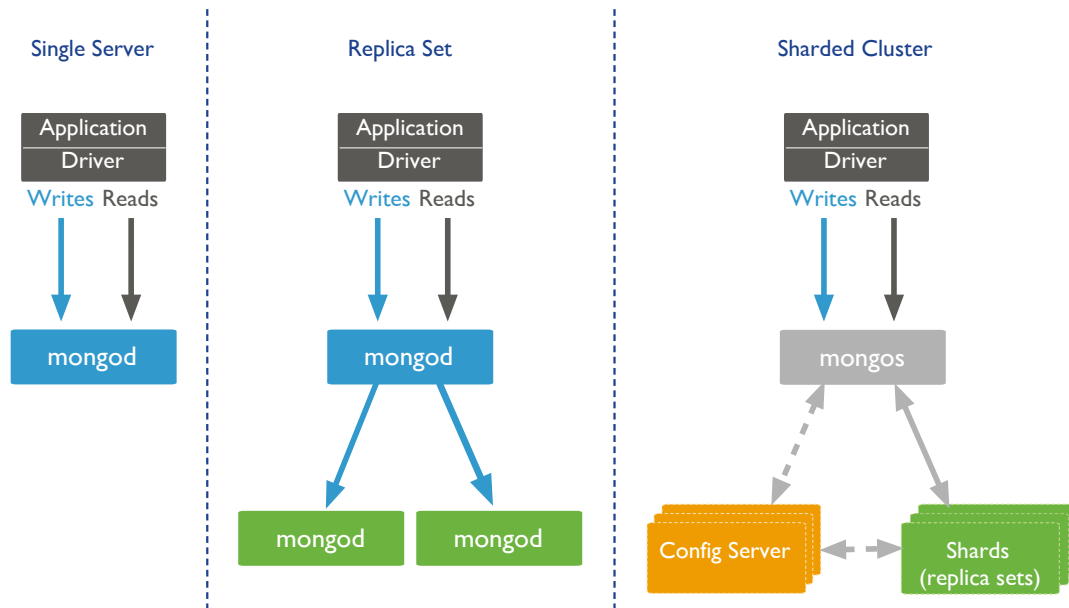
Database Landscape



Note:

- We've plotted each technology by scalability and functionality.
 - At the top left, are key/value stores like memcached.
 - These scale well, but lack features that make developers productive.
 - At the far right we have traditional RDBMS technologies.
 - These are full featured, but will not scale easily.
 - Joins and transactions are difficult to run in parallel.
 - MongoDB has nearly as much scalability as key-value stores.
 - Gives up only the features that prevent scaling.
 - We have compensating features that mitigate the impact of that design decision.
-

MongoDB Deployment Models



Note:

- MongoDB supports high availability through automated failover.
 - Do not use a single-server deployment in production.
 - Typical deployments use replica sets of 3 or more nodes.
 - The primary node will accept all writes, and possibly all reads.
 - Each secondary will replicate from another node.
 - If the primary fails, a secondary will automatically step up.
 - Replica sets provide redundancy and high availability.
 - In production, you typically build a fully sharded cluster:
 - Your data is distributed across several shards.
 - The shards are themselves replica sets.
 - This provides high availability and redundancy at scale.
-

1.3 MongoDB Stores Documents

Learning Objectives

Upon completing this module, students should understand:

- JSON
- BSON basics
- That documents are organized into collections
- ObjectIds
- Padding Factor

JSON

- JavaScript Object Notation
- Objects are associative arrays.
- They are composed of key-value pairs.

A Simple JSON Object

```
{  
  "firstname" : "Thomas",  
  "lastname" : "Smith",  
  "age" : 29  
}
```

JSON Keys and Values

- Keys must be strings.
- Values may be any of the following:
 - string (e.g., “Thomas”)
 - number (e.g., 29, 3.7)
 - true / false
 - null
 - array (e.g., [88.5, 91.3, 67.1])
 - object
- More detail at json.org¹.

¹<http://json.org/>

Example Field Values

```
{
  "first key" : "value" ,
  "second key" : {
    "first embedded key" : "first embedded value",
    "second embedded key" : "second embedded value"
  },
  "third key" : [
    "first array element",
    "second element",
    { "embedded key" : "embedded value" },
    [ 1, 2 ]
  ]
}
```

BSON

- MongoDB stores data as Binary JSON (BSON).
- MongoDB drivers send and receive data in this format.
- They map BSON to native data structures.
- BSON provides support for all JSON data types and several others.
- BSON was designed to be lightweight, traversable and efficient.
- See bsonspec.org².

Note: E.g., a BSON object will be mapped to a dictionary in Python.

BSON Hello World

```
// JSON
{ "hello" : "world" }

// BSON
"\x16\x00\x00\x00\x02hello\x00
\x06\x00\x00\x00world\x00\x00"
```

Note:

- \x16\x00\x00\x00 (document size)
- \x02 = string (data type of field value)
- hello\x00 (key/field name, \x00 is null and delimits the end of the name)
- \x06\x00\x00\x00 (size of field value including end null)
- world\x00 (field value)
- \x00 (end of the document)

²<http://bsonspec.org/#/specification>

A More Complex BSON Example

```
// JSON
{ "BSON" : [ "awesome", 5.05, 1986 ] }

// BSON
"\x3b\x00\x00\x00\x04BSON\x00\x26\x00
\x00\x00\x020\x00\x08\x00\x00
\x00awesome\x00\x011\x00\x33\x33\x33\x33\x33\x33
\x14\x40\x102\x00\xc2\x07\x00\x00
\x00\x00"
```

Documents, Collections, and Databases

- Documents are stored in collections.
- Collections are contained in a database.
- Example:
 - Database: products
 - Collections: books, movies, music
- Each database-collection combination defines a namespace.
 - products.books
 - products.movies
 - products.music

The `_id` Field

- All documents must have an `_id` field.
- The `_id` is immutable.
- If no `_id` is specified when a document is inserted, MongoDB will add the `_id` field.
- MongoDB assigns a unique ObjectId as the value of `_id`.
- Most drivers will actually create the ObjectId if no `_id` is specified.
- The `_id` field is unique to a collection (namespace).

ObjectId




Note:

- An ObjectId is a 12-byte value.
 - The first 4 bytes are a datetime reflecting when the ObjectId was created.
 - The next 3 bytes are the MAC address of the server.
 - Then a 2-byte process ID
 - Finally, 3 bytes that are monotonically increasing for each new ObjectId created within a collection.
-

Storing BSON Documents

- Each document may be a different size from the others.
- The maximum BSON document size is 16 megabytes.
- Documents are physically adjacent to each other on disk and in memory.
- If a document is updated in a way that makes it larger, MongoDB may move the document.
- This may cause fragmentation, resulting in unnecessary I/O.
- Strategies to reduce the effects of document growth:
 - Padding factor
 - usePowerOf2Sizes

Padding Factor



Note:

- Padding provides room for documents to grow into.
- As documents in a collection grow and need to be moved, MongoDB will begin to add padding.
- With padding, documents will not be as likely to move after update operations.
- The `padding factor` is a multiplier that defaults to 1 (no padding).

- At a padding factor of 2, the document will be inserted at twice its actual size.
 - This setting is not tunable; it is updated automatically.
-

usePowerOf2Sizes

- When a document must move to a new location this leaves a fragment.
 - MongoDB will attempt to fill this fragment with a new document eventually.
 - As of MongoDB 2.6, collections have a setting called `usePowerOf2Sizes` enabled by default for newly created collections.
 - This setting will round the size of the document up to the next power of 2.
 - E.g, a document that 118 bytes will be allocated 128 bytes.
 - If moved, the space can be filled with two 64-byte documents, four 32-byte documents, etc.
-

Note:

- Power of two sizes makes it easier for MongoDB to find new document to fill fragmented space.
 - Power of two sizing was introduced in MongoDB 2.4, but must be enabled using the `colMod` operation.
 - If a collection is read only, users should disable `usePowerof2Sizes` in MongoDB 2.6 and above.
-

1.4 Exercise: Installing MongoDB

Learning Objectives

Upon completing this exercise students should understand:

- How MongoDB is distributed
- How to install MongoDB
- Configuration steps for setting up a simple MongoDB deployment
- How to run MongoDB
- How to run the Mongo shell

Production Releases

64-bit production releases of MongoDB are available for the following platforms.

- Windows
- OSX
- Linux
- Solaris

Installing MongoDB

- Visit <http://www.mongodb.org/downloads>
- Download and install the appropriate package for your machine.
- Even-numbered builds are production releases, e.g., 2.4.x, 2.6.x.
- Odd-numbers indicate development releases, e.g., 2.5.x, 2.7.x.
- 32-bit versions should NOT be used in production (limited to 2GB of data). They are acceptable for training classes.

Setup

```
PATH=$PATH:path_to_mongodb/bin
```

```
sudo mkdir -p /data/db
```

```
sudo chmod -R 777 /data/db
```

Note:

- You might want to add the MongoDB bin directory to your path, e.g.
 - Once installed, create the MongoDB data directory.
 - Make sure you have write permission on this directory.
-

Launch a mongod

```
/<path_to_mongodb>/bin/mongod --help
```

```
/<path_to_mongodb>/bin/mongod
```

Note:

- Open a command shell and explore the mongod command.
- Then run mongod.

Note:

- Please verify that all students have successfully installed MongoDB.
 - Please verify that all can successfully launch a mongod.
-

Import Exercise Data

```
cd usb_drive
```

```
unzip sampledata.zip
```

```
cd sampledata
```

```
mongoimport -d sample -c tweets twitter.json
```

```
mongoimport -d sample -c zips zips.json
```

```
cd dump
```

```
mongorestore -d sample training
```

```
mongorestore -d sample digg
```

Note:

- Import the data provided on the USB drive into the *sample* database.
-

Launch a Mongo Shell

Open another command shell. Then type the following to start the Mongo shell.

```
mongo
```

Display available commands.

```
help
```

Explore Databases

Display available databases.

```
show dbs
```

To use a particular database we can type the following.

```
use <database_name>
```

```
db
```

Note:

- This assigns the variable `db` to a connection object for the selected database.
 - We can display the name of the database we are currently using by evaluating `db` in the mongo shell.
 - Highlight the power of the Mongo shell here.
 - It is a fully programmable JavaScript environment.
-

Exploring Collections

```
show collections
```

```
db.collection.help()
```

```
db.collection.find()
```

Note:

- Show the collections available in this database.
 - Show methods on the collection with parameters and a brief explanation.
 - Finally, we can query for the documents in a collection.
-

Admin Commands

- There are also a number of admin commands at our disposal.
- The following will shut down the mongod we are connected to through the Mongo shell.
- You can also just kill with Ctrl-C in the shell window from which you launched the mongod.

```
db.adminCommand( { shutdown : 1 } )
```

- Confirm that the mongod process has indeed stopped.
- Once you have, please restart it.

The MongoDB Data Directory

```
ls /data/db
```

- The mongod.lock file
 - This prevents multiple mongods from using the same data directory simultaneously.
 - Each MongoDB database directory has one .lock.
 - The lock file contains the process id of the mongod that is using the directory.
- Data files
 - The names of the files correspond to available databases.
 - A single database may have multiple files.

Note: Files for a single database increase in size as follows:

- sample.0 is 64 MB
 - sample.1 is 128 MB
 - sample.2 is 256 MB, etc.
 - This continues until sample.5, which is 2 GB
 - All subsequent data files are also 2 GB.
-

2 CRUD

Creating and Deleting Documents (page 25) Inserting documents into collections, deleting documents, and dropping collections.

Reading Documents (page 31) The `find()` command, query documents, dot notation, and cursors.

Query Operators (page 37) MongoDB query operators including: comparison, logical, element, and array operators.

Updating Documents (page 40) Using `update()` and associated operators to mutate existing documents.

2.1 Creating and Deleting Documents

Learning Objectives

Upon completing this module students should understand:

- How to insert documents into MongoDB collections.
- `_id` fields:
- How to remove documents from a collection
- How to remove a collection from a database
- How to remove a database from a MongoDB deployment

Creating New Documents

- Create documents using `insert()`.
- For example:

```
db.collection.insert( { "name" : "susan" } )
```

Exercise: Inserting a Document

Experiment with the following commands.

```
use sample
```

```
db.hellos.insert( { a : "hello, world!" } )
```

```
db.hellos.find()
```

Note:

- Make sure the students are performing the operations along with you.
 - Some students will have trouble starting things up, so be helpful at this stage.
-

Implicit `_id` Assignment

- We did not specify an `_id` in the document we inserted.
- If you do not assign one, MongoDB will create one automatically.
- The value will be of type `ObjectId`.

Exercise: Assigning `_ids`

Experiment with the following commands.

```
db.hellos.insert( { _id : 253, a : "a string" } )  
  
db.hellos.find()
```

Note:

- Note that you can assign an `_id` to be of almost any type.
 - It does not need to be an `ObjectId`.
-

Inserts will fail if...

- There is already a document in the collection with that `_id`.
- You try to assign an array to the `_id`.
- The argument is not a well-formed document.

Exercise: Inserts will fail if...

```
// fails because _id can't have an array value  
db.hellos.insert( { _id : [ 1, 2, 3 ] } )  
  
// succeeds  
db.hellos.insert( { _id : 3 } )  
  
// fails because of duplicate id  
db.hellos.insert( { _id : 3 } )  
  
// malformed document  
db.hellos.insert( { "hello" } )
```

Note:

- The following will fail because it attempts to use an array as an `_id`.

```
db.hellos.insert( { _id : [ 1, 2, 3 ] } )
```
 - The second insert with `_id : 3` will fail because there is already a document with `_id` of 3 in the collection.
 - The following will fail because it is a malformed document (i.e. no field name, just a value).
-

```
db.hellos.insert( { "hello" } )
```

Bulk Inserts

- MongoDB 2.6 introduced bulk inserts.
- You may bulk insert using an array of documents.
- The API has two core concepts:
 - Ordered bulk operations
 - Unordered bulk operations
- The main difference is in the way the operations are executed in bulk.

Note:

- In the case of an ordered bulk operation, every operation will be executed in the order they are added to the bulk operation.
 - In the case of an unordered bulk operation however there is no guarantee what order the operations are executed.
 - With an unordered bulk operation, the operations in the list may be reordered to increase performance.
-

Ordered Bulk Insert

- For ordered inserts MongoDB will stop processing inserts upon encountering an error.
- Meaning that only inserts occurring before an error will complete.
- The default setting for `db.collection.insert` is an ordered insert.
- See the next exercise for an example.

Exercise: Ordered Bulk Insert

Experiment with the following bulk insert.

```
db.things.insert( [ { _id : 19, type : "atom", symbol : "K" },  
                    { _id : 20, type : "car", color : "red" },  
                    { _id : 20, type : "planet", name : "Saturn" },  
                    { type : "office",  
                      street : "229 West 43rd Street, 5th Floor",  
                      city : "New York",  
                      state : "NY" } ] )  
  
db.things.find()
```

Note:

- This example has a duplicate key error.
 - Only the first 2 documents will be inserted.
-

Unordered Bulk Insert

- Pass `{ ordered : false }` to insert to perform unordered inserts.
- If any given insert fails, MongoDB will still attempt the others.
- The inserts may be executed in a different order from the way in which you specified them.
- The next exercise is very similar to the previous one.
- However, we are using `{ ordered : false }`
- One insert will fail, but all the rest will succeed.

Exercise: Unordered Bulk Insert

Experiment with the following bulk insert.

```
db.otherThings.insert( [ { _id : 19, type : "atom", symbol : "K" },
                          { _id : 20, type : "car", color : "red" },
                          { _id : 20, type : "planet", name : "Saturn" },
                          { type : "office",
                            street : "229 West 43rd Street, 5th Floor",
                            city : "New York",
                            state : "NY" } ],
                      { ordered : false } )

db.otherThings.find()
```

The Shell is a JavaScript Interpreter

- Sometimes it is convenient to create test data using a little JavaScript.
- The mongo shell is a fully-functional JavaScript interpreter. You may:
 - Define functions
 - Use loops
 - Assign variables
 - Perform inserts

Exercise: Creating Data in the Shell

Experiment with the following commands.

```
for (i=1; i<=10000; i++) {
  db.stuff.insert( { "a" : i } )
}

db.stuff.find()
```

Deleting Documents

You may delete documents from a MongoDB deployment in several ways.

- Use `remove()` to delete documents matching a specific set of conditions.
- Drop an entire collection.
- Drop a database.

Using `remove()`

- Remove documents from a collection using `remove()`.
- This command has one required parameter, a query document.
- All documents in the collection matching the query document will be removed.
- Pass an empty document to remove all documents.
- Prior to MongoDB 2.6 calling `remove()` with no parameters would remove all documents.
- Limit `remove()` to one document using `justOne`.

Exercise: Removing Documents

Experiment with removing documents. Do a `find()` after each `remove()` command below.

```
for (i=1; i<=20; i++) { db.testcol.insert( { _id : i, a : i } ) }

db.testcol.remove( { a : 1 } ) // Remove the first document

// $lt is a query operator that enables us to select documents that
// are less than some value. More on operators soon.
db.testcol.remove( { a : { $lt : 5 } } ) // Remove three more

db.testcol.remove( { a : { $lt : 10 } },
                  { justOne : true } ) // Remove one more

db.testcol.remove() // Error: requires a query document.

db.testcol.remove( { } ) // All documents removed
```

Dropping a Collection

- You can drop an entire collection with `db.collection.drop()`
- The collection and all documents will be deleted.
- It will also remove any metadata associated with that collection.
- Indexes are one type of metadata removed.
- More on meta data later.

Note: Mention that `drop()` is more performant than `remove` because of the lookup costs associated with `remove()`.

Exercise: Dropping a Collection

```
db.colToBeDropped.insert( { a : 1 } )
show collections // Shows the colToBeDropped collection

db.colToBeDropped.drop()
show collections // collection is gone
```

Dropping a Database

- You can drop an entire database with `db.dropDatabase()`
- This drops the database on which the method is called.
- It also deletes the associated data files from disk, freeing disk space.
- Beware that in the mongo shell, this does not change database context.

Exercise: Dropping a Database

```
use tempDB
db.testcol1.insert( { a : 1 } )
db.testcol2.insert( { a : 1 } )

show dbs // Here they are
show collections // Shows the two collections

db.dropDatabase()
show collections // No collections
show dbs // The db is gone

use sample // take us back to the sample db
```

2.2 Reading Documents

Learning Objectives

Upon completing this module students should understand:

- The query-by-example paradigm of MongoDB
- How to query on array elements
- How to query embedded documents using dot notation
- How the mongo shell and drivers use cursors
- Cursor methods: `.count()`, `.sort()`, `.skip()`, `.limit()`

The `find()` Method

- This is the fundamental method by which we read data from MongoDB.
- We have already used it in its basic form.
- `find()` returns a cursor that enables us to iterate through all documents matching a query.
- We will discuss cursors later.

Query by Example

- To query MongoDB, specify a document containing the key / value pairs you want to match
- You need only specify values for fields you care about.
- Other fields will not be used to exclude documents.
- The result set will include all documents in a collection that match.

Exercise: Querying by Example

Experiment with the following sequence of commands.

```
db.testcol.drop()
db.testcol.insert( [ { _id : 1, a : 5, b : 3 },
                    { _id : 3, b : 5, c : 12 },
                    { a : 7, b : 3 },
                    { c : 5, b : 7 } ] )
db.testcol.find()

db.testcol.find( { a : 5 } )

db.testcol.find( { b : 3, a : 7 } )
```

Note: Matching Rules:

- Any field specified in the query must be in each document returned.
- Values for returned documents must match the conditions specified in the query document.

- If multiple fields are specified, all must be present in each document returned.
 - Think of it as a logical AND for all fields.
-

Querying Arrays

- In MongoDB you may query array fields.
 - Specify a single value you expect to find in that array in desired documents.
 - Alternatively, you may specify an entire array in the query document.
 - As we will see later, there are also several operators that enhance our ability to query array fields.
-

Note: Students might find it helpful to think of an array field as having multiple values – one for each of its elements.

Exercise: Querying Arrays

Experiment with the following sequence of commands.

```
db.testcol.drop()
db.testcol.insert( [ { a : [ 1, 2, 3 ] },
                    { a : [ 3, 4, 5 ] },
                    { a : [ 5, 6, 7 ] } ] )

// These match documents where a contains the value specified
db.testcol.find( { a : 3 } )
db.testcol.find( { a : 5 } )

// These match documents where a equals the value specified
db.testcol.find( { a : [ 3, 5 ] } ) // no documents
db.testcol.find( { a : [ 3, 4, 5 ] } ) // only the second document
```

Note: Later, we'll see operators that will allow us to do things like match all documents where an array field contains any of a set of values.

Querying with Dot Notation

- Dot notation is used to query on fields in embedded documents.
- The syntax is:

```
"field1.field2" : value
```

- Put quotes around the field name when using dot notation.

Exercise: Querying with Dot Notation

```
db.buildings.insert(
  [ {
    type : "house",
    location : { streetNumber : 123,
                 street : "7th Ave" } },
    {
    type : "office",
    location : { streetNumber : 234,
                 street : "7th Ave",
                 floor : 7 } },
    {
    type : "apartment",
    location : { streetNumber : 335,
                 street : "43rd Street",
                 number : 745 } } ] )

db.buildings.find( { "location.street" : "7th Ave" } ) // Two matches
```

Exercise: Arrays and Dot Notation

Experiment with the following commands.

```
db.things.insert( [
  { type : "fruit",
    examples : [ { type : "banana", color : "yellow" },
                  { type : "apple", color : "red" },
                  { type : "mango", color : "red" } ] },
  { type : "cars",
    examples : [ { model : "Camaro", color : "red" },
                  { model : "Pinto", color : "yellow" },
                  { model : "Tacoma", color : "blue" } ] },
  { type : "planets",
    examples : [ { name : "Mars", color : "red" },
                  { name : "Venus", color : "blue" },
                  { name : "Earth", color : "blue" } ] } ] )

db.things.find( { "examples.color" : "blue" } ) // two documents
```

Note:

- This query finds documents where:
 - There is an `examples` field.
 - The `examples` field contains one or more embedded documents.
 - At least one embedded document has a field `color`.
 - The field `color` has the specified value (“blue”).
 - In this collection, `examples` is actually an array field.
 - The embedded documents we are matching are held within these arrays.
-

Cursors

- When you use `find()`, MongoDB returns a cursor.
- A cursor is a pointer to the result set
- You can get iterate through documents in the result using `next()`.
- By default, the mongo shell will iterate through 20 documents at a time.

Exercise: Introducing Cursors

Experiment with the following commands.

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
    db.testcol.insert( { a : Math.floor( Math.random() * 100 + 1 ),
                        b : Math.floor( Math.random() * 100 + 1 ) } )
}
db.testcol.find()

it
it
```

Note:

- With the `find()` above, the shell iterates over the first 20 documents.
 - `it` causes the shell to iterate over the next 20 documents.
 - Can continue issuing `it` commands until all documents are seen.
-

Exercise: Cursor Objects in the Mongo Shell

```
// Assigns the cursor returned by find() to a variable x
var x = db.testcol.find()

// Displays the first document in the result set.
x.next()

// True because there are more documents in the result set.
x.hasNext()

// Assigns the next document in the result set to the variable y.
y = x.next()

// Return value is the value of the a field of this document.
y.a

// Displaying a cursor prints the next 20 documents in the result set.
x
```

Cursor Methods

- `count()`: Returns the number of documents in the result set.
- `limit()`: Limits the result set to the number of documents specified.
- `skip()`: Skips the number of documents specified.

Exercise: Using `count()`

```
db.testcol.drop()
for (i=1; i<=100; i++) { db.testcol.insert( { a : i } ) }

// all 100
db.testcol.count()

// just 41 docs
db.testcol.count( { a : { $lt : 42 } } )

// Another way of writing the same query
db.testcol.find( { a : { $lt : 42 } } ).count( )
```

Note:

- You may pass a query document like you would to `find()`.
 - `count()` will count only the documents matching the query.
 - Will return the number of documents in the collection if you do not specify a query document.
 - The last query in the above achieves the same result because it operates on the cursor returned by `find()`.
-

Exercise: Using `sort()`

Experiment with the following sort commands.

```
db.testcol.drop()
for (i=1; i<=20; i++) {
  db.testcol.insert( { a : Math.floor( Math.random() * 10 + 1 ),
                      b : Math.floor( Math.random() * 10 + 1 ) } )
}

db.testcol.find()

// sort descending; use 1 for ascending
db.testcol.find().sort( { a : -1 } )

// sort by b, then a
db.testcol.find().sort( { b : 1, a : 1 } )

// $natural order is just the order on disk.
db.testcol.find().sort( { $natural : 1 } )
```

Note:

- Sort can be executed on a cursor until the point where the first document is actually read.
-

- If you never delete any documents or change their size, this will be the same order in which you inserted them.
 - Sorting two or more fields breaks the convention of javascript objects that key / value pairs are unordered.
 - In some drivers you may need to take special care with this.
 - For example, in Python, you would usually query with a dictionary.
 - But dictionaries are unordered in Python, so you would use an array of tuples instead.
-

The `skip()` Method

- Skips the specified number of documents in the result set.
- The returned cursor will begin at the first document beyond the number specified.
- Regardless of the order in which you specify `skip()` and `sort()` on a cursor, `sort()` happens first.

The `limit()` Method

- Limits the number of documents in a result set to the first `k`.
- Specify `k` as the argument to `limit()`
- Regardless of the order in which you specify `limit()`, `skip()`, and `sort()` on a cursor, `sort()` happens first.
- Helps reduce resources consumed by queries.

The `distinct()` Method

- Returns all values for a field found in a collection.
- Only works on one field at a time.
- Input is a string (not a document)

Exercise: Using `distinct()`

Experiment with the following commands and note what `distinct()` returns.

```
db.testcol.drop()
db.testcol.insert( [ { a : 2 , b : 3 },
                     { a : 2 },
                     { a : "hello" },
                     { a : "hello" },
                     { a : { hello : "world" } } ] )
db.testcol.distinct( "a" )
```

2.3 Query Operators

Learning Objectives

Upon completing this module students should understand the following types of MongoDB query operators:

- Comparison operators
- Logical operators
- Element query operators
- Operators on arrays

Comparison Query Operators

- `$lt`: Exists and is less than
- `$lte`: Exists and is less than or equal to
- `$gt`: Exists and is greater than
- `$gte`: Exists and is greater than or equal to
- `$ne`: Does not exist or does but is not equal to
- `$in`: Exists and is in a set
- `$nin`: Does not exist or is not in a set

Exercise: Comparison Operators

Experiment with the following.

```
db.testcol.drop()
for (i=1;i<=5;i++) { db.testcol.insert( { a : i } ) };
db.testcol.insert( { } ) // No "a" field
db.testcol.find()

db.testcol.find( { a : { $gte : 2 } } )

db.testcol.find( { a : { $ne : 2 } } )

db.testcol.find( { a : { $in : [ 3, 2 ] } } )

db.testcol.find( { a : { $nin : [ 3, 2 ] } } )
```

Logical Query Operators

- `$or`: Match either of two or more values
- `$not`: Used with other operators
- `$nor`: Match neither of two or more values
- `$and`: Match both of two or more values
 - This is the default behavior for queries specifying more than one condition.
 - Use `$and` if you need to include the same operator more than once in a query.

Exercise: Logical Operators (Setup)

Create a collection we can experiment with.

```
db.testcol.drop()
for (i=1; i<=3; i++) {
  for (j=1; j<=3; j++) {
    db.testcol.insert( { a : i, b : j } )
  }
};
db.testcol.insert( { b : 10 } ) // No "a" field
db.testcol.find()
```

Exercise: Logical Operators

Experiment with the following.

```
db.testcol.find( { $or : [ { a : 1 }, { b : 2 } ] } )

db.testcol.find( { a : { $not : { $gt : 3 } } } )

db.testcol.find( { $nor : [ { a : 3 }, { b : 3 } ] } )

db.testcol.find( { b : { $gt : 2 , $lte : 10 } } ) // and is implicit

db.testcol.find( { $and : [ { $or : [ { a : 1 }, { a : 2 } ] },
                             { $or : [ { b : 2 }, { b : 3 } ] } ] } )
```

Note:

- `db.testcol.find({ a : { $not : { $gt : 3 } } })`
 - Different from `db.testcol.find({ a : { $lte : 3 } })`
 - Returns all documents where `a` is less than or equal to 3
 - Also returns documents for which `a` does not exist
 - Without the use of `$and` in the last query above:
 - The second `$or` would replace the first.
 - The second `$or` would be the only condition evaluated in the query.
-

Element Query Operators

- `$exists`: Select documents based on the existence of a particular field.
- `$type`: Select documents based on their type.
- See [BSON types](http://docs.mongodb.org/manual/reference/bson-types)³ for reference on types.

Exercise: Element Operators

Experiment with the following.

```
db.testcol.drop()  
// by default, the mongo shell treats numbers as floating-point values  
db.testcol.insert( [ { a : 1 }, { b : 1 }, { a : NumberInt(2) },  
                    { b : "b" } ] )  
  
db.testcol.find( { a : { $exists : true } } )  
  
// type 1 is Double  
db.testcol.find( { b : { $type : 1 } } )  
  
// type 2 is String  
db.testcol.find( { b : { $type : 2 } } )  
  
// type 16 is 32-bit integer  
// use NumberInt(), NumberLong() to handle integers in the mongo shell  
db.testcol.find( { a : { $type : 16 } } )
```

Array Query Operators

- `$all`: Array field must contain all values listed.
- `$size`: Array must have a particular size. E.g., `$size : 2` means 2 elements in the array
- `$elemMatch`: All conditions must be matched by at least one element in the array

Exercise: Array Operators

Experiment with the following.

```
db.testcol.drop()  
db.testcol.insert( [ { a : [ 1, 2, 3, 4, 5 ] },  
                    { a : [ 1, 5 ] },  
                    { a : [ 1, 3, 5 ] } ] )  
  
db.testcol.find( { a : { $all : [ 1, 2 ] } } )  
  
db.testcol.find( { a : { $size : 3 } } )  
  
// at least one element must match both conditions  
db.testcol.find( { a : { $elemMatch : { $gte : 2, $lte : 4 } } } )
```

³<http://docs.mongodb.org/manual/reference/bson-types>

```
// at least one element must match either condition
// does not need to be the same element
db.testcol.find( { a : { $gte : 2, $lte : 4 } } )
```

Note:

- Comparing the last two queries demonstrates `$elemMatch`.
 - For the query using `$elemMatch` at least one element must match both conditions.
 - For the last query, there must be at least one element that matches each of the conditions. One element can match the `$gte` condition and another element can match the `$lte` condition.
-

2.4 Updating Documents

Learning Objectives

Upon completing this module students should understand

- The `update()` method
- The required parameters for `update()`
- Field update operators
- Array update operators
- The concept of an upsert and use cases.

The `update()` Method

- Mutate documents in MongoDB using `update()`.
- `update()` requires two parameters:
 - A query document used to select documents to be updated
 - An update document that specifies how selected documents will change
- `update()` cannot delete a document.

Parameters to update ()

- Keep the following in mind regarding the required parameters for update ()
- The query parameter:
 - Use the same syntax as with find () .
 - By default only the first document found is updated.
- The update parameter:
 - Take care to simply modify documents if that is what you intend.
 - Replacing documents in their entirety is easy to do by mistake.

\$set and \$unset

- Update one or more fields using the \$set operator.
- If the field already exists, using \$set will change its value.
- If the field does not exist, \$set will create it and set it to the new value.
- Any fields you do not specify will not be modified.
- You can remove a field using \$unset.

Exercise: \$set and \$unset

Experiment with the following. Do a find () after each update to view the results.

```
db.testcol.drop()
for (i=1; i<=5; i++) { db.testcol.insert( { _id : i, a : i, b : i } ) }

db.testcol.update( { _id : 3 }, { $set : { a : 6 } } )

db.testcol.update( { _id : 5 }, { $set : { c : 5 } } )

db.testcol.update( { _id : 5 }, { $set : { c : 7 , a : 7 } } )

db.testcol.update( { _id : 5 }, { d : 4 } )

db.testcol.update( { _id : 4 }, { $unset : { a : 1 } } )
```

Note:

- db.testcol.update({ _id : 3 }, { \$set : { a : 6 } }) just updates the a field.
 - db.testcol.update({ _id : 5 }, { \$set : { c : 5 } }) adds a c field to the matching document.
 - db.testcol.update({ _id : 5 }, { \$set : { c : 7 , a : 7 } }) modifies only the c and a fields.
 - db.testcol.update({ _id : 5 }, { d : 4 }) is the type of update that is probably a mistake.
 - It will replace the document with _id : 5 in its entirety.
 - The new document will be, { d : 4 }.
 - db.testcol.update({ _id : 4 }, { \$unset : a }) removes the a field.
-

Update Operators

- `$inc`: Increment a field's value by the specified amount.
- `$mul`: Multiply a field's value by the specified amount.
- `$rename`: Rename a field.
- `$set` (already discussed)
- `$unset` (already discussed)
- `$min`: Update only if value is smaller than specified quantity
- `$max`: Update only if value is larger than specified quantity
- `$currentDate`: Set the value of a field to the current date or timestamp.

Exercise: Update Operators

Experiment with the following update operators.

```
db.testcol.drop()
for (i=1; i<=5; i++) { db.testcol.insert( { _id : i, a : i, b : i } ) }
db.testcol.find()

db.testcol.update( { _id : 2 }, { $inc : { a : -3 } } )

db.testcol.update( { _id : 1 }, { $inc : { q : 1 } } )

db.testcol.update( { _id : 3 }, { $mul : { a : 4 } } )

db.testcol.update( { _id : 4 }, { $rename : { a : "xyz" } } )

db.testcol.update( { _id : 5 }, { $min : { a : 3 } } )

db.testcol.update( { _id : 1 },
                  { $currentDate : { c : { $type : "timestamp" } } } )
```

`update()` Defaults to one Document

- By default, `update()` modifies the first document found that matches the query.
- The default use case is one where there is only one document that fits the query.
- This is to reduce the chances of unintended collection scans for updates.

Updating Multiple Documents

- In order to update multiple documents, we use the third (optional) parameter to `update()`.
- The third parameter is an options document.
- Specify `multi: true` as one field in this document.
- Bear in mind that without an appropriate index, you may scan every document in the collection.

Exercise: Multi-Update

Use `db.testcol.find()` after each of these updates.

```
db.testcol.drop()
for (i=1; i<=5; i++) { db.testcol.insert( { _id : i, a : i } ) }

db.testcol.update( { _id : { $lt : 5 } }, { $set : { a : 6 } } )

db.testcol.update( { _id : { $lt : 5 } }, { $set : { a : 7 } },
                  { multi : true } )
```

Note:

- `db.testcol.update({ _id : { $lt : 5 } }, { $set : { a : 6 } })` only updates one document.
 - `db.testcol.update({ _id : { $lt : 5 } }, { $set : { a : 7 } }, { multi : true })` updates four documents.
-

Array Operators

- `$push`: Appends an element to the end of the array.
- `$pushAll`: Appends multiple elements to the end of the array.
- `$pop`: Removes one element from the end of the array.
- `$pull`: Removes all elements in the array that match a specified value.
- `$pullAll`: Removes all elements in the array that match any of the specified values.
- `$addToSet`: Appends an element to the array if not already present.

Note:

- These operators may be applied to array fields.
-

Exercise: Array Operators

Experiment with the following updates.

```
db.testcol.drop()
for (i=1; i<=5; i++) { db.testcol.insert(
  { _id : i, a : i, b: [ 1, 2, 3, 3, "NA", 6 ] } ) }

db.testcol.update( { _id : 1 }, { $push : { b : 3 } } )
db.testcol.update( { _id : 2 }, { $pushAll : { b : [ 1, 2, 3 ] } } )

db.testcol.update( { _id : 2 }, { $pop : { b : "" } } )

db.testcol.update( { _id : 3 }, { $pull : { b : 3 } } )
db.testcol.update( { _id : 4 },
  { $pullAll : { b : [ 1, 2, "NA", 4 ] } } )

db.testcol.update( { _id : 5 }, { $addToSet : { b : 2 } } )
db.testcol.update( { _id : 5 }, { $addToSet : { b : 4 } } )
```

Note:

- If you have time you might want to show the students:

```
db.testcol.update( { _id : 1 }, { $push : { b : [ 3, 12 ] } } )
```

- In { \$pop : { b : "" } }, the value passed for b is just a placeholder. Most values will work fine.
 - In { \$pullAll : { b : [1, 2, "NA", 4] } } the value 4 is not present in the array found in b. This has no effect on the update.
-

The Positional \$ Operator

- ⁴\$ is a positional operator that specifies an element in an array to update.
- It acts as a placeholder for the first element that matches the query document.
- \$ replaces the element in the specified position with the value given.
- Example:

```
db.collection.update(
  { <array> : value ... },
  { <update operator> : { "<array>.$" : value } }
)
```

⁴<http://docs.mongodb.org/manual/reference/operator/update/postional>

Exercise: The Positional \$ Operator

Experiment with the following commands.

```
db.testcol.drop()
for (i=1; i<=5; i++) { db.testcol.insert(
  { _id: i, a: i, b: [ 1, 2, 3, 3, "NA", 6 ] } ) }
db.testcol.find()

db.testcol.update( { b: "NA" }, { $set: { "b.$" : 11 } },
  { multi: true } )
db.testcol.find()
```

Upserts

- By default, if no document matches an update query, the `update()` method does nothing.
- By specifying `upsert: true`, `update()` will insert a new document if no matching document exists.
- The `db.collection.save()` method is syntactic sugar that performs an upsert if the `_id` is not yet present
- Syntax:

```
db.collection.update( <query document>, <update document>,
  { upsert: true } )
```

Upsert Mechanics

- Will update as usual if documents matching the query document exist.
- Will be an upsert if no documents match the query document.
 - MongoDB creates a new document using equality conditions in the query document.
 - Adds an `_id` if the query did not specify one.
 - Performs an update on the new document.

Exercise: Upserts

Experiment with the following upserts.

```
db.testcol.drop()
for (i=1; i<=5; i++) {
  db.testcol.insert( { _id: i, a: i, b: i } ) }
db.testcol.find()

db.testcol.update( { a: 4 }, { $inc: { b : 3 } }, { upsert: true } )

db.testcol.update( { a: 12 }, { $inc: { b : 3 } }, { upsert: true } )

db.testcol.update( { _id: 6, a: 6 }, { c: 155 }, { upsert: true } )
```

Note:

```
// updates the document with a: 4 by incrementing b
db.testcol.update( { a: 4 }, { $inc: { b : 3 } }, { upsert: true } )

// 1) creates a new document, 2) assigns an _id, 3) sets a to 12,
// 4) performs the update
db.testcol.update( { a: 12 }, { $inc: { b : 3 } }, { upsert: true } )

// 1) creates a new document, 2) sets _id: 6 and a: 6,
// 3) update deletes a and sets c to 155.
db.testcol.update( { _id: 6, a: 6 }, { c: 155 }, { upsert: true } )
```

save()

- Updates the document if the _id is found, inserts it otherwise
- Syntax:

```
db.collection.save( document )
```

Exercise: save()

```
db.testcol.drop()
for (i=1; i<=5; i++) { db.testcol.save( { _id : i, a : i, b : i } ) }
db.testcol.find()

// Look at the code for save. Note that it involves an upsert.
db.testcol.save

// new document, _id created
db.testcol.save( { a : 3 } )

db.testcol.save( { _id : 6, a : 6 } ) // new document

db.testcol.save( { _id : 3, a : 12, b : 12 } ) // update!!
```

Be Careful with save()

Be careful that you are not modifying stale data when using save(). For example:

```
db.testcol.drop()
db.testcol.insert( { _id : 2, a : 2, b : 2 } )

db.testcol.find( { _id : 2 } )
doc = db.testcol.findOne( { _id: 2 } )

db.testcol.update( { _id: 2 }, { $inc: { b : 1 } } )
db.testcol.find()

doc.c = 11
doc
```

```
db.testcol.save(doc)  // just lost our incrementing of b.  
db.testcol.find()
```

3 Indexes

Index Fundamentals (page 48) An introduction to MongoDB indexes.

Compound Indexes (page 52) Indexes on two or more fields.

Multikey Indexes (page 58) Indexes on array fields.

Hashed Indexes (page 62) Hashed Indexes.

Geospatial Indexes (page 64) Geospatial indexes: both those on legacy coordinate pairs and those supporting queries that calculate geometries on an earth-like sphere.

TTL Indexes (page 72) Time-To-Live Indexes.

Text Indexes (page 73) Free text indexes on string fields.

3.1 Index Fundamentals

Learning Objectives

Upon completing this module students should understand:

- The impact of indexing on read performance
- The impact of indexing on write performance
- How to choose effective indexes
- The utility of specific indexes for particular query patterns

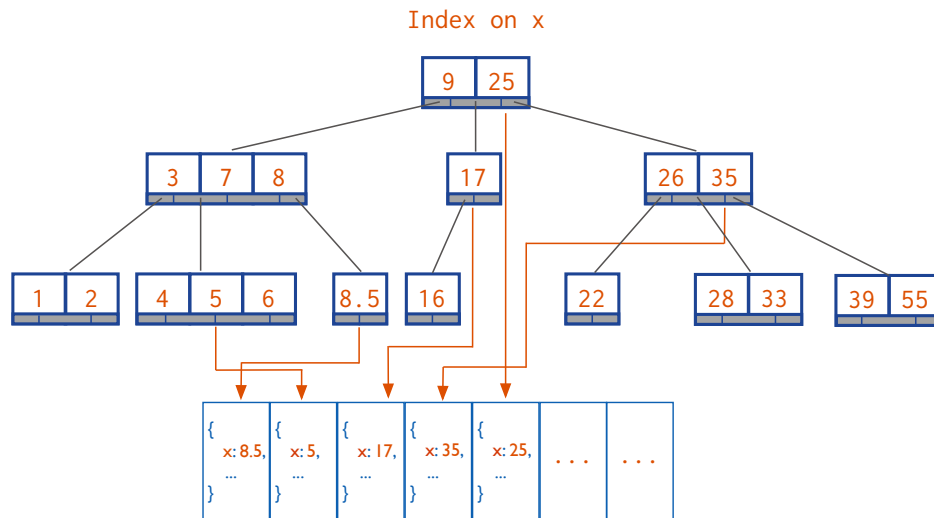
Note:

- Ask how many people in the room are familiar with indexes in a relational database.
 - If the class is already familiar with indexes, just explain that they work the same way in MongoDB.
-

Why Indexes?

Note:

- Without an index, in order to find all documents matching a query, MongoDB must scan every document in the collection.
 - This is murder for read performance, and often write performance, too.
 - If all your documents do not fit into memory, the system will page data in and out in order to scan the entire collection.
 - An index enables MongoDB to locate exactly which documents match the query and where they are located on disk.
 - MongoDB indexes are based on B-trees.
-



Types of Indexes

- Single-field indexes
- Compound indexes
- Multikey indexes
- Geospatial indexes
- Text indexes

Note:

- There are also hashed indexes and TTL indexes.
 - We will discuss those elsewhere.
-

Exercise: Using `explain()`

- Let's explore what MongoDB does for the following query by using `explain()`.
- We are projecting only `user.name` so that the results are easy to read.

```
db.tweets.find( { "user.followers_count" : 1000 },
               { "_id" : 0, "user.name": 1 } )
```

```
db.tweets.find( { "user.followers_count" : 1000 } ).explain()
```

Note:

- Make sure the students are using the sample database.
 - Review the structure of documents in the tweets collection by doing a `find()`.
 - We'll be looking at the user subdocument for documents in this collection.
-

Results of `explain()`

You will see results similar to the following.

```
{
  "cursor" : "BasicCursor",
  "isMultiKey" : false,
  "n" : 8,
  "nscannedObjects" : 51428,
  "nscanned" : 51428,
  "nscannedObjectsAllPlans" : 51428,
  "nscannedAllPlans" : 51428,
  "scanAndOrder" : false,
  "indexOnly" : false,
  "nYields" : 401,
  "nChunkSkips" : 0,
  "millis" : 161,
  "server" : "new-host-3.home:27017",
  "filterSet" : false
}
```

Understanding `explain()` Output

- `n` displays the number of documents that match the query.
- `nscannedObjects` displays the number of documents the retrieval engine considered during the query.
- `nscanned` displays how many documents in an existing index were scanned.
- An `nscanned` value much higher than `nreturned` indicates we need a different index.
- Given `nscannedObjects`, this query will benefit from an index.

Single-Field Indexes

- Based on a single field of the documents in a collection
- The field may be a top-level field
- You may also create an index on fields in embedded documents

Creating an Index

- The following creates a single-field index on `user.followers_count`.
- `explain()` indicated there will be a substantial performance improvement in handling this type of query.

```
db.tweets.ensureIndex( { "user.followers_count" : 1 } )  
db.tweets.find( { "user.followers_count" : 1000 } ).explain()
```

Note:

- `nscannedObjects` should now be a much smaller number, e.g., 8.
 - Operations teams are accustomed to thinking about indexes.
 - With MongoDB, developers need to be more involved in the creation and use of indexes.
-

Indexes and Read/Write Performance

- Indexes improve read performance for queries that are supported by the index.
- Inserts will be slower when there are indexes that MongoDB must also update.
- The speed of updates may be improved because MongoDB will not need to do a collection scan to find target documents.
- An index is modified any time a document:
 - Is inserted
 - Is deleted
 - Is updated in such a way that its indexed field changes
 - If an update causes a document to move on disk

Index Limitations

- You can have up to 64 indexes per collection.
- You should NEVER be anywhere close to that upper bound.
- Write performance will degrade to unusable at somewhere between 20-30.

Use Indexes with Care

- Every query should use an index.
- Every index should be used by a query.
- Any write operation that touches an indexed field will require each index to be updated.
- Indexes require RAM.
- Be judicious about the choice of key.

Note:

- If your system has limited RAM, then using the index will force other data out of memory.
 - When you need to access those documents, they will need to be paged in again.
-

3.2 Compound Indexes

Learning Objectives

Upon completing this module students should understand:

- What a compound index is.
- How compound indexes are created.
- The importance of considering field order when creating compound indexes.
- How to efficiently handle queries involving some combination of equality matches, ranges, and sorting.
- Some limitations on compound indexes.

Introduction to Compound Indexes

- It is common to create indexes based on more than one field.
- These are called `compound indexes`.
- You may use up to 31 fields in a compound index.
- You may not use hashed index fields.

The Order of Fields Matters

Specifically we want to consider how the index will be used for:

- Equality tests, e.g.,

```
db.example.find( { a : 15, b : 17 } )
```

- Range queries, e.g.,

```
db.example.find( { a : 15, b : { $lt : 85 } } )
```

- Sorting, e.g.,

```
db.example.find( { a : 15, b : 17 } ).sort( { b : -1 } )
```

Note:

- The order in which the fields are specified is of critical importance.
 - It is especially important to consider query patterns that require two or more of these operations.
-

Designing Compound Indexes

- Let's look at some guiding principles for building compound indexes.
- These will generally produce a good if not optimal index.
- You can optimize after a little experimentation.
- We will explore this in the context of a running example.

Example: A Simple Message Board

Requirements:

- Find all messages in a specified timestamp range.
- Select for whether the messages are anonymous or not.
- Sort by rating from highest to lowest.

Load the Data

```
a = [ { "timestamp" : 1, "username" : "anonymous", "rating" : 3 },
      { "timestamp" : 2, "username" : "anonymous", "rating" : 5 },
      { "timestamp" : 3, "username" : "sam", "rating" : 1 },
      { "timestamp" : 4, "username" : "anonymous", "rating" : 2 },
      { "timestamp" : 5, "username" : "martha", "rating" : 5 } ]
db.messages.insert(a)
```

Start with a Simple Index

Start by building an index on { timestamp : 1 }

```
db.messages.ensureIndex( { timestamp : 1 }, { name : "myindex" } )
```

Now let's query for messages with timestamp in the range 2 through 4 inclusive.

```
db.messages.find( { timestamp : { $gte : 2, $lte : 4 } } ).explain()
```

Note:

- Explain plan shows good performance, i.e. nscanned = n.
 - However, this does not satisfy our query.
 - Need to query again with { username : "anonymous" } as part of the query.
-

Query Adding username

```
db.messages.find( { timestamp : { $gte : 2, $lte : 4 },
                  username : "anonymous" } ).explain()
```

Note:

- Let's add the user field to our query.
 - Now nscanned > n.
-

Include username in Our Index

```
db.messages.dropIndex( "myindex" );
db.messages.ensureIndex( { timestamp : 1, username : 1 },
                        { name : "myindex" } )
db.messages.find( { timestamp : { $gte : 2, $lte : 4 },
                  username : "anonymous" } ).explain()
```

Note:

- nscanned is still greater than n.
 - Why?
-

ncanned > n

timestamp	username
1	"anonymous"
2	"anonymous"
3	"sam"
4	"anonymous"
5	"martha"

Note:

- The index we have created stores the range values before the equality values.
 - The documents with timestamp values 2, 3, and 4 were found first.
 - Then the associated anonymous values had to be evaluated.
-

A Different Compound Index

Drop the index and build a new one with user.

```
db.messages.dropIndex( "myindex" );
db.messages.ensureIndex( { username : 1 , timestamp : 1 },
                        { name : "myindex" } )
db.messages.find( { timestamp : { $gte : 2, $lte : 4 },
                  username : "anonymous" } ).explain()
```

Note:

- Now nscanned is 2 and n is 2.
-

nscanned == n == 2

username	timestamp
"anonymous"	1
"anonymous"	2
"anonymous"	4
"sam"	2
"martha"	5

Note:

- This illustrates why.
 - There is fundamental difference in the way the index is structured.
 - This supports a more efficient treatment of our query.
-

Let Selectivity Drive Field Order

- Order fields in a compound index from most selective to least selective.
- Usually, this means equality fields before range fields.
- When dealing with multiple equality values, start with the most selective.
- If a common range query is more selective instead (rare), specify the range component first.

Adding in the Sort

Finally, let's add the sort and run the query.

```
db.messages.find( {  
    timestamp : { $gte : 2, $lte : 4 },  
    username  : "anonymous"  
} ).sort( { rating : -1 } ).explain();
```

Note:

- Note that the `scanAndOrder` field is set to true.
 - This means that MongoDB had to perform a sort in memory.
 - In memory sorts for queries that retrieve large numbers of documents can degrade performance significantly.
 - Especially, if they are used frequently.
-

In-Memory Sorts

Let's modify the index again to allow the database to sort for us.

```
db.messages.dropIndex( "myindex" );
db.messages.ensureIndex( { username : 1 , timestamp : 1, rating : 1 },
                        { name : "myindex" } );
db.messages.find( {
    timestamp : { $gte : 2, $lte : 4 },
    username : "anonymous"
} ).sort( { rating : -1 } ).explain();
```

Note:

- The explain plan remains unchanged.
 - The field being sorted on comes after the range fields.
 - The index does not store entries in order by rating.
 - To have the index structured this way we need to specify the field for sorting (rating) before the range field (timestamp).
 - Note that this requires us to consider a tradeoff.
-

Avoiding an In-Memory Sort

Rebuild the index as follows.

```
db.messages.dropIndex( "myindex" );
db.messages.ensureIndex( { username : 1, rating : 1, timestamp : 1 },
                        { name : "myindex" } );
db.messages.find( {
    timestamp : { $gte : 2, $lte : 4 },
    username : "anonymous"
} ).sort( { rating : -1 } ).explain();
```

Note:

- We have a tradeoff between `nscanned` and `scanAndOrder`
 - Now `scanAndOrder` is set to `false`.
 - However, `nscanned` is 3 and `n` is 2.
 - This is the best we can do in this case and in this situation is fine.
 - However, if `nscanned` is much larger than `n`, this might not be the best index.
 - You will have to evaluate your use case to determine how to make this tradeoff.
-

General Rules of Thumb

- Equality before range.
- Equality before sorting.
- Sorting before range.

3.3 Multikey Indexes

Learning Objectives

Upon completing this module, students should understand:

- What a multikey index is
- When MongoDB will use a multikey index to satisfy a query
- How multikey indexes work
- How multikey indexes handle sorting
- Some limitations on multikey indexes

Introduction to Multikey Indexes

- A multikey index is an index on an array.
- An index entry is created on each value found in the array.
- Multikey indexes can support primitives, documents, or sub-arrays.
- There is nothing special that you need to do to create a multikey index.
- You created them using `ensureIndex()` just as you would with an ordinary single-field index.
- If there is an array as a value for an indexed field, the index will be multikey on that field.

Example: Array of Numbers

```
db.testcol.drop()
db.testcol.ensureIndex( { x : 1 } )
a = [ { x : [ 1, 2, 3 ], y : [ "a", "b" ] },
      { x : [ 3, 4 ], y : [ "a", "b" ] },
      { x : [ 4, 5 ], y : [ "a", "b" ] },
      { x : 3, y : [ "a", "b" ] }, { x : 4, y : [ "a", "b" ] } ]
db.testcol.insert( a )
db.testcol.find( { x : 3 } )
db.testcol.find( { x : 3 } ).explain()
db.testcol.find( { "x.2" : 3 } )
db.testcol.find( { "x.2" : 3 } ).explain()
```

Note:

```
// Used the index
db.testcol.find( { x : 3 } ).explain()

// One document found.
// Index not used, because it is naive to position.
db.testcol.find( { "x.2" : 3 } )
```

Exercise: Array of Documents, Part 1

Create a collection and add an index on the x field:

```
db.testcol.drop()
b = [ { x : [ { name : "Alice", number : 1 }, { name : "Bob", number : 2 },
             { name : "Cherry", number : 3 } ] },
      { x : [ { name : "Cherry", number : 3 },
             { name : "Dan", number : 4 } ] },
      { x : [ { name : "Dan", number : 4 },
             { name : "Erica", number : 5 } ] },
      { x : { name : "Cherry", number : 3 } },
      { name : "Dan", number : 4 } ]
db.testcol.insert(b)
db.testcol.ensureIndex( { x : 1 } )
db.testcol.find()
```

Note:

- In this collection there are four documents.
 - In each document is an array, x, containing subdocuments.
 - Each subdocument has a name and number in it
 - Always the same number for each name.
-

Exercise: Array of Documents, Part 2

For each of the three queries below:

- How many documents will be returned?
- Will it use our multi-key index? Why or why not?
- If a query will not use the index, which index will it use?

```
db.testcol.find( { x : { name : "Cherry", number : 3 } } )
db.testcol.find( { x : { number : 3 } } )
db.testcol.find( { "x.number" : 3 } )
```

Note:

```
// 3 documents
db.testcol.find( { x : { name : "Cherry", number : 3 } } )
// Used the multi-key index. We pass a complete document for ``x``.
db.testcol.find( { x : { name : "Cherry", number : 3 } } ).explain()
```

```
// 3 documents
db.testcol.find( { "x.number" : 3 } )
// Does not use the multi-key index.
// `x.number` is only part of the document that is indexed.
db.testcol.find( { "x.number" : 3 } ).explain()

// We would need to add an index such as this.
db.testcol.ensureIndex( { "x.number" : 1 } )
db.testcol.find( { "x.number" : 3 } ).explain()
```

Exercise: Array of Arrays, Part 1

Add some documents and create an index:

```
db.testcol.drop()
db.testcol.ensureIndex( { x : 1 } )
c = [ { x : [ [ 1, 2 ], [ 2, 3 ], [ 3, 4 ] ] },
      { x : [ [ 3, 4 ], [ 4, 5 ] ] },
      { x : [ [ 4, 5 ], [ 5, 6 ] ] },
      { x : [ 3, 4 ] },
      { x : [ 4, 5 ] } ]
db.testcol.insert(c)
db.testcol.find()
```

Exercise: Array of Arrays, Part 2

For each of the queries below:

- How many documents will be returned?
- Does the query use the multi-key index? Why or why not?
- If the query does not use the index, what is an index it could use?

```
db.testcol.find( { x : [ 3, 4 ] } )
db.testcol.find( { x : 3 } )
db.testcol.find( { "x.1" : [ 4, 5 ] } )
db.testcol.find( { "x.1" : 4 } )
```

Note:

```
// 3 documents
db.testcol.find( { x : [ 3, 4 ] } )
// Uses the multi-key index
db.testcol.find( { x : [ 3, 4 ] } ).explain()

// One document found, where the element of x is just a number.
db.testcol.find( { x : 3 } )
// Used the index
db.testcol.find( { x : 3 } ).explain()

db.testcol.find( { "x.1" : [ 4, 5 ] } ).explain()
// Does not use the multi-key index, because it is naive to position.
db.testcol.find( { "x.1" : 4 } ).explain()
```

How Multikey Indexes Work

- Each array element is given one entry in the index.
- So an array with 17 elements will have 17 entries – one for each element.
- Multikey indexes can take up much more space than standard indexes.

Multikey Indexes and Sorting

- If you sort using a multikey index:
 - A document will appear at the first position where a value would place the document.
 - It does not appear multiple times.
- This applies to array values generally.
- It is not a specific property of multikey indexes.

Exercise: Multikey Indexes and Sorting

```
db.testcol.drop()
a = [ { x : [ 1, 11 ] }, { x : [ 2, 10 ] }, { x : [ 3 ] },
      { x : [ 4 ] }, { x : [ 5 ] } ]
db.testcol.insert(a)

db.testcol.ensureIndex( { x : 1 } )

// x : [ 1, 11 ] array comes first. It contains the lowest value.
db.testcol.find().sort( { x : 1 } )

// x : [ 1, 11 ] array still comes first. Contains the highest value.
db.testcol.find().sort( { x : -1 } )
```

Note:

```
// x : [ 1, 11 ] array comes first, because it contains the lowest value
db.testcol.find().sort( { x : 1 } )

// x : [ 1, 11 ] array still comes first, because it contains the highest value
db.testcol.find().sort( { x : -1 } )
```

Limitations on Multikey Indexes

- You cannot create a compound index using more than one array-valued field.
- This is because of the combinatorics.
- For a compound index on two array-valued fields you would end up with $N * M$ entries for one document.
- You cannot have a hashed multikey index.
- You cannot have a shard key use a multikey index
- We discuss shard keys in another module.
- The index on the `_id` field cannot become a multikey index.

Example: Multikey Indexes on Multiple Fields

```
db.testcol.drop()
db.testcol.ensureIndex( { x : 1, y : 1 } )

// no problems yet
db.testcol.insert( { _id : 1, x : 1, y : 1 } )

// still OK
db.testcol.insert( { _id : 2, x : [ 1, 2 ], y : 1 } )

// still OK
db.testcol.insert( { _id : 3, x : 1, y : [ 1, 2 ] } )

// Won't work
db.testcol.insert( { _id : 4, x : [ 1, 2 ], y : [ 1, 2 ] } )
```

3.4 Hashed Indexes

Learning Objectives

Upon completing this module, students should understand:

- What a hashed index is.
- When to use one.

What is a Hashed Index?

- Hashed indexes are based on field values like any other index.
- The difference is that the values are hashed and it is the hashed value that is indexed.
- The hashing function collapses sub-documents and computes the hash for the entire value.
- MongoDB can use the hashed index to support equality queries.
- Hashed indexes do not support multi-key indexes, i.e. indexes on array fields.
- Nor do they support range queries.

Why Hashed Indexes?

- In MongoDB, the primary use for hashed indexes is to support sharding a collection using a hashed shard key.
- In some cases, the field we would like to use to shard data would make it difficult to scale using sharding.
- Using a hashed shard key to shard a collection ensures an even distribution of data and overcomes this problem.
- See [Shard a Collection Using a Hashed Shard Key](#)⁵ for more details.
- We discuss sharding in detail in another module.

Limitations

- You may not create compound indexes that have hashed index fields
- You may not specify a unique constraint on a hashed index
- You can create both a hashed index and a non-hashed index on the same field.

Note:

- For a field on which there is both a hashed index and a non-hashed index, MongoDB will use the non-hashed index for range queries.
-

Floating Point Numbers

- MongoDB hashed indexes truncate floating point numbers to 64-bit integers before hashing.
- Do not use a hashed index for floating point numbers that cannot be reliably converted to 64-bit integers.
- MongoDB hashed indexes do not support floating point values larger than 2^{53} .

⁵<http://docs.mongodb.org/manual/tutorial/shard-collection-with-a-hashed-shard-key/>

Creating a Hashed Index

Create a hashed index using an operation that resembles the following. This operation creates a hashed index for the active collection on the a field.

```
db.active.ensureIndex( { a: "hashed" } )
```

3.5 Geospatial Indexes

Learning Objectives

Upon completing this module, students should understand:

- Use cases of geospatial indexes
- The two types of geospatial indexes
- How to create 2d geospatial indexes
- How to query for documents in a region
- How to create 2dsphere indexes
- Types of geoJSON objects
- How to query using 2dsphere indexes

Introduction to Geospatial Indexes

We can use geospatial indexes to quickly determine geometric relationships:

- All points within a certain radius of another point.
- Whether or not points fall within a polygon
- Whether or not two polygons intersect

Easiest to Start with 2 Dimensions

- Initially, it is easiest to think about geospatial indexes in two dimensions.
- And one type of geospatial index in MongoDB is a flat 2d index.
- With a geospatial index we can, for example, search for nearby items.
- This is the type of service that many phone apps provide when, say, searching for a nearby cafe.
- We might have a query location identified by an X in a 2d coordinate system.

Note:

- Instructor, please draw a 2d coordinate system with axes for lat and lon.
 - Draw a red (or some other color) x to represent the query document.
-

Location Field

- A geospatial index is based on a location field within documents in a collection.
- The structure of location values depends on the type of geospatial index.
- We will go into more detail on this in a few minutes.
- We can identify other documents in this collection with Xs in our 2d coordinate system.

Note:

- Draw several Xs to represent other documents.
-

Find Nearby Documents

- A geospatial index enables us to efficiently query a collection based on geometric relationships between documents and the query.
- For example, we can quickly locate all documents within a certain radius of our query location.
- In this example, we've illustrated a `$near` query in a 2d geospatial index.

Flat vs. Spherical Indexes

There are two types of geospatial indexes:

- Flat, made with a `2d` index
- Two-dimensional spherical, made with the `2dsphere` index
 - Takes into account the curvature of the earth.
 - Joins any two points using a geodesic or “great circle arc”.
 - Deviates from flat geometry as you get further from the equator, and as your points get further apart.

Flat Geospatial Index

- This is a Cartesian treatment of coordinate pairs.
- E.g., would NOT know that the shortest path from Canada to Siberia is over the North Pole (if units are degrees).
- Can be used to describe any flat surface.
- Recommended if:
 - You have legacy coordinate pairs (MongoDB 2.2 or earlier).
 - You do not plan to use geoJSON objects such as LineStrings or Polygons.
 - You are not going to use points far enough North or South to worry about the Earth's curvature.

Spherical Geospatial Index

- Knows about the curvature of the Earth.
- If you want to plot the shortest path from the Klondike to Siberia, this will know to go over the North Pole.
- Uses geoJSON objects (Points, LineString, and Polygons).
- Coordinate pairs are converted into geoJSON Points.

Creating a 2d Index

Creating a 2d index:

```
db.collection.ensureIndex(  
  { field_name : "2d", <optional additional field> : <value> },  
  { <optional options document> } )
```

Possible options key-value pairs:

- min : <lower bound>
- max : <upper bound>
- bits : <bits of precision for geohash>

Exercise: Creating a 2d Index

Create a 2d index on the collection `testcol` with:

- A min value of -20
- A max value of 20
- 10 bits of precision
- The field indexed should be `xy`.

Note: Answer:

```
db.testcol.ensureIndex( { xy : "2d" }, { min : -20, max : 20, bits : 10 } )
```

Inserting Documents with a 2d Index

There are two accepted formats:

- Legacy coordinate pairs
- Document with the following fields specified:
 - lng (longitude)
 - lat (latitude)

Exercise: Inserting Documents with 2d Fields

- Insert 2 documents into the 'twoD' collection.
- Assign 2d coordinate values to the xy field of each document.
- Longitude values should be -3 and 3 respectively.
- Latitude values should be 0 and 0.4 respectively.

Note: Answer:

```
db.twoD.insert( { xy : [ -3, 0 ] } ) // legacy coordinate pairs
db.twoD.insert( { xy : { lng : 3, lat : 0.4 } } ) // document with lng, lat
db.twoD.find() // both went in OK
db.twoD.insert( { xy : 5 } ) // insert works fine
// Keep in mind that the index doesn't apply to this document.
db.twoD.insert( { xy : [ 0, -500 ] } )
// Generates an error because -500 isn't between +/-20.
db.twoD.insert( { xy : [ 0, 0.00003 ] } )
db.twoD.find()
// last insert worked fine, even though the position resolution is below
// the resolution of the Geohash.
```

Querying Documents Using a 2d Index

- Use \$near to retrieve documents close to a given point.
- Use \$geoWithin to find documents with a shape contained entirely within the query shape.
- Use the following operators to specify a query shape:
 - \$box
 - \$polygon
 - \$center (circle)

Example: Find Based on 2d Coords

Write a query to find all documents in the testcol collection that have an xy field value that falls entirely within the circle with center at [-2.5, -0.5] and a radius of 3.

```
db.testcol.find( { xy : { $geoWithin : { $center : [ [ -2.5, -0.5 ], 3 ] } } } )
```

Creating a 2dsphere Index

You can index one or more 2dsphere fields in an index.

```
db.collection.ensureIndex( { <location field> : "2dsphere" } )
```

The geoJSON Specification

- The geoJSON format encodes location data on the earth.
- The spec is at <http://geojson.org/geojson-spec.html>
- This spec is incorporated in MongoDB 2dsphere indexes.
- Includes Point, LineString, Polygon, and combinations of these.

geoJSON Considerations

- The coordinates of points are given in degrees (latitude then longitude)
- The LineString that joins two points will always be a geodesic.
- Short lines (around a few hundred kilometers or less) will go about where you would expect them to.
- Polygons are made of a closed set of LineStrings.

Note:

- A geodesic may not go where you think.
 - E.g., the LineString that joins the points [90, 5] and [-90, 5]:
 - Does NOT go through the point [0, 5]
 - DOES go through the point [0, 90] (i.e., the North Pole).
-

Simple Types of 2dsphere Objects

Point: A single point on the globe

```
{ <field_name> : { type : "Point",  
  coordinates : [ <longitude>, <latitude> ] } }
```

LineString: A geodesic line that is defined by its two end Points

```
{ <field_name> : { type : "LineString",  
  coordinates : [ [ <longitude 1>, <latitude 1> ],  
    [ <longitude 2>, <latitude 2> ],  
    ...,  
    [ <longitude n>, <latitude n> ] ] } }
```

Note:

- Legacy coordinate pairs are treated as Points by a 2dsphere index.
-

Polygons

Simple Polygon:

```
{ <field_name> : { type : "Polygon",  
  coordinates : [ [ [ <Point1 coordinate pair> ],  
    [ <Point2 coordinate pair> ],  
    ...,  
    [ <Point1 coordinate pair again> ] ] ]  
  } }
```

Polygon with One Hole:

```
{ <field_name> : { type : "Polygon",  
  coordinates : [ [ <Points that define outer polygon> ],  
    [ <Points that define inner polygon> ] ]  
  } }
```

Other Types of 2dsphere Objects

- **MultiPoint:** One or more Points in one document
- **MultiLine:** One or more LineStrings in one document
- **MultiPolygon:** One or more Polygons in one document
- **GeometryCollection:** One or more geoJSON objects in one document

Exercise: Inserting geoJSON Objects (1)

Create a coordinate pair for each the following airports. Create one variable per airport.

- LaGuardia (New York): 40.7772° N, 73.8726° W
- JFK (New York): 40.6397° N, 73.7789° W
- Newark (New York): 40.6925° N, 74.1686° W
- Heathrow (London): 52.4775° N, 0.4614° W
- Gatwick (London): 51.1481° N, 0.1903° W
- Stansted (London): 51.8850° N, 0.2350° E
- Luton (London): 51.9000° N, 0.4333° W

Note:

```
laguardia = [ -73.8726, 40.7772 ]
jfk = [ -73.7789, 40.6397 ],
newark = [ -74.1686, 40.6925 ]
heathrow = [ -0.4614, 52.4775 ]
gatwick = [ -0.1903, 51.1481 ]
stansted = [ 0.2350, 51.8850 ]
luton = [-0.4333, 51.9000 ]
```

- Remember, we use [latitude, longitude].
 - In this example, we have made North (latitude) and East (longitude) positive.
 - West and South are negative.
-

Exercise: Inserting geoJSON Objects (2)

- Now let's make arrays of these.
- Put all the New York area airports into an array called `nyPorts`.
- Put all the London area airports into an array called `londonPorts`.
- Create a third array for flight numbers: "AA4453", "VA3333", "UA2440"

Note:

```
nyPorts = [ laguardia, jfk, newark ]
londonPorts = [ heathrow, gatwick, stansted, luton ]
flightNumbers = [ "AA4453", "VA3333", "UA2440" ]
```

Exercise: Inserting geoJSON Objects (3)

- Create documents for every possible New York to London flight.
- Include a `flightNumber` field for each flight.

Note:

```
for (takeoff in ny_ports) {
  for (landing in london_ports) {
    db.flights.insert(
      { origin : { type : "Point",
                    coordinates : ny_ports[takeoff] },
        destination : { type : "Point",
                        coordinates : london_ports[landing] },
        flightNumber : flightNumbers[takeoff] } )
  }
}
```

Exercise: Creating a 2dsphere Index

- Create two indexes on the collection `flights`.
- Make the first a compound index on the fields:
 - `origin`
 - `destination`
 - `flightNumber`
- Specify 2dsphere indexes on both `origin` and `destination`.
- Specify a simple index on `name`.
- Make the second index just a 2dsphere index on `destination`.

Note:

```
db.flights.ensureIndex( { origin : "2dsphere",
                          destination : "2dsphere",
                          flightNumber : 1 } )

db.flights.ensureIndex( { destination : "2dsphere" } )

db.flights.getIndexes() // see the indexes.
```

Querying 2dsphere Objects

`$geoNear`: Finds all points, orders them by distance from a position.

```
{ <field name> : { $near : { $geometry : {  
    type : "Point",  
    coordinates : [ lng, lat ] },  
    $maxDistance : <meters> } } } }
```

`$near`: Just like `$geoNear`, except in very edge cases; check the docs.

`$geoWithin`: Only returns documents with a location completely contained within the query.

`$geoIntersects`: Returns documents with their indexed field intersecting any part of the shape in the query.

3.6 TTL Indexes

Learning Objectives

Upon completing this module students should understand:

- How to create a TTL index.
- When a TTL indexed document will get deleted.
- Limitations of TTL indexes.

TTL Index Basics

- TTL is short for “Time To Live”.
- This must index a field of type “Date” (including `ISODate`) or “Timestamp”
- Any Date field older than `expireAfterSeconds` will get deleted at some point

Creating a TTL Index

Create with:

```
db.collection.ensureIndex( { field_name : 1 },  
    { expireAfterSeconds : some_number } )
```


Exercise: Creating a TTL Index

Let's create a TTL index on the `ttl` collection that will delete documents older than 30 seconds. Write a script that will insert documents at a rate of one per second.

```
db.testcol.drop()
db.testcol.ensureIndex( { a : 1 }, { expireAfterSeconds : 30 } )

i = 0
while (true) {
  i += 1;
  db.testcol.insert( { a : ISODate(), b : i } );
  sleep(1000); // Sleep for 1 second
}
```

Exercise: Check the Collection

Then, leaving that window open, open up a new terminal and connect to the database with the mongo shell. This will allow us to verify the TTL behavior.

```
// look at the output and wait. After a ramp-up of up to a minute or so,
// count() will be reset to 30 once/minute.
while (true) {
  print(db.testcol.count());
  sleep(100);
}
```

3.7 Text Indexes

Learning Objectives

Upon completing this module, students should understand:

- The purpose of a text index
- How to create text indexes
- How to search using text indexes
- How to rank search results by relevance score

What is a Text Index?

- A text index is based on the tokens (words, etc.) used in string fields.
- MongoDB supports text search for a number of languages.
- Text indexes drop language-specific stop words (e.g. in English “the”, “an”, “a”, “and”, etc.)
- Text indexes use simple, language-specific suffix stemming (e.g., “running” to “run”)

Creating a Text Index

You create a text index a little bit differently than you create a standard index.

```
db.collection.ensureIndex( { <field name> : "text" } )
```

Exercise: Creating a Text Index

Create a text index on the “dialog” field of the montyPython collection.

```
db.montyPython.ensureIndex( { dialog : "text" } )
```

Text Indexes are Similar to Multikey Indexes

- Continuing our example, you can treat the dialog field as a multikey index.
- A multikey index with each of the words in dialog as values.
- You can query the field using the \$text operator.

Exercise: Inserting Texts

Let’s add some documents to our montyPython collection.

```
db.montyPython.insert( [
  { _id : 1,
    dialog : "What is the air-speed velocity of an unladen swallow?" },
  { _id : 2,
    dialog : "What do you mean? An African or a European swallow?" },
  { _id : 3,
    dialog : "Huh? I... I don't know that." },
  { _id : 45,
    dialog : "You're using coconuts!" },
  { _id : 55,
    dialog : "What? A swallow carrying a coconut?" } ] )
```

Querying a Text Index

Next, let's query the collection. The syntax is:

```
db.collection.find( { $text : { $search : "query terms go here" } } )
```

Exercise: Querying a Text Index

Using the text index, find all documents in the `montyPython` collection with the word “swallow” in it.

```
// Returns 3 documents.  
db.montyPython.find( { $text : { $search : "swallow" } } )
```

Exercise: Querying Using Two Words

- Find all documents in the `montyPython` collection with either the word ‘coconut’ or ‘swallow’.
- By default MongoDB ORs query terms together.
- E.g., if you query on two words, results include documents using either word.

```
// Finds 4 documents, 3 of which contain only one of the two words.  
db.montyPython.find( { $text : { $search : "coconut swallow" } } )
```

Search for a Phrase

- To match an exact phrase, include search terms in quotes (escaped).
- The following query selects documents containing the phrase “coffee cake”:

```
db.montyPython.find( { $text: { $search: "\"European swallow\"" } } )
```

Text Search Score

- The search algorithm assigns a relevance score to each search result.
- The score is generated by a vector ranking algorithm.
- The documents can be sorted by that score.

```
db.collection.find(  
  { $text : { $search : "swallow coconut" } },  
  { textScore: { $meta : "textScore" } }  
)  
.sort(  
  { textScore: { $meta: "textScore" } }  
)
```

4 Replica Sets

Introduction to Replica Sets (page 76) An introduction to replication and replica sets.

Elections in Replica Sets (page 80) The process of electing a new primary (automated failover) in replica sets.

Replica Set Roles and Configuration (page 86) Configuring replica set members for common use cases.

The Oplog: Statement Based Replication (page 88) The process of replicating data from one node of a replica set to another.

Write Concern (page 93) Balancing performance and durability of writes.

Read Preference (page 97) Configuring clients to read from specific members of a replica set.

Exercise: Setting up a Replica Set (page 98) Launching members, configuring, and initiating a replica set.

4.1 Introduction to Replica Sets

Learning Objectives

Upon completing this module, students should understand:

- Striking the right balance between cost and redundancy
- The many scenarios replication addresses and why
- How to avoid downtime and data loss using replication

Use Cases for Replication

- High Availability
- Disaster Recovery
- Functional Segregation

High Availability (HA)

- Data still available following:
 - Equipment failure (e.g. server, network switch)
 - Datacenter failure
- This is achieved through automatic failover.

Note: If we lose a server and MongoDB is correctly configured:

- Our database system can still service reads and writes, but by default not during failover period when the election takes place and there is no primary.
 - Without manual intervention as long as there is still a majority of nodes available.
-

Disaster Recovery (DR)

- We can duplicate data across:
 - Multiple database servers
 - Storage backends
 - Datacenters
- Can restore data from another node following:
 - Hardware failure
 - Service interruption

Functional Segregation

There are opportunities to exploit the topology of a replica set.

- Based on physical location (e.g. rack or datacenter location)
- For analytics, reporting, data discovery, system tasks, etc.
- For backups

Note:

- You may direct particular queries to specific nodes (which may have different indexes or hardware) to increase overall performance.
 - Backup data from secondaries to avoid performance penalties on the primary, especially when using tools like `mongodump` which are I/O intensive and evict the working set from memory (significant when data size is larger than RAM and disks are slow).
 - Dedicate secondaries for other purposes such as analytics jobs.
-

Replication is Not Designed for Scaling


- Can be used for scaling reads, but generally not recommended.
- Drawbacks include:
 - Eventual consistency
 - Not scaling writes
 - Potential system overload when secondaries are unavailable
- Consider sharding for scaling reads and writes.

Note:

- Eventual consistency: This is often tricky to handle as most developers are used to the strong consistency that most databases (and MongoDB) have by default. It also raises the question of how stale the data can be, and what to do when it crosses the threshold (e.g. fall back to reading from the primary).
- Potential system overload: For example, consider a 3 data node replica set using secondaries to scale reads. Each node is serving reads at 70% of its capacity, which is a reasonable utilization rate. What happens if one of

the secondaries fail or is intentionally taken down for maintenance (e.g. upgrading the OS or MongoDB)? Even if the load splits evenly between the 2 remaining nodes, they will be at $70 + (70/2) = 105\%$ capacity.

Replica Sets



Note:

- MongoDB implements replication in the form of replica sets. Don't use the term master-slave as that is what we had before replica sets. It still exists for some corner cases (e.g. > 12 replicas) but should otherwise be avoided.
 - A replica set consists of one or more `mongod` servers. Maximum 12 nodes in total and up to 7 with votes.
 - There is at most one `mongod` that is "primary" at any one time (though there are edge cases/bugs when there is more than one).
 - There are usually two or more other `mongod` instances that are secondaries.
 - Secondaries may become primary if there is a failover event of some kind.
 - Failover is automatic when correctly configured and a majority of nodes remain.
 - The secondaries elect a new primary automatically. A primary may also voluntarily step down, like when it can no longer reach the majority of nodes to avoid a potential split brain scenario.
-

Primary Server


- Clients send writes the primary only.
- MongoDB, Inc. maintains client drivers in many programming languages like Java, C#, Python, Ruby, and PHP.
- MongoDB drivers are replica set aware.

Note: If the primary for a replica set changes from one node to another, the driver will automatically route writes to the correct `mongod`.

Secondaries

- A secondary replicates operations from another node in the replica set.
- Secondaries usually replicate from the primary.
- Secondaries may also replicate from other secondaries. This is called replication chaining.
- A secondary may become primary as a result of a failover scenario.

Heartbeats



Note:

- The members of a replica set use heartbeats to determine if they can reach every other node.
 - The heartbeats are sent every two seconds.
 - If a node is unreachable, this may indicate server failure, a network partition, or simply too slow to respond. The heartbeat will timeout and retried several times before the state is updated.
-

The Oplog

- The operations log, or oplog, is a special capped collection that is the basis for replication.
- The oplog maintains one entry for each document affected by every write operation.
- Secondaries copy operations from the oplog of their sync source.


4.2 Elections in Replica Sets

Learning Objectives

Upon completing this module students should understand:

- That elections enable automated failover in replica sets
- How votes are distributed to members
- What prompts an election
- How a new primary is selected

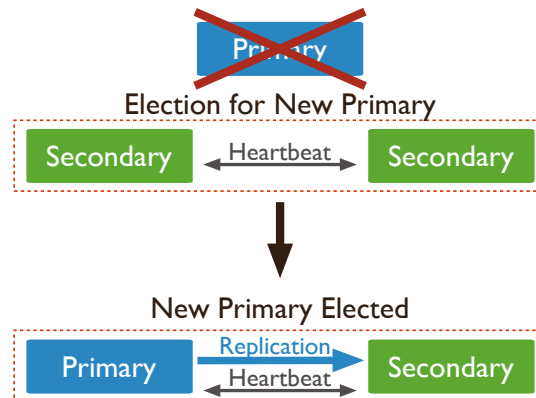
Members and Votes



Note:

- In order for writes to occur, one member of a replica set must be primary.
 - In the event the current primary becomes unavailable, the remaining members elect a new primary.
 - Voting members of replica set each get one vote.
 - Up to seven members may be voting members.
 - This enables MongoDB to ensure elections happen quickly, but enables distribution of votes to different data centers.
 - In order to be elected primary a server must have a true majority of votes.
 - A member must have greater than 50% of the votes in order to be elected primary.
-

Calling Elections



Note:

- MongoDB uses a consensus protocol to determine when an election is required.
 - Essentially, an election will occur if there is no primary.
 - Upon initiation of a new replica set the members will elect a primary.
 - If a primary steps down the set will hold an election.
 - A secondary will call for an election if it does not receive a response to a heartbeat sent to the primary after waiting for 10 seconds.
 - If other members agree that the primary is not available, an election will be held.
-

Selecting a New Primary

Three factors are important in the selection of a primary:

- Priority
- Optime
- Connections

Priority

- The higher its priority, the more likely a member is to become primary.
- The default is 1.
- Servers with a priority of 0 will never become primary.
- Priority values are floating point numbers 0 - 1000 inclusive.

Note:

- Priority is a configuration parameter for replica set members.
 - Use priority to determine where writes will be directed by default.
 - And where writes will be directed in case of failover.
-

- Generally all identical nodes in a datacenter should have the same priority to avoid unnecessary failovers. For example, when a higher priority node rejoins the replica set after a maintenance or failure event, it will trigger a failover (during which by default there will be no reads and writes) even though it is unnecessary.
 - More on this in a later module.
-

Optime

- Optime: Operation time, which is the timestamp of the last operation the member applied from the oplog.
- To be elected primary, a member must have the most recent optime.
- Only optimes of visible members are compared.

Connections

- Must be able to connect to a majority of the members in the replica set.
- Majority refers to the total number of votes.
- Not the total number of members.

Note: To be elected primary, a replica set member must be able to connect to a majority of the members in the replica set.

When will a primary step down?


- After receiving the `replSetStepDown` or `rs.stepDown()` command.
- If a secondary is eligible for election and has a higher priority.
- If it cannot contact a majority of the members of the replica set.

Exercise: Elections in Failover Scenarios

- We have learned about electing a primary in replica sets
- Let's look at some scenarios in which failover might be necessary.

Scenario A: 3 Data Nodes in 1 DC

Which secondary will become the new primary?




Note:

- It depends on the priorities of the secondaries.
 - And on the optime.
-

Scenario B: 3 Data Nodes in 2 DCs

Which member will become primary following this type of network partition?




Note:

- The current primary is likely to remain primary.
 - It probably has the highest priority.
 - If DC2 fails, we still have a primary.
 - If DC1 fails, we won't have a primary automatically. The remaining node in DC2 needs to be manually promoted by reconfiguring the replica set.
-

Scenario C: 4 Data Nodes in 2 DCs

What happens following this network partition?




Note:

- We enter a state with no primary.
 - Each side of the network partition has only 2 votes (not a majority).
 - All the servers assume secondary status.
 - This is avoidable.
 - One solution is to add another member to the replica set.
 - If another data node can not be provisioned, MongoDB has a special alternative called an arbiter that requires minimal resources.
 - An arbiter is a `mongod` instance without data and performs only heartbeats, votes, and vetoes.
-

Scenario D: 5 Nodes in 2 DCs

The following is similar to Scenario C, but with the addition of an arbiter in Data Center 1. What happens here?




Note:

- The current primary is likely to remain primary.
 - The arbiter helps ensure that the primary can reach a majority of the replica set.
-

Scenario E: 3 Data Nodes in 3 DCs

- What happens here if any one of the nodes/DCs fail?
- What about recovery time?




Note:

- The intent is to explain the advantage of deploying to 3 DCs - it's the minimum number of DCs in order for MongoDB to automatically failover if any one DC fails. This is generally what we recommend to customers in our consult and health check reports, though many continue to use 2 DCs due to costs and legacy reasons.
 - To have automated failover in the event of single DC level failure, there must be at least 3 DCs. Otherwise the DC with the minority of nodes must be manually reconfigured.
 - One of the data nodes can be replaced by an arbiter to reduce costs.
-

Scenario F: 5 data nodes in 3 DCs

What happens here if any one of the nodes/DCs fail? What about recovery time?



Note:

- Adds another data node to each “main” DC to reduce typically slow and costly cross DC network traffic if an initial sync or similar recovery is needed, as the recovering node can pull from a local replica instead.
 - Depending on the data sizes, operational budget, and requirements, this can be overkill.
 - The data node in DC3 can be replaced by an arbiter to reduce costs.
-

4.3 Replica Set Roles and Configuration

Learning Objectives

Upon completing this module students should understand:

- The use of priority to preference certain members or datacenters as primaries.
- Hidden members.
- The use of hidden secondaries for data analytics and other purposes (when secondary reads are used).
- The use of slaveDelay to protect against operator error.

Example: A Five-Member Replica Set Configuration

- For this example application, there are two datacenters.
- We name the hosts accordingly: dc1-1, dc1-2, dc2-1, etc.
 - This is just a clarifying convention for this example.
 - MongoDB does not care about host names except to establish connections.
- The nodes in this replica set have a variety of roles in this application.

Configuration

```
conf = {                                // 5 data-bearing nodes
  _id: "mySet",
  members: [
    { _id : 0, host : "dc1-1.example.net:27017", priority : 5 },
    { _id : 1, host : "dc1-2.example.net:27017", priority : 5 },
    { _id : 2, host : "dc2-1.example.net:27017" },
    { _id : 3, host : "dc1-3.example.net:27017", hidden : true },
    { _id : 4, host : "dc2-2.example.net:27017", hidden : true,
      slaveDelay: 7200 }
  ]
}
```

Principal Data Center

```
{ _id : 0, host : "dc1-1.example.net", priority : 5 },
{ _id : 1, host : "dc1-2.example.net", priority : 5 },
```

Note:

- The objective with the priority settings for these two nodes is to prefer to DC1 for writes.
 - The highest priority member that is up to date will be elected primary.
 - Up to date means the member's copy of the oplog is within 10 seconds of the primary.
 - If a member with higher priority than the primary is a secondary because it is not up to date, but eventually catches up, it will force an election and win.
-

Data Center 2

```
{ _id : 2, host : "dc2-1.example.net:27017" },
```

Note:

- Priority is not specified, so it is at the default of 1.
 - dc2-1 could become primary, but only if both dc1-1 and dc1-2 are down.
 - If there is a network partition and clients can only reach DC2, we can manually failover to dc2-1.
-

What about dc1-3 and dc2-2?

```
// Both are hidden.  
// Clients will not distribute reads to hidden members.  
// We use hidden members for dedicated tasks.  
{ _id : 3, host : "dc1-3.example.net:27017", hidden : true },  
{ _id : 4, host : "dc2-2.example.net:27017", hidden : true,  
  slaveDelay: 7200 }
```

Note:

- Will replicate writes normally.
 - We would use this node to pull reports, run analytics, etc.
 - We can do so without paying a performance penalty in the application for either reads or writes.
-

What about dc2-2?

```
{ _id : 4, host : "dc2-2.example.net:27017", hidden : true,  
  slaveDelay : 7200 }
```

Note:

- slaveDelay permits us to specify a time delay (in seconds) for replication.
 - In this case it is 7200 seconds or 2 hours.
 - slaveDelay allows us to use a node as a short term protection against operator error:
 - Fat fingering – for example, accidentally dropping a collection in production.
 - Other examples include bugs in an application that result in corrupted data.
 - Not recommended. Use proper backups instead as there is no optimal delay value. E.g. 2 hours might be too long or too short depending on the situation.
-

4.4 The Oplog: Statement Based Replication

Learning Objectives

Upon completing this module students should understand:

- Binary vs. statement-based replication.
- How the oplog is used to support replication.
- How operations in MongoDB are translated into operations written to the oplog.
- Why oplog operations are idempotent.
- That the oplog is a capped collection and the implications this holds for syncing members.

Binary Replication

- MongoDB replication is statement based.
- Contrast that with binary replication.
- With binary replication we would keep track of:
 - The data files
 - The offsets
 - How many bytes were written for each change
- In short, we would keep track of actual bytes and very specific locations.
- We would simply replicate these changes across secondaries.

Tradeoffs

- The good thing is that figuring out where to write, etc. is very efficient.
- But we must have a byte-for-byte match of our data files on the primary and secondaries.
- The problem is that this couples our replica set members in ways that are inflexible.
- Binary replication may also replicate disk corruption.

Note:

- Some deployments might need to run different versions of MongoDB on different nodes.
 - Different versions of MongoDB might write to different file offsets.
 - We might need to run a compaction or repair on a secondary.
 - In many cases we want to do these types of maintenance tasks independently of other nodes.
-

Statement-Based Replication

- Statement-based replication facilitates greater independence among members of a replica set.
- MongoDB stores a statement for every operation in a capped collection called the `oplog`.
- Secondaries do not simply apply exactly the operation that was issued on the primary.

Example

Suppose the following remove is issued and it deletes 100 documents:

```
db.foo.remove({ age : 30 })
```

This will be represented in the `oplog` with records such as the following:

```
{ "ts" : Timestamp(1407159845, 5), "h" : NumberLong("-704612487691926908"),  
  "v" : 2, "op" : "d", "ns" : "bar.foo", "b" : true, "o" : { "_id" : 65 } }  
{ "ts" : Timestamp(1407159845, 1), "h" : NumberLong("6014126345225019794"),  
  "v" : 2, "op" : "d", "ns" : "bar.foo", "b" : true, "o" : { "_id" : 333 } }  
{ "ts" : Timestamp(1407159845, 4), "h" : NumberLong("8178791764238465439"),  
  "v" : 2, "op" : "d", "ns" : "bar.foo", "b" : true, "o" : { "_id" : 447 } }  
{ "ts" : Timestamp(1407159845, 3), "h" : NumberLong("-1707391001705528381"),  
  "v" : 2, "op" : "d", "ns" : "bar.foo", "b" : true, "o" : { "_id" : 1033 } }  
{ "ts" : Timestamp(1407159845, 2), "h" : NumberLong("-6814297392442406598"),  
  "v" : 2, "op" : "d", "ns" : "bar.foo", "b" : true, "o" : { "_id" : 9971 } }
```

Replication Based on the Oplog

- One statement per document affected by each write: insert, update, or delete.
- Provides a level of abstraction that enables independence among the members of a replica set:
 - With regard to MongoDB version.
 - In terms of how data is stored on disk.
 - Freedom to do maintenance without the need to bring the entire set down.

Note:

- Can do maintenance without bringing the set down because statement-based replication does not depend on all nodes running the same version of MongoDB or other restrictions that may be imposed by binary replication.
 - In the next exercise, we will see that the `oplog` is designed so that each statement is idempotent.
 - This feature has several benefits for independent operation of nodes in replica sets.
-

Create a Replica Set

Let's take a look at a concrete example. Launch mongo shell as follows.

```
mongo --nodb
```

Create a replica set by running the following command in the mongo shell.

```
replicaSet = new ReplSetTest( { nodes : 3 } )
```

ReplSetTest

- ReplSetTest is useful for experimenting with replica sets as a means of hands-on learning.
- It should never be used in production. Never.
- The command above will create a replica set with three members.
- It does not start the mongods, however.
- You will need to issue additional commands to do that.

Start the Replica Set

Start the mongod processes for this replica set.

```
replicaSet.startSet()
```

Issue the following command to configure replication for these mongods. You will need to issue this while output is flying by in the shell.

```
replicaSet.initiate()
```

Status Check

- You should now have three mongods running on ports 31000, 31001, and 31002.
- You will see log statements from all three printing in the current shell.
- To complete the rest of the exercise, open a new shell.

Connect to the Primary

Open a new shell, connecting to the primary.

```
mongo --port 31000
```

Create some Inventory Data

Use the `store` database:

```
use store
```

Add the following inventory:

```
inventory = [ { _id: 1, inStock: 10 }, { _id: 2, inStock: 20 },
              { _id: 3, inStock: 30 }, { _id: 4, inStock: 40 },
              { _id: 5, inStock: 50 }, { _id: 6, inStock: 60 } ]
db.products.insert(inventory)
```

Perform an Update

Issue the following update. We might issue this update after a purchase of three items.

```
db.products.update({ _id: { $in: [ 2, 5 ] } },
                  { $inc: { inStock : -1 } },
                  { multi: true })
```

View the Oplog

The oplog is a capped collection in the `local` database of each replica set member:

```
use local
db.oplog.rs.find()
{ "ts" : Timestamp(1406944987, 1), "h" : NumberLong(0), "v" : 2, "op" : "n", "ns" : "",
  "o" : { "msg" : "initiating set" } }
...
{ "ts" : Timestamp(1406945076, 1), "h" : NumberLong("-9144645443320713428"), "v" : 2,
  "op" : "u", "ns" : "store.products", "o2" : { "_id" : 2 }, "o" : { "$set" : { "inStock" : 19 } } }
{ "ts" : Timestamp(1406945076, 2), "h" : NumberLong("-7873096834441143322"), "v" : 2,
  "op" : "u", "ns" : "store.products", "o2" : { "_id" : 5 }, "o" : { "$set" : { "inStock" : 49 } } }
```

Note:

- Note the last two entries in the oplog.
 - These entries reflect the update command issued above.
 - Note that there is one operation per document affected.
 - More specifically, one operation for each of the documents with the `_id` values 2 and 5.
-

Operations in the Oplog are Idempotent

- Each operation in the oplog is idempotent.
- Whether applied once or multiple times it produces the same result.
- Necessary if you want to be able to copy data while simultaneously accepting writes.

Note: We need to be able to copy while accepting writes when:

- Doing an initial sync for a new replica set member.
 - When a member rejoins a replica set after a network partition a member might end up writing operations it had already received prior to the partition.
-

The Oplog Window

- Oplogs are capped collections.
- Capped collections are fixed-size.
- They guarantee preservation of insertion order.
- They support high-throughput operations.
- Like circular buffers, once a collection fills its allocated space:
 - It makes room for new documents.
 - By overwriting the oldest documents in the collection.

Sizing the Oplog

- The oplog should be sized to account for latency among members.
- The default size oplog is usually sufficient.
- But you want to make sure that your oplog is large enough:
 - So that the oplog window is large enough to support replication
 - To give you a large enough history for any diagnostics you might wish to run.

4.5 Write Concern

Learning Objectives

Upon completing this module students should understand:

- How and when rollback occurs in MongoDB.
- The tradeoffs between durability and performance.
- Write concern as a means of ensuring durability in MongoDB.
- The different levels of write concern.

What happens to the write?

- A write is sent to a primary.
- The primary acknowledges the write to the client.
- And then the primary becomes unavailable before a secondary can replicate the write

Answer

- Another member might be elected primary.
- It will not have the last write that occurred before the previous primary became unavailable.
- When the previous primary becomes available again:
 - It will note it has writes that were not replicated.
 - It will put these writes into a `rollback file`.
 - A human will need to determine what to do with this data.
- This is default behavior in MongoDB and can be controlled using `write concern`.

Balancing Durability with Performance

- The previous scenario is a specific instance of a common distributed systems problem.
- For some applications it might be acceptable for writes to be rolled back.
- Other applications may have varying requirements with regard to durability.
- Tunable write concern:
 - Make critical operations persist to an entire MongoDB deployment.
 - Specify replication to fewer nodes for less important operations.

Note:

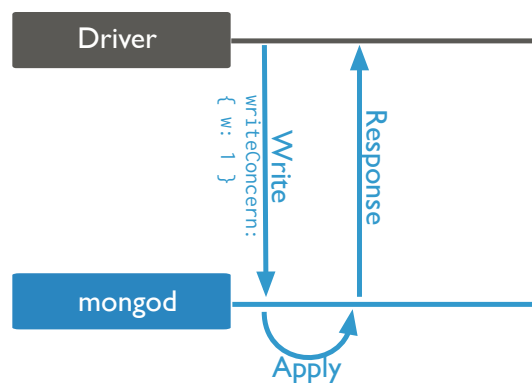
- MongoDB provides tunable write concern to better address the specific needs of applications.
- Clients may adjust write concern to ensure that the most important operations persist successfully to an entire MongoDB deployment.

- For other less critical operations, clients can adjust write concern to ensure faster performance.
-

Defining Write Concern

- Clients may define the write concern per write operation, if necessary.
- Standardize on specific levels of write concerns for different classes of writes.
- In the discussion that follows we will look at increasingly strict levels of write concern.

Write Concern: { w : 1 }



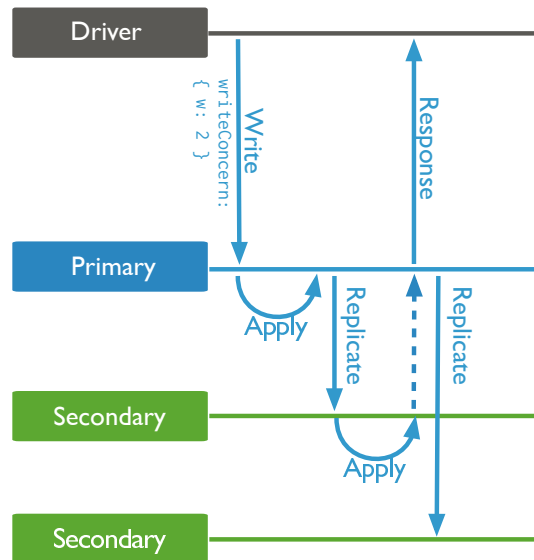
Note:

- We refer to this write concern as “Acknowledged”.
 - This is the default.
 - The primary sends an acknowledgement back to the client that it received the write operation (in RAM).
 - Allows clients to catch network, duplicate key, and other write errors.
-

Example: { w : 1 }

```
db.edges.insert({ from : "tom185", to : "mary_p" }, { w : 1 })
```

Write Concern: { w : 2 }



Note:

- Called “Replica Acknowledged”
 - Ensures the primary completed the write.
 - Ensures at least one secondary replicated the write.
-

Example: { w : 2 }

```
db.customer.update({ user : "mary_p" },
  { $push : { shoppingCart:
    { _id : 335443, name : "Brew-a-cup",
      price : 45.79 } } },
  { w : 2 })
```

Other Write Concerns

- You may specify any integer as the value of the w field for write concern.
- This guarantees that write operations have propagated to the specified number of members.
- E.g., { w : 3 }, { w : 4 }, etc.

Write Concern: { w : "majority" }

- Ensures the primary completed the write (in RAM).
- Ensures write operations have propagated to a majority of a replica set's members.
- Avoids hard coding assumptions about the size of your replica set into your application.
- Using majority trades off performance for durability.
- It is suitable for critical writes and to avoid rollbacks.

Example: { w : "majority" }

```
db.products.update({ _id : 335443 },
  { $inc : { inStock : -1 } },
  { w : "majority" })
```

Quiz: Which write concern?

Suppose you have a replica set with 7 data nodes. Your application has critical inserts for which you do not want rollbacks to happen. Secondaries may be taken down from to time for maintenance, leaving you with a potential 4 server replica set. Which write concern is best suited for these critical inserts?

- { w : 1 }
- { w : 2 }
- { w : 3 }
- { w : 4 }
- { w : "majority" }

Note: Answer: { w : "majority" }. This is the same as 4 for a 7 member replica set.

Further Reading

See [Write Concern Reference](http://docs.mongodb.org/manual/reference/write-concern)⁶ for more details on write concern configurations, including setting timeouts and identifying specific replica set members that must acknowledge writes (i.e. [tag sets](http://docs.mongodb.org/manual/tutorial/configure-replica-set-tag-sets/#replica-set-configuration-tag-sets)⁷).

⁶<http://docs.mongodb.org/manual/reference/write-concern>

⁷<http://docs.mongodb.org/manual/tutorial/configure-replica-set-tag-sets/#replica-set-configuration-tag-sets>

4.6 Read Preference

What is Read Preference?

- Read preference allows you to specify the nodes in a replica set to read from.
- Clients only read from the primary by default.
- There are some situations in which a client may want to read from:
 - Any secondary
 - A specific secondary
 - A specific type of secondary
- Only read from a secondary if you can tolerate possibly stale data, as not all writes might have replicated.

Use Cases

- Running systems operations without affecting the front-end application.
- Providing local reads for geographically distributed applications.
- Maintaining availability during a failover.

Note:

- If you have application servers in multiple data centers, you may consider having a [geographically distributed replica set](#)⁸ and using a read preference of `nearest`.
 - This allows the client to read from the lowest-latency members.
 - Use `primaryPreferred` if you want an application to read from the primary under normal circumstances, but to allow possibly stale reads from secondaries during failover.
-

Not for Scaling

- In general, do *not* read from secondaries to provide extra capacity for reads.
- [Sharding](#)⁹ increases read and write capacity by distributing operations across a group of machines.
- Sharding is a better strategy for adding capacity.

⁸<http://docs.mongodb.org/manual/core/replica-set-geographical-distribution>

⁹<http://docs.mongodb.org/manual/sharding>

Read Preference Modes

MongoDB drivers support the following read preferences. Note that `hidden` nodes will never be read from when connected via the replica set.

- **primary**: Default. All operations read from the primary.
- **primaryPreferred**: Read from the primary but if it is unavailable, read from secondary members.
- **secondary**: All operations read from the secondary members of the replica set.
- **secondaryPreferred**: Read from secondary members but if no secondaries are available, read from the primary.
- **nearest**: Read from member of the replica set with the least network latency, regardless of the member's type.

Tag Sets

- There is also the option to use tag sets.
- You may tag nodes such that queries that contain the tag will be routed to one of the servers with that tag.
- This can be useful for running reports, say for a particular data center or nodes with different hardware (e.g. hard disks vs SSDs).

For example, in the mongo shell:

```
conf = rs.conf()
conf.members[0].tags = { dc : "east", use : "production" }
conf.members[1].tags = { dc : "east", use : "reporting" }
conf.members[2].tags = { use : "production" }
rs.reconfig(conf)
```

4.7 Exercise: Setting up a Replica Set

Overview

- In this exercise we will setup a 3 data node replica set on a single machine.
- In production, each node should be run on a dedicated host:
 - To avoid any potential resource contention
 - To provide isolation against server failure.

Create Data Directories

Since we will be running all nodes on a single machine, make sure each has its own data directory.

On Linux or Mac OS, run the following in the terminal to create the 3 directories `~/data/rs1`, `~/data/rs2`, and `~/data/rs3`:

```
mkdir -p ~/data/rs{1,2,3}
```

On Windows, run the following command instead in Command Prompt or PowerShell:

```
md c:\data\rs1 c:\data\rs2 c:\data\rs3
```

Launch Each Member

Now start 3 instances of `mongod` in the foreground so that it is easier to observe and shutdown.

On Linux or Mac OS, run *each* of the following commands in its *own terminal window*:

```
mongod --replSet myReplSet --dbpath ~/data/rs1 --port 27017 --oplogSize 200 --smallfiles
mongod --replSet myReplSet --dbpath ~/data/rs2 --port 27018 --oplogSize 200 --smallfiles
mongod --replSet myReplSet --dbpath ~/data/rs3 --port 27019 --oplogSize 200 --smallfiles
```

On Windows, run *each* of the following commands in its *own Command Prompt or PowerShell window*:

```
mongod --replSet myReplSet --dbpath c:\data\rs1 --port 27017 --oplogSize 200 --smallfiles
mongod --replSet myReplSet --dbpath c:\data\rs2 --port 27018 --oplogSize 200 --smallfiles
mongod --replSet myReplSet --dbpath c:\data\rs3 --port 27019 --oplogSize 200 --smallfiles
```

Status

- At this point, we have 3 `mongod` instances running.
- They were all launched with the same `replSet` parameter of “myReplSet”.
- Despite this, the members are not aware of each other yet.
- This is fine for now.

Note:

- In production, each member would run on a different machine and use service scripts. For example on Linux, modify `/etc/mongod.conf` accordingly and run:

```
sudo service mongod start
```
 - To simplify this exercise, we run all members on a single machine.
 - The same configuration process is used for this deployment as for one that is distributed across multiple machines.
-

Connect to a MongoDB Instance

- Connect to the one of the MongoDB instances with the mongo shell.
- To do so run the following command in the terminal, Command Prompt, or PowerShell:

```
mongo --port 27017
```

Configure the Replica Set

Note the port number of the primary from the output of `rs.status()`, which we'll need for the next step.

```
var config = {
  _id: "mySet",
  members: [
    { _id: 0, host: "localhost:27017" },
    { _id: 1, host: "localhost:27018" },
    { _id: 2, host: "localhost:27019" }
  ]
}
rs.initiate(config)

# Keep running rs.status() until there's a primary and 2 secondaries
rs.status()

exit      # or Ctrl-d
```

Write to the Primary

Connect to the primary with mongo shell by issuing a command such as the following:

```
mongo --port <PRIMARY_PORT>    # e.g. 27017
```

Now insert a simple test document via mongo shell. Once the insert succeeds, exit the mongo shell.

```
db.testcol.insert({ a: 1 })
db.testcol.count()
```

Read from a Secondary

Connect to one of the secondaries. E.g.:

```
mongo --port <NON_PRIMARY_PORT>    # e.g. 27018
```

Read from the secondary

```
rs.slaveOk()
db.testcol.find()
```

Review the Oplog

```
use local
db.oplog.rs.find()
```

Changing Replica Set Configuration

To change the replica set configuration, first connect to the primary via mongo shell:

```
mongo --port <PRIMARY_PORT> # e.g. 27017
```

Let's raise the priority of one of the secondaries. Assuming it is the 3rd node (e.g. on port 27019):

```
cfg = rs.conf()
cfg["members"][2]["priority"] = 10
rs.reconfig(cfg)
```

Note:

- Note that `cfg["members"][2]["priority"] = 10` does not actually change the priority.
 - `rs.reconfig(cfg)` does.
-

Verifying Configuration Change

You will see errors like the following, which are expected:

```
2014-10-07T17:01:34.610+0100 DBClientCursor::init call() failed
2014-10-07T17:01:34.613+0100 trying reconnect to 127.0.0.1:27017 (127.0.0.1) failed
2014-10-07T17:01:34.617+0100 reconnect 127.0.0.1:27017 (127.0.0.1) ok
reconnected to server after rs command (which is normal)
```

Verify that the replica set configuration is now as expected:

```
rs.conf()
```

The secondary will now become a primary. Check by running:

```
rs.status()
```

Further Reading

- [Replica Configuration](#)¹⁰
- [Replica States](#)¹¹

¹⁰<http://docs.mongodb.org/manual/reference/replica-configuration/>

¹¹<http://docs.mongodb.org/manual/reference/replica-states/>

5 Sharding

Introduction to Sharding (page 103) An introduction to sharding.

Balancing Shards (page 112) Chunks, the balancer, and their role in a sharded cluster.

Shard Tags (page 115) How tag-based sharding works.

Exercise: Setting Up a Sharded Cluster (page 117) Deploying a sharded cluster.

5.1 Introduction to Sharding

Learning Objectives

Upon completing this module, students should understand:

- What problems sharding solves
- When sharding is appropriate
- The importance of the shard key and how to choose a good one
- Why sharding increases the need for redundancy

Contrast with Replication

- In an earlier module, we discussed Replication.
- This should never be confused with sharding.
- Replication is about high availability and durability.
 - Taking your data and constantly copying it
 - Being ready to have another machine step in to field requests.


Sharding is Concerned with Scale

- What happens when a system is unable to handle the application load?
- It is time to consider scaling.
- There are 2 types of scaling we want to consider:
 - Vertical scaling
 - Horizontal scaling

Vertical Scaling

- Adding more RAM, faster disks, etc.
- When is this the solution?
- First, consider a concept called the *working set*.

The Working Set



Note:

- The working set contains the documents and indexes that are currently being used by an application.
 - It is best if the working set fits in RAM.
 - In this figure there are many more documents in use than can fit in RAM.
 - This would result in page faults affecting performance.
 - More RAM would cover a larger area of the database.
-


Limitations of Vertical Scaling

- There is a limit to how much RAM one machine can support.
- There are other bottlenecks such as I/O, disk access and network.
- Cost may limit our ability to scale up.
- There may be requirements to have a large working set that no single machine could possibly support.
- This is when it is time to scale horizontally.

Sharding Overview

- MongoDB enables you to scale horizontally through sharding.
- Sharding is about adding more capacity to your system.
- MongoDB's sharding solution is designed to perform well on commodity hardware.
- The details of sharding are abstracted away from applications.
- Queries are performed the same way as if sending operations to a single server.
- Connections work the same by default.


A Model that Does Not Scale



Note:

- Load and latency are related.
 - On a single server, latency tends to increase linearly with load.
 - Eventually latency is too great and the database has become a bottleneck.
-

A Scalable Model




Note:

- Latency increases, but not linearly
 - We're not saying there is no increase in latency.
 - Just that it is much smaller than the increase in the server load.
 - Relational databases can scale vertically (i.e., by buying a bigger box).
-

- MongoDB scales vertically and horizontally (i.e., spreading out the load across several boxes)
-

Sharding Basics




Note:

- When you shard a collection it is distributed across several servers.
 - When you perform a read or write operation it will go to a router that will then direct the query to the appropriate server.
 - Depending on the operation and how your cluster is configured you may need to touch only one server to complete the operation.
-

Sharded Cluster Architecture

Note:

- This figure illustrates one possible architecture for a sharded cluster.
 - Each shard is a self contained replica set.
 - Each replica set holds a partition of the data.
 - As many new shards could be added to this sharded cluster as scale requires.
 - At this point our cluster can handle a load approaching twice that of a single replica set using the same hardware.
 - As mentioned, read/write operations go through a router.
 - The server that routes requests is the mongos.
-




Mongos

- A mongos is responsible for accepting requests and returning results to an application driver.
- In a sharded cluster, nearly all operations go through a mongos.
- A sharded cluster can have as many mongos routers as required.
- It is typical for each application server to have one mongos.
- Always use more than one mongos to avoid a single point of failure.

Note:

- A mongos is typically deployed on an application server.
 - There should be one mongos per app server.
 - Scale with your app server.
 - Very little latency between the application and the router.
-

Config Servers



Note:

- The previous diagram was incomplete; it was missing config servers.
 - Use three config servers in production.
 - These hold only metadata about the sharded collections.
 - Where your mongos servers are
 - Any hosts that are not currently available
 - What collections you have
 - How your collections are partitioned across the cluster
 - Mongos processes use them to retrieve the state of the cluster.
 - You can access cluster metadata from a mongos by looking at the `config db`.
-

Config Server Hardware Requirements

- Quality network interfaces
- A small amount of disk space (typically a few GB)
- A small amount of RAM (typically a few GB)
- The larger the sharded cluster, the greater the config server hardware requirements.

When to Shard

- If you have more data than one machine can hold on its drives
- If your application is write heavy and you experiencing too much latency.
- If your working set outgrows the memory you can allocate to a single machine.

Possible Imbalance?

- Depending on how you configure sharding, data can become unbalanced on your sharded cluster.
 - Some shards might receive more inserts than others.
 - Some shards might have documents that grow more than those in other shards.
- This may result in too much load on a single shard.
 - Reads and writes
 - Disk activity
- This would defeat the purpose of sharding.

Balancing Shards

- MongoDB divides data into `chunks`.
- This is bookkeeping metadata.
 - There is nothing in a document that indicates its chunk.
 - The document does not need to be updated if its assigned chunk changes.
- If a chunk grows too large MongoDB will split it into two chunks.
- The MongoDB balancer keeps chunks distributed across shards in equal numbers.
- However, a balanced sharded cluster depends on a good shard key.


What is a Shard Key?

- You must define a shard key for a sharded collection.
- Based on one or more fields that every document must contain.
- Is immutable.
- The shard key determines where documents are located in the cluster.
- It is used to route operations to the appropriate shard.
- For reads and writes

Note:

- For reads and updates, the shard key determines which shard holds the document and will handle the operation.
 - When you insert a document, the shard key determines which server you will write to.
-

Targeted Query Using Shard Key



With a Good Shard Key

You might easily see that:

- Reads hit only 1 or 2 shards per query.
- Writes are distributed across all servers.
- Your disk usage is evenly distributed across shards.
- Things stay this way as you scale.

With a Bad Shard Key

You might see that:

- Your reads hit every shard.
- Your writes are concentrated on one shard.
- Most of your data is on just a few shards.
- Adding more shards to the cluster will not help.

Choosing a Shard Key

Generally, you want a shard key:

- That has high cardinality
- That is used in the majority of read queries
- For which the values read and write operations use are randomly distributed
- For which the majority of reads are routed to a particular server

More Specifically

- Your shard key should be consistent with your query patterns.
- If reads usually find only one document, you only need good cardinality.
- If reads retrieve many documents:
 - Your shard key supports locality
 - Matching documents will reside on the same shard.

Cardinality

- A good shard key will have high cardinality.
- A relatively small number of documents should have the same shard key.
- Otherwise operations become isolated to the same server.
- Because documents with the same shard key reside on the same shard.
- Adding more servers will not help.
- Hashing will not help.

Non-Monotonic

- A good shard key will generate new values non-monotonically.
- Datetimes, counters, and ObjectIds make bad shard keys.
- Monotonic shard keys cause all inserts to happen on the same shard.
- Hashing will solve this problem.
- However, doing range queries with a hashed shard key will perform a scatter-gather query across the cluster.

Note:

- Documents will eventually move as chunks are balanced.
 - But in the meantime one server gets hammered while others are idle.
 - And moving chunks has its own performance costs.
-

Shards Should be Replica Sets

- As the number of shards increases, the number of servers in your deployment increases.
- This increases the probability that one server will fail on any given day.
- With redundancy built into each shard you can mitigate this risk.

5.2 Balancing Shards

Learning Objectives

Upon completing this module students should understand:

- Chunks and the balancer
- The status of chunks in a newly sharded collection
- How chunk splits automatically occur
- Advantages of pre-splitting chunks
- How the balancer Works

Chunks and the Balancer


- Chunks are groups of documents.
- The shard key determines which chunk a document will be contained in.
- Chunks can be split when they grow too large.
- The balancer decides where chunks go.
- It handles migrations of chunks from one server to another.

Chunks in a Newly Sharded Collection

- The range of a chunk is defined by the shard key values of the documents the chunk contains.
- When a collection is sharded it starts with just one chunk.
- The first chunk for a collection will have the range:

```
{ $minKey : 1 } to { $maxKey : 1 }
```
- All shard key values from the smallest possible to the largest fall in this chunk's range

Chunk Splits



Note:

- When a chunk grows larger than the chunk size it will be split in half.
 - The default chunk size is 64MB.
 - A chunk can only be split between two values of a shard key.
 - If every document on a chunk has the same shard key value, it cannot be split.
 - This is why the shard key's cardinality is important
 - Chunk splitting is just a bookkeeping entry in the metadata.
 - No data bearing documents are altered.
-

Pre-Splitting Chunks

- You may pre-split data before loading data into a sharded cluster.
 - Pre-splitting is useful if:
 - You plan to do a large data import early on
 - You expect a heavy initial server load and want to ensure writes are distributed.
-

Note:

- A large data import will take time to split and balance without pre-splitting.
-

Start of a Balancing Round

- A balancing round may be initiated by any mongos in the cluster.
- This happens when the difference in the number of chunks between two shards becomes too large.
- Specifically, the difference between the shard with the most chunks and the shard with the fewest.
- A balancing round starts when the imbalance reaches:
 - 2 when the cluster has < 20 chunks
 - 4 when the cluster has 20-79 chunks
 - 8 when the cluster has 80+ chunks

Balancing is Resource Intensive

- Chunk migration requires copying all the data in the chunk from one shard to another.
- MongoDB can migrate only a single chunk at a time.
- MongoDB creates splits only after an insert operation.
- For these reasons, it is possible to define a balancing window to ensure the balancer will only run during scheduled times.

Chunk Migration Steps

1. The balancer process sends the moveChunk command to the source shard.
2. The source shard continues to process reads/writes for that chunk during the migration.
3. The destination shard requests documents in the chunk and begins receiving copies.
4. After receiving all documents, the destination shard receives any changes to the chunk.
5. Then the destination shard tells the config db that it has the chunk.
6. The destination shard will now handle all reads/writes.
7. The source shard deletes its copy of the chunk.

Concluding a Balancing Round

- Each chunk will move:
 - From the shard with the most chunks
 - To the shard with the fewest
- A balancing round ends when all shards differ by at most one chunk.

5.3 Shard Tags

Learning Objectives

Upon completing this module students should understand:

- The purpose for shard tags
- Advantages of using shard tags
- Potential drawbacks of shard tags

Tags - Overview

- Shard tags allow you to “tie” data to one or more shards.
- A shard tag describes a range of shard key values.
- If a chunk is in the shard tag range, it will live on a shard with that tag.

Example: DateTime

- Documents older than one year need to be kept, but are rarely used.
- You tag those ranges as “LTS” for Long Term Storage.
- Tag specific shards to hold LTS documents.
- These shards can be on cheaper, slower machines.
- Invest in high-performance servers for more frequently accessed data.

Example: Location

- You are required to keep certain data in its home country.
- You include the country in the shard tag.
- Maintain data centers within each country that house the appropriate shards.
- Meets the country requirement but allows all servers to be part of the same system.

Example: Premium Tier

- You have customers who want to pay for a “premium” tier.
- The shard key permits you to distinguish one customer’s documents from all others.
- Tag the document ranges for each customer so that their documents will be located on shards of the appropriate tier.
- Shards tagged as premium tier run on high performance servers.
- Other shards run on commodity hardware.
- See [Manage Shard Tags](#)¹²

Note:

- As customers move from one tier to another it will be necessary to execute commands that either add a given customer’s shard key range to the premium tag or remove that range from those tagged as “premium”.
- During balancing rounds, if the balancer detects that any chunks are not on the correct shards per configured tags, the balancer migrates chunks in tagged ranges to shards associated with those tags.
- After re-configuring tags with a shard key range, and associating it with a shard or shards, the cluster may take some time to balance the data among the shards.

See: [Tiered Storage Models in MongoDB: Optimizing Latency and Cost](#)¹³.

Tags - Caveats

- Because tagged chunks will only be on certain servers, if you tag more than those servers can handle, you’ll have a problem.
 - You’re not only worrying about your overall server load, you’re worrying about server load for each of your tags.
- Your chunks will evenly distribute themselves across the available chunks. You cannot control things more fine grained than your tags.

¹²<http://docs.mongodb.org/manual/tutorial/administer-shard-tags/>

¹³<http://blog.mongodb.org/post/85721044164/tiered-storage-models-in-mongodb-optimizing-latency>

5.4 Exercise: Setting Up a Sharded Cluster

Learning Objectives

Upon completing this module students should understand:

- How to set up a sharded cluster including:
 - Replica Sets as Shards
 - Config Servers
 - Mongos processes
- How to enable sharding for a database
- How to shard a collection
- How to determine where data will go

Our Sharded Cluster

- In this exercise, we will set up a cluster with 3 shards.
- Each shard will be a replica set with 3 members (including one arbiter).
- We will insert some data and see where it goes.

Sharded Cluster Configuration

- Three shards:
 1. A replica set on ports 27107, 27108, 27109
 2. A sharded replica set on ports 27117, 27118, 27119
 3. A sharded replica set on ports 27127, 27128, 27129
- Three config servers on ports 27217, 27218, 27219
- Two mongos servers at ports 27017 and 27018

Build Our Data Directories

On Linux or MacOS, run the following in the terminal to create the data directories we'll need.

```
mkdir -p ~/data/cluster/config/{c0,c1,c2}
mkdir -p ~/data/cluster/shard0/{m0,m1,arb}
mkdir -p ~/data/cluster/shard1/{m0,m1,arb}
mkdir -p ~/data/cluster/shard2/{m0,m1,arb}
mkdir -p ~/data/cluster/{s0,s1}
```

On Windows, run the following commands instead:

```
md c:\data\cluster\config\c0 c:\data\cluster\config\c1 c:\data\cluster\config\c2
md c:\data\cluster\shard0\m0 c:\data\cluster\shard0\m1 c:\data\cluster\shard0\arb
md c:\data\cluster\shard1\m0 c:\data\cluster\shard1\m1 c:\data\cluster\shard1\arb
md c:\data\cluster\shard2\m0 c:\data\cluster\shard2\m1 c:\data\cluster\shard2\arb
md c:\data\cluster\s0 c:\data\cluster\s1
```

Initiate a Replica Set

```
mongod --replSet shard0 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard0/m0 --logpath ~/data/cluster/shard0/m0/mongod.log \
  --fork --port 27107

mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard0/m1 --logpath ~/data/cluster/shard0/m1/mongod.log \
  --fork --port 27108

mongod --replSet shard0 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard0/arb --logpath ~/data/cluster/shard0/arb/mongod.log \
  --fork --port 27109

echo "cfg = {'_id': 'shard0', 'version': 1,
  'members': [{'_id' : 0, 'host' : 'localhost:27107'},
    {'_id': 1, 'host': 'localhost:27108' },
    {'_id': 2, 'host': 'localhost:27109', 'arbiterOnly': true}]};
rs.initiate(cfg);" | mongo --port 27107
```

Spin Up a Second Replica Set

```
mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard1/m0 --logpath ~/data/cluster/shard1/m0/mongod.log \
  --fork --port 27117

mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard1/m1 --logpath ~/data/cluster/shard1/m1/mongod.log \
  --fork --port 27118

mongod --replSet shard1 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard1/arb --logpath ~/data/cluster/shard1/arb/mongod.log \
  --fork --port 27119

echo "cfg = {'_id': 'shard1', 'version': 1,
  'members': [{'_id': 0, 'host': 'localhost:27117'},
```

```

        {'_id': 1, 'host': 'localhost:27118'},
        {'_id': 2, 'host': 'localhost:27119', 'arbiterOnly': true}]]];
rs.initiate(cfg);" | mongo --port 27117

```

A Third Replica Set

```

mongod --replSet shard2 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard2/m0 --logpath ~/data/cluster/shard2/m0/mongod.log \
  --fork --port 27127

mongod --replSet shard2 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard2/m1 --logpath ~/data/cluster/shard2/m1/mongod.log \
  --fork --port 27128

mongod --replSet shard2 --smallfiles --nojournal --noprealloc \
  --dbpath ~/data/cluster/shard2/arb --logpath ~/data/cluster/shard2/arb/mongod.log \
  --fork --port 27129

echo "cfg = {'_id': 'shard2', 'version': 1,
  'members': [{'_id': 0, 'host': 'localhost:27127'},
    {'_id': 1, 'host': 'localhost:27128'},
    {'_id': 2, 'host': 'localhost:27129', 'arbiterOnly': true}]]];
rs.initiate(cfg);" | mongo --port 27127

```

Status Check

- Now we have three replica sets running.
- We have one for each shard.
- They do not know about each other yet.
- To make them a sharded cluster we will:
 - Build our config databases
 - Launch our mongos processes
 - Add each shard to the cluster
- To benefit from this configuration we also need to:
 - Enable sharding for a database
 - Shard at least one collection within that database

Launch Config Servers

```
mongod --smallfiles --nojournal --noprealloc \
--dbpath $CONFIG/c0 --logpath ~/data/cluster/config/c0/mongod.log \
--fork --port 27227 --configsvr

mongod --smallfiles --nojournal --noprealloc \
--dbpath $CONFIG/c1 --logpath ~/data/cluster/config/c1/mongod.log \
--fork --port 27228 --configsvr

mongod --smallfiles --nojournal --noprealloc \
--dbpath $CONFIG/c2 --logpath ~/data/cluster/config/c2/mongod.log \
--fork --port 27229 --configsvr
```

Launch the Mongos Processes

Now our mongos's. We need to tell them about our config servers.

```
mongos --logpath ~/data/cluster/s0/mongos.log --fork --port 27017 \
--configdb localhost:27227,localhost:27228,localhost:27229

mongos --logpath ~/data/cluster/s1/mongos.log --fork --port 27018 \
--configdb localhost:27227,localhost:27228,localhost:27229
```

Add All Shards

```
echo 'sh.addShard( "shard0/localhost:27107" ); sh.addShard("shard1/localhost:27117" );
sh.addShard( "shard2/localhost:27127" ); sh.status()' | mongo
```

Note: Instead of doing this through a bash (or other) shell command, you may prefer to launch a mongo shell and issue each command individually.

Enable Sharding and Shard a Collection

Enable sharding for the test database, shard a collection, and insert some documents.

```
echo 'sh.enableSharding("test"); sh.shardCollection("test.testcol", { a : 1, b : 1 })' | mongo
```

```
echo 'for (i=0; i<10000; i++) { docArr = []; for (j=0; j<1000; j++) {  
docArr.push( { a : i, b : j, c : "Filler String 00000000000000000000  
0000000000000000000000000000000000000000000000000000000000000000  
0000000000000000000000000000000000000000000000000000000000000000"  
} ) }; db.testcol.insert(docArr) }' | mongo
```


Observe What Happens

Connect to either mongos using a mongo shell and frequently issue:

```
sh.status()
```

Note:

- Point out to the students that you can see chunks get created and moved to different shards.
- Also useful to have students run a query or two.

```
db.testcol.find( { a : { $lte : 100 } } ).explain()
```

6 Security

Security (page 122) An overview of security options for MongoDB.

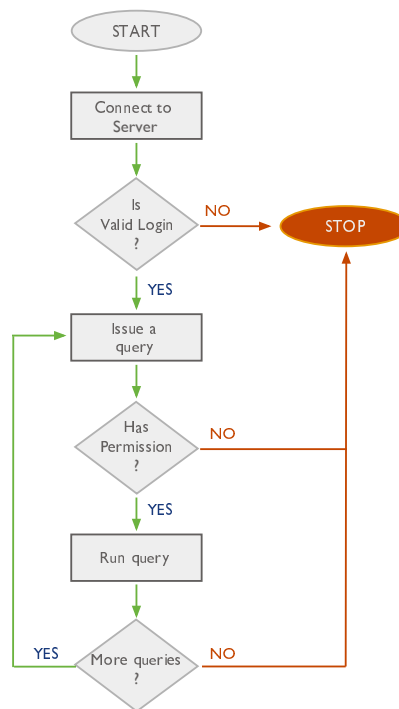
6.1 Security

Learning Objectives

Upon completing this module students should understand:

- Security options for MongoDB
- Basics of native auth for MongoDB
- User roles in MongoDB
- How to manage user roles in MongoDB

Overview



Note:

- You should only run MongoDB in a trusted environment.
- However, MongoDB offers security features from several angles:
 - Authentication: Is the client who they say they are?
 - Authorization: Is the client allowed to do this?
 - Network Exposure: Is this query or login coming from the place we expect?

- You are welcome to use any features you desire, or none.
 - All security is off by default.
-

Authentication Options

- MONGODB-CR Authentication (username & password)
- x.509 Authentication (using x.509 Certificates)
- Kerberos (through an Enterprise subscription)
- LDAP

Authorization via MongoDB

- Each user has a set of potential roles
 - read, readWrite, dbAdmin, etc.
- Each role applies to *one* database
 - A single user can have roles on each database
 - Some roles apply to all databases
 - You can also create custom roles.

Network Exposure Options

- bindIp limits the ip addresses the server listens on.
- Using a non-standard port can provide a layer of obscurity.
- MongoDB should still be run only in a trusted environment.

Encryption (SSL)

- MongoDB can be configured at build time to run with SSL.
- To get it, build from the source code with `–ssl`.
- Alternatively, use MongoDB Enterprise.
- Allows you to use public key encryption.
- You can also validate with x.509 certificates.

Native MongoDB Auth

- Uses the Challenge/Response mechanism
- Sometimes called MongoDB-CR
- Start a mongod instance with `--auth` to enable this feature
- You can initially login using localhost
 - Called the “localhost exception”.
 - Stops working when you create a user.

Note:

- Be careful to create a user who can create other users.
 - Otherwise you’ll be stuck with the users you initially create.
-

Exercise: Create an Admin User, Part 1

- Launch a mongo shell.
- Create a user with the role, `userAdminAnyDatabase`
- Use name “roland” and password “12345”.
- Enable this user to login on the admin database.

Note:

```
use admin
var role = "userAdminAnyDatabase"
var name = "roland"
var password = "12345"
newUserDoc = { user : name,
               pwd : password,
               roles : [ { role : role,
                           db : "admin" } ] }
db.createUser( newUserDoc )
exit
```

Exercise: Create an Admin User, Part 2

- Launch a mongo shell without logging in.
- Attempt to create a user.
- Exit the shell.
- Log in again as roland.
- Ensure that you can create a user.

Note:

```
mongo -u roland admin -p
```

Remember:

- Once a user is created, the localhost exception no longer applies.
 - If that first user can't create other users, you will be extremely limited in your ability
-

Using MongoDB Roles

- Each user logs in on *one* database.
- The user inputs their password on login.
 - Use the -u flag for username.
 - Use the -p flag to enter the password.
- userAdmins may create other users
- But they cannot read/write without other roles.

Note:

- Users must specify the db when logging in.
 - Trying to log into another db won't work
 - Users may follow -p with the password
 - They may also enter the password after hitting enter.
 - The `mongo` command will prompt for the password before launching the shell.
-

Exercise: Creating a readWrite User, Part 1

- Create a user named *vespa*.
- Give *vespa* readWrite access on the *test* and *druidia* databases.
- Create this user so that the login database is *druidia*.

Note:

```
mongo -u roland admin -p

use druidia
var name = "vespa"
var password = "12345"
vespa = { user : name,
          pwd : password,
          roles : [ { role : "readWrite", db : "druidia" },
                    { role : "readWrite", db : "test" } ] }
db.createUser(vespa)
exit
```

Exercise: Creating a readWrite User, Part 2

Log in with the user you just created.

Note:

```
mongo -u vespa test -p # won't work.
mongo -u vespa druidia -p # will work.

show collections // should be empty
db.foo.insert ( { a : 1 } )
db.foo.find() // see the doc
use test
db.foo.insert( { a : 1 } )
db.foo.find()
use test2
show collections // can't
db.foo.find() // can't
```

MongoDB Custom User Roles

- You can create custom user roles in MongoDB.
- You do this by modifying the `system.roles` collection.
- You can also inherit privileges from other roles into a given role.
- You won't remember how to do this, so if you need it, consult the docs¹⁴.

¹⁴<http://docs.mongodb.org/manual/core/security-introduction/>

7 Performance Troubleshooting

Performance Troubleshooting (page 127) An introduction to reporting and diagnostic tools for MongoDB.

7.1 Performance Troubleshooting

Learning Objectives

Upon completing this module students should understand basic performance troubleshooting techniques and tools including:

- `mongostat`
- `mongotop`
- `db.setProfilingLevel()`
- `db.currentOp()`
- `db.collection.stats()`
- `db.serverStatus()`

mongostat and mongotop

- `mongostat` samples a server every second.
 - See current ops, pagefaults, network traffic, etc.
 - Does not give a view into historic performance; use MMS for that.
- `mongotop` looks at the time spent on reads/writes in each collection.

Exercise: mongostat (setup)

In one window, perform the following commands.

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  arr = [];
  for (j=1; j<=1000; j++) {
    doc = { _id : (1000 * (i-1) + j), a : i, b : j, c : (1000 * (i-1) + j) };
    arr.push(doc)
  };
  db.testcol.insert(arr);
  var x = db.testcol.find( { b : 255 } );
  x.next();
  var x = db.testcol.find( { _id : 1000 * (i-1) + 255 } );
  x.next();
  var x = "asdf";
  db.testcol.update( { a : i, b : 255 }, { $set : { d : x.pad(1000) } });
  print(i)
}
```

Exercise: mongostat (run)

- In another window/tab, run mongostat.
- You will see:
 - Inserts
 - Queries
 - Updates

Exercise: mongostat (create index)

- In a third window, create an index when you see things slowing down:

```
db.testcol.ensureIndex( { a : 1, b : 1 } )
```

- Look at mongostat.
- Notice that things are going significantly faster.
- Then, let's drop that and build another index.

```
db.testcol.dropIndexes()  
db.testcol.ensureIndex( { b : 1, a : 1 } )
```

Exercise: mongotop

Perform the following then, in another window, run mongotop.

```
db.testcol.drop()  
for (i=1; i<=10000; i++) {  
  arr = [];  
  for (j=1; j<=1000; j++) {  
    doc = { _id: (1000*(i-1)+j), a: i, b: j, c: (1000*(i-1)+j) };  
    arr.push(doc)  
  };  
  db.testcol.insert(arr);  
  var x = db.testcol.find( {b: 255} ); x.next();  
  var x = db.testcol.find( { _id: 1000*(i-1)+255 } ); x.next();  
  var x = "asdf";  
  db.testcol.update( {a: i, b: 255}, { $set: {d: x.pad(1000)}});  
  print(i)  
}
```

Note: Direct the students to the fact that you can see the activity on the server for reads/writes/total.

db.currentOp()

- currentOp is a tool that asks what the db is doing at the moment.
- currentOp is useful for finding long-running processes.
- Fields of interest:
 - microsecs_running
 - op
 - query
 - lock
 - waitingForLock

Exercise: db.currentOp()

Do the following then, connect with a separate shell, and repeatedly run db.currentOp().

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  arr = [];
  for (j=1; j<=1000; j++) {
    doc = {_id: (1000*(i-1)+j), a: i, b: j, c: (1000*(i-1)+j)};
    arr.push(doc)
  };
  db.testcol.insert(arr);
  var x = db.testcol.find( {b: 255} ); x.next();
  var x = db.testcol.find( {_id: 1000*(i-1)+255 } ); x.next();
  var x = "asdf";
  db.testcol.update( {a: i, b: 255}, {$set: {d: x.pad(1000)}});
  print(i)
}
```

Note: Point out to students that the running time gets longer & longer, on average.

db.collection.stats()

- Used to view the current stats for a collection.
- Everything is in bytes; use the multiplier parameter to view in KB, MB, etc
- You can also use db.stats() to do this at scope of the entire database

Exercise: Using Collection Stats

Look at the output of the following:

```
db.testcol.drop()
db.testcol.insert( { a : 1 } )
db.testcol.stats()
var x = "asdf"
db.testcol2.insert( { a : x.pad(10000000) } )
db.testcol2.stats()
db.stats()
```

The Profiler

- Off by default.
- To reset, `db.setProfilerLevel(0)`
- At setting 1, it captures “slow” queries.
- You may define what “slow” is.
- Default is 100ms: `db.setProfilerLevel(1)`
- E.g., to capture 20 ms: `db.setProfilerLevel(1, 20)`

The Profiler (continued)

- If the profiler level is 2, it captures all queries.
 - This will severely impact performance.
 - Turns all reads into writes.
- Always turn the profiler off when done (set level to 0)
- Creates `db.system.profile` collection

Exercise: Exploring the Profiler

Perform the following, then look in your `db.system.profile`.

```
db.setProfilingLevel(0)
db.testcol.drop()
db.system.profile.drop()
db.setProfilingLevel(2)
db.testcol.insert( { a : 1 } )
db.testcol.find()
var x = "asdf"
db.testcol.insert( { a : x.pad(10000000) } ) // ~10 MB
db.setProfilingLevel(0)
db.system.profile.find().pretty()
```

Note:

- Mention to the students what the fields mean.

- Things to keep in mind:
 - op can be command, query, or update
 - ns is sometimes the db.collection namespace
 - * but sometimes db.\$cmd for commands
 - key updates refers to index keys
 - ts (timestamp) is useful for some queries if problems cluster.
-

db.serverStatus()

- Takes a snapshot of server status.
- By taking diffs, you can see system trends.
- Most of the data that MMS gets is from here.

Exercise: Using db.serverStatus()

- Open up two windows. In the first, type:

```
db.testcol.drop()  
var x = "asdf"  
for (i=0; i<=10000000; i++) {  
    db.testcol.insert( { a : x.pad(100000) } )  
}
```

- In the second window, type periodically:

```
var x = db.serverStatus(); x.metrics.document
```

Analyzing profiler data

- Enable the profiler at default settings.
- Run for 5 seconds.
- Slow operations are captured.
- The issue is there is not a proper index on the message field.
- Allow class to discover this as the data is examined.
- You will see how fast documents are getting inserted.
- It will be slow b/c the documents are big.

Performance Improvement Techniques

- Appropriate write concerns
- Bulk operations
- Good schema design
- Good Shard Key choice
- Good indexes

Performance Tips: Write Concern

- Increasing the write concern increases data safety.
- This will have an impact on performance, however.
- This is especially true when there are network issues.

Bulk Operations

- Using bulk operations can improve performance, especially when using write concern greater than 1.
- These enable the server to bulk write and bulk acknowledge.
- Can be done with both inserts and updates.

Exercise: Comparing bulk inserts with `mongostat`

Let's spin up a 3-member replica set:

```
mkdir -p /data/replset/{1,2,3}
mongod --logpath /data/replset/1/mongod.log \
--dbpath /data/replset/1 --replSet mySet --port 27017 --fork
mongod --logpath /data/replset/2/mongod.log \
--dbpath /data/replset/2 --replSet mySet --port 27018 --fork
mongod --logpath /data/replset/3/mongod.log \
--dbpath /data/replset/3 --replSet mySet --port 27019 --fork
```

```
echo "conf = {_id: 'mySet', members: [{_id: 0, host: 'localhost:27017'}, {_id: 1, host: 'localhost:27018'}, {_id: 2, host: 'localhost:27019'}]}; rs.initiate(conf)" | mongo
```

mongostat, bulk inserts with {w: 1}

- Perform the following, with writeConcern : 1 and no bulk inserts:

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  for (j=1; j<=1000; j++) {
    db.testcol.insert( { _id : (1000 * (i-1) + j),
                        a : i, b : j, c : (1000 * (i-1)+ j) },
                      { writeConcern : { w : 1 } } );
  };
  print(i);
}
```

Run mongostat and see how fast that happens.

Bulk inserts with {w: 3}

Increase the write concern to 3 (safer but slower):

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  for (j=1; j<=1000; j++) {
    db.testcol.insert(
      { _id: (1000 * (i-1) + j), a: i, b: j, c: (1000 * (i-1)+ j) },
      { writeConcern: { w: 3 } }
    );
  };
  print(i);
}
```

Again, run mongostat.

mongostat, bulk inserts with {w: 3}

- Finally, let's use bulk inserts to our advantage:
- Note that writeConcern is still { w: 3 }

```
db.testcol.drop()
for (i=1; i<=10000; i++) {
  arr = []
  for (j=1; j<=1000; j++) {
    arr.push(
      { _id: (1000 * (i-1) + j), a: i, b: j, c: (1000 * (i-1)+ j) }
    );
  };
  db.testcol.insert( arr, { writeConcern : { w : 3 } } );
  print(i);
}
```

Schema Design

- The structure of documents affects performance.
- Optimize for your application's read/write patterns.
- We want as few requests to the database as possible to perform a given application task.
- See the data modeling section for more information.

Shard Key Considerations

- Choose a shard key that distributes load across your cluster.
- Create a shard key such that only a small number of documents will have the same value.
- Create a shard key that has a high degree of randomness.
- Your shard key should enable a mongos to target a single shard for a given query.

Indexes and Performance

- Reads and writes that don't use an index will cripple performance.
- In compound indexes, order matters:
 - Sort on a field that comes before any range used in the index.
 - You can't skip fields; they must be used in order.
 - Revisit the indexing section for more detail.

8 Backup and Recovery

Backup and Recovery (page 135) An overview of backup options for MongoDB.

8.1 Backup and Recovery

Disasters Do Happen



Human Disasters



Terminology: RPO vs. RTO

- **Recovery Point Objective (RPO):** How much data can you afford to lose?
- **Recovery Time Objective (RTO):** How long can you afford to be off-line?

Terminology: DR vs. HA

- **Disaster Recovery (DR)**
- **High Availability (HA)**
- Distinct business requirements
- Technical solutions may converge

Quiz

- Q: What's the hardest thing about backups?
- A: Restoring them!
- **Regularly test that restoration works!**

Backup Options

- Document Level
 - Logical
 - mongodump, mongorestore
- File system level
 - Physical
 - Copy files
 - Volume/disk snapshots

Document Level Backups - mongodump

- Dumps collection to BSON files
- Mirrors your structure
- Can be run live or in offline mode
- Does not include indexes (rebuilt during restore)
- `--dbpath` for direct file access
- `--oplog` to record oplog while backing up
- `--query/filter` selective dump

mongodump

```
$ mongodump --help
Export MongoDB data to BSON files.
```

options:

<code>--help</code>	produce help message
<code>-v [--verbose]</code>	be more verbose (include multiple times for more verbosity e.g. <code>-vvvvv</code>)
<code>--version</code>	print the program's version and exit
<code>-h [--host] arg</code>	mongo host to connect to (/s1,s2 for
<code>--port arg</code>	server port. Can also use <code>--host hostname</code>
<code>-u [--username] arg</code>	username
<code>-p [--password] arg</code>	password
<code>--dbpath arg</code>	directly access mongod database files in path
<code>-d [--db] arg</code>	database to use
<code>-c [--collection] arg</code>	collection to use (some commands)
<code>-o [--out] arg</code>	(=dump) output directory or "-" for stdout

```
-q [ --query ] arg      json query
--oplog                  Use oplog for point-in-time snapshotting
```

File System Level

- **Must use journaling!**
- Copy /data/db files
- Or snapshot volume (e.g., LVM, SAN, EBS)
- *Seriously, always use journaling!*

Ensure Consistency

Flush RAM to disk and stop accepting writes:

- `db.fsyncLock()`
- Copy/Snapshot
- `db.fsyncUnlock()`

File System Backups: Pros and Cons

- Entire database
- Backup files will be large
- Fastest way to create a backup
- Fastest way to restore a backup

Document Level - mongorestore

- `mongorestore`
- `--oplogReplay` replay oplog to point-in-time

File System Restores

- All database files
- Selected databases or collections
- Replay Oplog

Backup Sharded Cluster

1. Stop Balancer (and wait) or no balancing window
2. Stop one config server (data R/O)
3. Backup Data (shards, config)
4. Restart config server
5. Resume Balancer

Restore Sharded Cluster

1. Dissimilar # shards to restore to
2. Different shard keys?
3. Selective restores
4. Consolidate shards
5. Changing addresses of config/shards

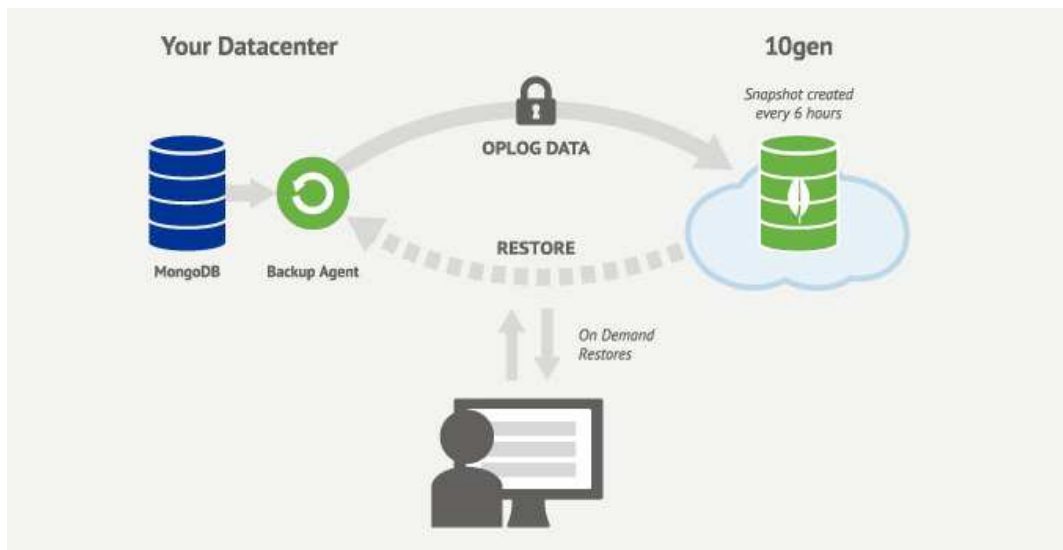
Tips and Tricks

- mongodump/mongorestore
 - --oplog[Replay]
 - --objcheck/--repair
 - --dbpath
 - --query/--filter
- bsondump
 - inspect data at console
- LVM snapshot time/space tradeoff
 - Multi-EBS (RAID) backup
 - clean up snapshots

Backup Options

- You can do it yourself as outlined in this section so far
- Or have the people who created MongoDB run your backups

MMS Backup



Sharded Clusters

- Balancer paused every 6 hours
- A no-op token is inserted across all shards, mongos instances, and config servers
- Oplog applied to replica sets until point in which token was inserted
- Provides a consistent state of database across shards

Under the Hood

- From the initial sync, we rebuild your data in our datacenters and take a snapshot
- We take snapshots every 6 hours
- Oplog is stored for 48 hours

Key Benefits

- Point in time backups
- Easy to restore
- Unlimited resources
- Fully managed

Point in Time Backups

- Oplog stored for 48 hours
- Restore your replica set to any point-in-time in the last 48 hours by creating a custom snapshot

Easy to Restore

- Pull from custom URL
- Push via scp

Unlimited Restores

- Confidence in your restore process
- Build development, QA, analytics environments without impacting production

Fully Managed

- Created by the engineers that build MongoDB
- No need to write or maintain custom backup scripts

Getting Started

- Go to <https://mms.mongodb.com> and sign up